SYSTEM V APPLICATION BINARY INTERFACE

MIPS® RISC Processor Supplement 3rd Edition

© 1990-1996 The Santa Cruz Operation, Inc. All rights reserved.

No part of this publication may be reproduced, transmitted, stored in a retrieval system, nor translated into any human or computer language, in any form or by any means, electronic, mechanical, magnetic, optical, chemical, manual, or otherwise, without the prior written permission of the copyright owner, The Santa Cruz Operation, Inc., 400 Encinal Street, Santa Cruz, California, 95060, USA. Copyright infringement is a serious matter under the United States and foreign Copyright Laws.

Information in this document is subject to change without notice and does not represent a commitment on the part of The Santa Cruz Operation, Inc.

SCO, the SCO logo, The Santa Cruz Operation, and UnixWare are trademarks or registered trademarks of The Santa Cruz Operation, Inc. in the USA and other countries. UNIX is a registered trademark in the USA and other countries, licensed exclusively through X/Open Company Limited. All other brand and product names are or may be trademarks of, and are used to identify products or services of, their respective owners.

SCO® UnixWare® is commercial computer software and, together with any related documentation, is subject to the restrictions on US Government use as set forth below. If this procurement is for a DOD agency, the following DFAR Restricted Rights Legend applies:

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the Government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of Rights in Technical Data and Computer Software Clause at DFARS 252.227-7013. Contractor/Manufacturer is The Santa Cruz Operation, Inc., 400 Encinal Street, Santa Cruz, CA 95060.

If this procurement is for a civilian government agency, this FAR Restricted Rights Legend applies:

RESTRICTED RIGHTS LEGEND: This computer software is submitted with restricted rights under Government Contract No. ______, if appropriate). It may not be used, reproduced, or disclosed by the Government except as provided in paragraph (g)(3)(i) of FAR Clause 52.227-14 alt III or as otherwise expressly stated in the contract. Contractor/Manufacturer is The Santa Cruz Operation, Inc., 400 Encinal Street, Santa Cruz, CA 95060.

If any copyrighted software accompanies this publication, it is licensed to the End User only for use in strict accordance with the End User License Agreement, which should be read carefully before commencing use of the software.

Document Version: 3 February 1996

Table of Contents

The MIPS Processor and System V ABI	1-1
How to Use the MIPS ABI Supplement	1-2
Evolution of the ABI Specification	1-2
Software Distribution Formats	2-1
Physical Distribution Media	2-1
Machine Interface	3-1
Processor Architecture	3-1
Data Representation	3-2
Byte Ordering	3-2
Fundamental Types	3-4
Aggregates and Unions	3-4
Bit-Fields	3-7
Function Calling Sequence	3-11
CPU Registers	3-11
Floating-Point Registers	3-13
The Stack Frame	3-15
Standard Called Function Rules	3-16
Argument Passing	3-17
Function Return Values	3-21
Operating System Interface	3-22
Virtual Address Space	3-22
Page Size	3-22
Virtual Address Assignments	3-22
Managing the Process Stack	3-24
Coding Guidelines	3-25
Exception Interface	3-25
Stack Backtracing	3-27
Process Initialization	3-28
Special Registers	3-29
Process Stack	3-30
Coding Examples	3-36
Code Model Overview	3-37
Position–Independent Function Prologue	3-38

TABLE OF CONTENTS

i

Data Objects	3-38
Position-Independent Load and Store	
Function Calls	3-40
Branching	3-42
C Stack Frame	3-46
Variable Argument List	3-46
Dynamic Allocation of Stack Space	3-49
ELF Header	4-1
Machine Information	4-1
Sections	4-3
Special Sections	4-6
Symbol Table	4-10
Symbol Values	4-10
Global Data Area	4-11
Register Information	4-14
Relocation	4-16
Relocation Types	4-16
Program Loading	5-1
Program Header	5-4
Segment Contents	5-4
Dynamic Linking	5-6
Dynamic Section	5-6
Shared Object Dependencies	5-8
Global Offset Table	5-8
Calling Position–Independent Functions	5-12
Symbols	5-13
Relocations	5-13
Ordering	5-14
Quickstart	5-15
Shared Object List	5-15
Conflict Section	5-17
System Library	6-1
Additional Entry Points	6-1
Support Routines	6-2
Global Data Symbols	6-3
Application Constraints	6-4
System Data Interfaces	6-5

TABLE OF CONTENTS

ii

Data Definitions	6-5
X Window Data Definitions	6-87
TCP/IP Data Definitions	6-152
Development Environment	7-1
Development Commands	7-1
PATH Access to Development Tools	7-1
Software Packaging Tools	7-1
System Headers	7-1
Static Archives	7-2
Execution Environment	8-1
Application Environment	8-1
The /dev Subtree	8-1

The MIPS Processor and System V ABI

The System V Application Binary Interface (ABI) defines a system interface for compiled application programs. It establishes a standard binary interface for application programs on systems that implement the interfaces defined in the *System V Interface Definition, Third Edition*. This includes systems that have implemented UNIX® System V, Release 4.

This document supplements the generic $System\ V\ ABI$, and it contains information specific to System V implementations built on the MIPS® RISC processor architecture. These two documents constitute the complete System V Application Binary Interface specification for systems that implement the MIPS RISC processor architecture.

INTRODUCTION 1-1

How to Use the MIPS ABI Supplement

This document contains information referenced in the generic *System V ABI* that may differ when System V is implemented on different processors. Therefore, the generic Application Binary Interface is the prime reference document, and this supplement is provided to fill gaps in that specification.

As with the *System V ABI*, this specification references other available reference documents, especially *MIPS RISC Architecture* (Copyright © 1990, MIPS Computer Systems, Inc., ISBN 0-13-584749-4). All the information referenced by this supplement is part of this specification, and just as binding as the requirements and data explicitly included here.

Evolution of the ABI Specification

The *System V Application Binary Interface* will evolve over time to address new technology and market requirements, and will be reissued at three-year intervals. Each new edition will contain extensions and additions to increase the capabilities of applications that conform to the *ABI*.

As with the *System V Interface Definition*, the *ABI* implements Level 1 and Level 2 support for its constituent parts. Level 1 support indicates a portion of the specification that will be supported indefinitely, while Level 2 support indicates a portion of the specification that may be withdrawn or altered when the next edition of the *System V ABI* is made available.

All components of this document and the generic *System V ABI* have Level 1 support unless they are explicitly labeled as Level 2.

Software Distribution Formats

Physical Distribution Media

The approved media for physical distribution of *ABI*-conforming software are listed below. *ABI*-conforming systems are not required to accept any of these media. A conforming system can install all software through its network connection.

- 60 MByte 1/4-inch cartridge tape in QIC-24 format¹
- 20 MByte 1/4-inch cartridge tape in QIC-120 format ²
- 1/2-inch, 9-track magnetic tape recorded at 1600 bpi
- 1.44 MByte 3 1/2-inch floppy disk: double-sided, 80 cylinders/side, 18 sectors/cylinder, 512 bytes/sector
- DDS Recording Format for Digital Audio Tape (DAT) DDS01 Rev E January, 1990 ³
- CD-ROM, ISO 9660 with Rockridge extensions

The QIC-24 cartridge tape data format is described in Serial Recorded Magnetic Tape Cartridge for Information Interchange (9 tracks, 10,000 FTPI, GCR, 60MB), Revision D, April 22, 1983. This document is available from the Quarter-Inch Committee (QIC) through Freeman Associates, 311 East Carillo St., Santa Barbara, CA 93101.

^{2.} The QIC-120 cartridge tape data format is described in *Serial Magnetic Tape Cartridge for Information Interchange, Fifteen Track, 0.250 in (6.30mm), 10,000 bpi (394 bpmm) Streaming Mode Group Code Recording,* Revision D, February 12, 1987. This document is available from the Quarter-Inch Committee (QIC) through Freeman Associates, 311 East Carillo St., Santa Barbara, CA 93101

^{3.} The DDS recording format is specified in ANSI Standard X3B5/88-185A, DDS Recording Format.

Machine Interface

Processor Architecture

MIPS RISC Architecture processor (Copyright © 1990, MIPS Computer Systems, Inc., ISBN 0-13-584749-4) defines the processor architecture for two separate Instruction Set Architectures (ISA), MIPS I and MIPS II. The MIPS I Instruction Set Architecture provides the architectural basis for this processor supplement to the generic ABI. Programs intended to execute directly on a processor that implements this ISA use the instruction set, instruction encodings, and instruction semantics of the architecture. Extensions available in the MIPS II ISA are explicitly not a part of this specification.

Three points deserve explicit mention.

- A program can assume all documented instructions exist.
- A program can assume all documented instructions work.
- A program can use only the instructions defined by the MIPS I ISA. In other words, *from a program's perspective*, the execution environment provides a complete and working implementation of the MIPS I ISA.

This does not mean that the underlying implementation provides all instructions in hardware, only that the instructions perform the specified operations and produce the specified results. The ABI neither places performance constraints on systems nor specifies what instructions must be implemented in hardware.

Some processors might support the MIPS I ISA as a subset, providing additional instructions or capabilities, e.g., the R6000 processor. Programs that use those capabilities explicitly do not conform to the MIPS ABI. Executing those programs on machines without the additional capabilities gives undefined behavior.

Data Representation

Byte Ordering

The architecture defines an 8-bit **byte**, 16-bit **halfword**, a 32-bit **word**, and a 64-bit **doubleword**. By convention there is also a 128-bit **quadword**. Byte ordering defines how the bytes that make up halfwords, words, doublewords, and quadwords are ordered in memory. Most significant byte (MSB) byte ordering, or big endian as it is sometimes called, means that the most significant byte is located in the lowest addressed byte position in a storage unit (byte 0).

Although the MIPS processor supports either big endian or little endian byte ordering, an ABI-conforming system must support big endian byte ordering.

The figures below illustrate the conventions for bit and byte numbering within various width storage units. These conventions hold for both integer data and floating-point data, where the most significant byte of a floating-point value holds the sign and at least the start of the exponent.

Figure 3-1: Bit and Byte Numbering in Halfwords

0	msb		1	lsb	
15		8	7		0

Figure 3-2: Bit and Byte Numbering in Words

0	msb	1	2	3 lsb
31	24	23 16	15 8	7 0

Figure 3-3: Bit and Byte Numbering in Doublewords

0	msb	1		2		3		
31	24	23	16	15	8	7		0
4		5		6		7	lsb	
31	24	23	16	15	8	7		0

Figure 3-4: Bit and Byte Numbering in Quadwords

0	msb	1		2		3		
31	24	23	16	15	8	7		0
4		5		6		7		
31	24	23	16	15	8	7		0
8		9		10		11		
31	24	23	16	15	8	7		0
12		13		14		15	lsb	
31	24	23	16	15	8	7	130	0

Fundamental Types

Figure 3-5 shows the correspondence between ANSI C's scalar types and the processor's.

Figure 3-5: Scalar Types

			Alignment	
Type	C	sizeof	(bytes)	MIPS
	char unsigned char	1	1	unsigned byte
	signed char	1	1	signed byte
	short signed short	2	2	signed halfword
	unsigned short	2	2	unsigned halfword
Integral	int signed int long signed long enum	4	4	signed word
	unsigned int unsigned long	4	4	unsigned word
Pointer	<pre>any-type * any-type (*)()</pre>	4	4	unsigned word
Floating-	float	4	4	single-precision
point	double	8	8	double-precision
	long double	8	8	double-precision

A null pointer (for all types) has the value zero.

Aggregates and Unions

Aggregates (structures and arrays) and unions assume the alignment of their most strictly aligned components. The size of any object, including aggregates and unions, is always a multiple of the alignment of the object. An array uses the same alignment as its elements. Structure and union objects can require padding to meet size and alignment constraints. The contents of any padding is undefined.

■ An entire structure or union object is aligned on the same boundary as its

most strictly aligned member.

- Each member is assigned to the lowest available offset with the appropriate alignment. This may require *internal padding*, depending on the previous member.
- If necessary, a structure's size is increased to make it a multiple of the alignment. This may require *tail padding*, depending on the last member.

In the following examples, byte offsets of the members appear in the upper left corners.

Figure 3-6: Structure Smaller Than a Word

```
struct {
      char c;
};
Byte aligned, sizeof is 1
```

Figure 3-7: No Padding

```
struct {
                     Word aligned, sizeof is 8
      char
              c;
                                          2
                         С
                                    d
                                                    s
      char
              d;
      short
             s;
                     4
      long
                                         n
};
```

Figure 3-8: Internal Padding

```
struct {
    char c;
    short s;
};
```

Halfword aligned, sizeof is 4

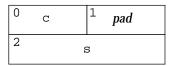


Figure 3-9: Internal and Tail Padding

```
struct {
    char c;
    double d;
    short s;
};
```

Doubleword aligned, sizeof is 24

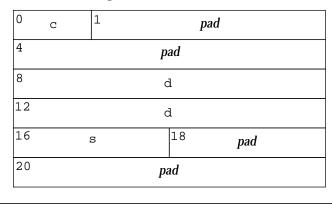
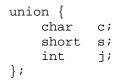
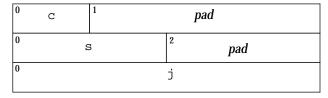


Figure 3-10: union Allocation



Word aligned, sizeof is 4



Bit-Fields

C struct and union definitions can have *bit-field*s, defining integral objects with a specified number of bits. Figure 3-11 lists the bit-field ranges.

Figure 3-11: Bit-Field Ranges

Bit-field Type	Width w	Range
signed char char unsigned char	1 to 8	-2 ^{w-1} to 2 ^{w-1} -1 0 to 2 ^{w-1} 0 to 2 ^{w-1}
signed short short unsigned short	1 to 16	-2^{-1} to 2^{w-1} -1 -2^{w-1} to 2^{w-1} -1 0 to 2^{w-1}
signed int int unsigned int	1 to 32	-2^{w-1} to 2^{w-1} -1 -2^{w-1} to 2^{w-1} -1 0 to 2^{w-1}
signed long long unsigned long	1 to 32	-2^{w-1} to 2^{w-1} -1 -2^{w-1} to 2^{w-1} -1 0 to 2^{w-1}

Plain bit-fields always have signed or unsigned values depending on whether the basic type is signed or unsigned. In particular, char bit-fields are unsigned while short, int, and long bit-fields are signed. A signed or unsigned modifier overrides the default type.

In a signed bit-field, the most significant bit is the sign bit; sign bit extension occurs when the bit-field is used in an expression. Unsigned bit-fields are treated as simple unsigned values.

Bit-fields follow the same size and alignment rules as other structure and union members, with the following additions:

■ Bit-fields are allocated from left to right (most to least significant).

- A bit-field must reside entirely in a storage unit that is appropriate for its declared type. Thus a bit-field never crosses its unit boundary. However, an unnamed bit-field of non-zero width is allocated in the smallest storage unit sufficient to hold the field, regardless of the defined type.
- Bit-fields can share a storage unit with other struct/union members, including members that are not bit-fields. Of course, struct members occupy different parts of the storage unit.
- Unnamed types of bit-fields do not affect the alignment of a structure or union, although member offsets of individual bit-fields follow the alignment constraints.



NOTE The X3J11 ANSI C specification only allows bit–fields of type int, with or without a signed or unsigned modifier.

Figures 3-12 through 3-17 provide examples that show the byte offsets of struct and union members in the upper left corners.

Figure 3-12: Bit Numbering

0x01020304

 $\begin{bmatrix} 0 & 01 & & & 1 & 02 & & 2 & & 3 & & 3 & & 04 \\ 31 & & & 24 & 23 & & & 16 & 15 & & & 8 & 7 & & 0 \end{bmatrix}$

Figure 3-13: Left-to-Right Allocation

```
struct {
    int    j:5;
    int    k:6;
    int    m:7;
};
```

Word aligned, sizeof is 4

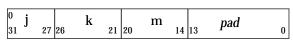


Figure 3-14: Boundary Alignment

```
struct {
    short s:9;
    int j:9;
    char c;
    short t:9;
    short u:9;
    char d;
};
```

Word aligned, sizeof is 12

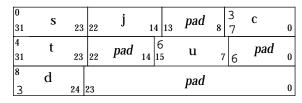


Figure 3-15: Storage Unit Sharing

```
struct {
    char c;
    short s:8;
};
```

Halfword aligned, sizeof is 2

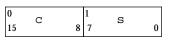


Figure 3-16: union Allocation

```
union {
   char c;
   short s:8;
};
```

Halfword aligned, sizeof is 2

0 15	С	8 7	pad	0
0 15	ន	8 7	pad	0

Figure 3-17: Unnamed Bit-Fields

```
Byte aligned, sizeof is 9
struct {
     char
               c;
                                                       :0
                                  С
                             31
                                      24 23
     int
               :0;
                             4
     char
               d;
                                                        :9
                                                                   pad
                                  d
                                            pad
                                                  16 15
                             31
                                      24 23
                                                              7 6
     short
               :9;
     char
               e;
                                  е
                             31
                                      24
     char
               :0;
};
```

As the examples show, int bit-fields (including signed and unsigned) pack more densely than smaller base types. One can use char and short bit-fields to force particular alignments, but int generally works better.

Function Calling Sequence

This section describes the standard function calling sequence, including stack frame layout, register usage, parameter passing, etc. The system libraries described in Chapter 6 require this calling sequence.

CPU Registers

The MIPS I ISA specifies 32 general purpose 32-bit registers; two special 32-bit registers that hold the results of multiplication and division instructions; and a 32-bit program counter register. The general registers have the names *\$0..\$31*. By convention, there is also a set of software names for some of the general registers. Figure 3-18 describes the conventions that constrain register usage. Figure 3-19 describes special CPU registers.



Not all register usage conventions are described. In particular, register usage con ventions in languages other than C are not included, nor are the effects of high optimization levels. These conventions do not affect the interface to the system libraries described in Chapter 6.

Figure 3-18: General CPU Registers

Register Name	Software Name	Use
<i>\$0</i>	zero	always has the value 0.
\$at	AT	temporary generally used by assembler.
\$2\$3	v0-v1	used for expression evaluations and to hold the integer and pointer type function return values.
\$4\$7	a0-a3	used for passing arguments to functions; values are not preserved across function calls. Additional arguments are passed on the stack, as described below.
\$8-\$15	t0-t7	temporary registers used for expression evaluation; values are not preserved across function calls.
\$16-\$23	s0-s7	saved registers; values are preserved across function calls.
\$24\$25	t8-t9	temporary registers used for expression evaluations; values are not preserved across function calls. When calling position independent functions $\$25$ must contain the address of the called function.
\$26-\$27	kt0-kt1	used only by the operating system.
\$28 or \$gp	gp	global pointer and context pointer.
\$29 or \$sp	sp	stack pointer.
\$30	s8	saved register (like s0-s7).
\$31	ra	return address. The return address is the location to which a function should return control.

Figure 3-19: Special CPU Registers

Register Name	Use
pc	program counter
hi	multiply/divide special register. Holds the most significant 32 bits of multiply or the remainder of a divide
lo	multiply/divide special register. Holds the least significant 32 bits of multiply or the quotient of a divide



Only registers \$16..\$23 and registers \$28.\$30 are preserved across a function call. Register \$28 is not preserved, however, when calling position independent code.

Floating-Point Registers

The MIPS ISA provides instruction encodings to move, load, and store values for up to four co-processors. Only co-processor 1 is specified in a *MIPS ABI* compliant system; the effect of moves, loads and stores to the other co-processors (0, 2, and 3) is unspecified.

Co-processor 1 adds 32 32-bit floating-point general registers and a 32-bit control/status register. Each even/odd pair of the 32 floating-point general registers can be used as either a 32-bit single-precision floating-point register or as a 64-bit double-precision floating-point register. For single-precision values, the even-numbered floating-point register holds the value. For double-precision values, the even-numbered floating-point register holds the least significant 32 bits of the value and the odd-numbered floating-point register holds the most significant 32 bits of the value. This is always true, regardless of the byte ordering conventions in use (big endian or little endian).

Floating-point data representation is that specified in *IEEE Standard for Binary Floating-Point Arithmetic*, ANSI/IEEE Standard 754-1985.

Figure 3-20 describes the conventions for using the floating-point registers.

Figure 3-20: Floating Point Registers

Register	
Name	Use
Sf0Sf2	used to hold floating-point type function results; single-precision uses <i>\$f0</i> and double-precision uses the register pair <i>\$f0\$f1</i> . <i>\$f2\$f3</i> return values that are not used in any part of this specification.
\$f4\$f10	temporary registers.
\$f12\$f14	used to pass the first two single- or double-precision actual arguments.
\$f16\$f18	temporary registers.
\$f20\$f30	saved registers; their values are preserved across function calls.
fcr31	control/status register. Contains control and status data for floating-point operations, including arithmetic rounding mode and the enabling of floating-point exceptions; it also indicates floating-point exceptions that occurred in the most recently executed instruction and all floating-point exceptions that have occurred since the register was cleared. This register is read/write and is described more fully in the



Only registers \$120.\$130 are preserved across a function call. All other floating-point registers can change across a function call. However, functions that use any of \$120.\$130 for single-precision operations only must still save and restore the corresponding odd-numbered register since the odd-numbered register contents are left undefined by single-precision operations.



There are other user visible registers in some implementations of the architecture, but these are explicitly not part of this processor supplement. A program that uses these registers is not *ABI* compliant and its behavior is undefined.

The Stack Frame

Each called function in a program allocates a stack frame on the run-time stack, if necessary. A frame is allocated for each non-leaf function and for each leaf function that requires stack storage. A non-leaf function is one that calls other function(s); a leaf function is one that does not itself make any function calls. Stack frames are allocated on the run-time stack; the stack grows downward from high addresses to low addresses.

Each stack frame has sufficient space allocated for:

- local variables and temporaries.
- saved general registers. Space is allocated only for those registers that need to be saved. For non-leaf function, \$31 must be saved. If any of \$16..\$23 or \$29..\$31 is changed within the called function, it must be saved in the stack frame before use and restored from the stack frame before return from the function. Registers are saved in numerical order, with higher numbered registers saved in higher memory addresses. The register save area must be doubleword (8 byte) aligned.
- saved floating-point registers. Space is allocated only for those registers that need to be saved. If any of *Sf20..Sf30* is changed within the called function, it must be saved in the stack frame before use and restored from the stack frame before return from the function. Both even- and odd-numbered registers must be saved and restored, even if only single-precision operations are performed since the single-precision operations leave the odd-numbered register contents undefined. The floating-point register save area must be doubleword (8 byte) aligned.
- function call argument area. In a non-leaf function the maximum number of bytes of arguments used to call other functions from the non-leaf function must be allocated. However, at least four words (16 bytes) must always be reserved, even if the maximum number of arguments to any called function is fewer than four words.
- alignment. Although the architecture requires only word alignment, soft-

ware convention and the operating system require every stack frame to be doubleword (8 byte) aligned.

A function allocates a stack frame by subtracting the size of the stack frame from *Ssp* on entry to the function. This *Ssp* adjustment must occur before *Ssp* is used within the function and prior to any jump or branch instructions.

Figure 3-21: Stack Frame

Base	Offset	Contents	Frame
		unspecified	High addresses
		variable size	
		(if present)	
		incoming arguments	Previous
	+16	passed in stack frame	
		space for incoming	
old \$sp	0+	arguments 1-4	
		locals and	
		temporaries	
		general register	
		save area	Current
		floating-point	
		register save area	
		argument	
\$sp	+0	build area	Low addresses

The corresponding restoration of *Ssp* at the end of a function must occur after any jump or branch instructions except prior to the jump instruction that returns from the function. It can also occupy the branch delay slot of the jump instruction that returns from the function.

Standard Called Function Rules

By convention, there is a set of rules that must be followed by every function that allocates a stack frame. Following this set of rules ensures that, given an arbitrary program counter, return address register \$31, and stack pointer, there is a deterministic way of performing stack backtracing. These rules also make possible programs that translate already compiled absolute code into position-independent

code. See Coding Examples in this chapter.

Within a function that allocates a stack frame, the following rules must be observed:

- In position-independent code that calculates a new value for the *gp* register, the calculation must occur in the first three instructions of the function. One possible optimization is the total elimination of this calculation; a local function called from within a position-independent module guarantees that the context pointer *gp* already points to the global offset table. The calculation must occur in the first basic block of the function.
- The stack pointer must be adjusted to allocate the stack frame before any other use of the stack pointer register.
- At most, one frame pointer can be used in the function. Use of a frame pointer is identified if the stack pointer value is moved into another register, after the stack pointer has been adjusted to allocate the stack frame. This use of a frame pointer must occur within the first basic block of the function before any branch or jump instructions, or in the delay slot of the first branch or jump instruction in the function.
- There is only one exit from a function that contains a stack adjustment: a jump register instruction that transfers control to the location in the return address register \$31. This instruction, including the contents of its branch delay slot, mark the end of function.
- The deallocation of the stack frame, which is done by adjusting the stack pointer value, must occur once and in the last basic block of the function. The last basic block of a function includes all of the non control-transfer instructions immediately prior to the function exit, including the branch delay slot.

Argument Passing

Arguments are passed to a function in a combination of integer general registers, floating-point registers, and the stack. The number of arguments, their type, and their relative position in the argument list of the calling function determines the mix of registers and memory used to pass arguments. General registers \$4..\$7 and floating-point registers \$f12 and \$f14 pass the first few arguments in registers. Double-precision floating-point arguments are passed in the register pairs \$f12, \$f13 and \$f14, \$f15; single-precision floating-point arguments are passed in registers \$f12 and \$f14.



These argument passing rules apply only to languages such as C that do not do dynamic stack allocation of structures and arrays. Ada is an example of a language that does dynamic stack allocation of structures and arrays.

In determining which register, if any, an argument goes into, take into account the following considerations:

- All integer-valued arguments are passed as 32-bit words, with signed or unsigned bytes and halfwords expanded (promoted) as necessary.
- If the called function returns a structure or union, the caller passes the address of an area that is large enough to hold the structure to the function in *\$4*. The called function copies the returned structure into this area before it returns. This address becomes the first argument to the function for the purposes of argument register allocation and all user arguments are shifted down by one.
- Despite the fact that some or all of the arguments to a function are passed in registers, always allocate space on the stack for all arguments. This stack space should be a structure large enough to contain all the arguments, aligned according to normal structure rules (after promotion and structure return pointer insertion). The locations within the stack frame used for arguments are called the home locations.
- At the call site to a function defined with an ellipsis in its prototype, the normal calling conventions apply up until the first argument corresponding to where the ellipsis occurs in the parameter list. If, in the absence of the prototype, this argument and any following arguments would have been passed in floating-point registers, they are instead passed in integer registers. Arguments passed in integer registers are not affected by the ellipsis.

This is the case only for calls to functions which have prototypes containing an ellipsis. A function without a prototype or without an ellipsis in a prototype is called using the normal argument passing conventions.

3-18

- When the first argument is integral, the remaining arguments are passed in the integer registers.
- Structures are passed as if they were very wide integers with their size rounded up to an integral number of words. The fill bits necessary for rounding up are undefined.
- A structure can be split so a portion is passed in registers and the remainder passed on the stack. In this case, the first words are passed in \$4, \$5, \$6, and \$7 as needed, with additional words passed on the stack.
- Unions are considered structures.

The rules that determine which arguments go into registers and which ones must be passed on the stack are most easily explained by considering the list of arguments as a structure, aligned according to normal structure rules. Mapping of this structure into the combination of stack and registers is as follows: up to two leading floating-point arguments can be passed in *\$f12* and *\$f14*; everything else with a structure offset greater than or equal to 16 is passed on the stack. The remainder of the arguments are passed in *\$4..\$7* based on their structure offset. Holes left in the structure for alignment are unused, whether in registers or in the stack.

The following examples in Figure 3-22 give a representative sampling of the mix of registers and stack used for passing arguments, where d represents double-precision floating-point values, s represents single-precision floating-point values, and n represents integers or pointers. This list is not exhaustive.

See the section "Variable Argument List" later in this section for more information about variable argument lists.

Figure 3-22: Examples of Argument Passing

Argument List	Register and Stack Assignments
d1, d2	\$f12, \$f14
s1, s2	\$f12, \$f14
s1, d1	\$f12, \$f14
d1, s1	\$f12, \$f14
n1, n2, n3, n4	\$4, \$5, \$6, \$7
d1, n1, d2	\$f12, \$6, stack
d1, n1, n2	\$f12, \$6, \$7
s1, n1, n2	\$f12, \$5, \$6
n1, n2, n3, d1	\$4, \$5, \$6, stack
n1, n2, n3, s1	\$4, \$5, \$6, \$7
n1, n2, d1	\$4, \$5, (\$6, \$7)
n1, d1	\$4, (\$6, \$7)
s1, s2, s3, s4	\$f12, \$f14, \$6, \$7
s1, n1, s2, n2	\$f12, \$5, \$6, \$7
d1, s1, s2	\$f12, \$f14, \$6
s1, s2, d1	\$f12, \$f14, (\$6, \$7)
n1, s1, n2, s2	\$4, \$5, \$6, \$7
n1, s1, n2, n3	\$4, \$5, \$6, \$7
n1, n2, s1, n3	<i>\$4, \$5, \$6, \$7</i>

In the following examples, an ellipsis appears in the second argument slot.

```
      n1, d1, d2
      $4, ($6, $7), stack

      $1, n1
      $f12, $5

      $1, n1, d1
      $f12, $5, ($6, $7)

      $1, n1
      $f12, $6

      $12, $6, $12, $6, $12, $6, $12, $6, $12, $16
```

Function Return Values

A function can return no value, an integral or pointer value, a floating-point value (single- or double-precision), or a structure; unions are treated the same as structures.

A function that returns no value (also called procedures or void functions) puts no particular value in any register.

A function that returns an integral or pointer value places its result in register \$2.

A function that returns a floating-point value places its result in floating-point register *\$f0*. Floating-point registers can hold single- or double-precision values.

The caller to a function that returns a structure or a union passes the address of an area large enough to hold the structure in register *\$4*. Before the function returns to its caller, it will typically copy the return structure to the area in memory pointed to by *\$4*; the function also returns a pointer to the returned structure in register *\$2*. Having the caller supply the return object's space allows re-entrancy.



Structures and unions in this context have fixed sizes. The *ABI* does not specify how to handle variable sized objects.

Both the calling and the called function must cooperate to pass the return value successfully:

- The calling function must supply space for the return value and pass its address in the stack frame.
- The called function must use the address from the frame and copy the return value to the object so supplied.

Failure of either side to meet its obligations leads to undefined program behavior.



These rules for function return values apply to languages such as C, but do not necessarily apply to other languages. Ada is one language to which the rules do not apply.

Operating System Interface

Virtual Address Space

Processes execute in a 31-bit virtual address space with addresses from 0 to 2^{31} - 1. Memory management hardware translates virtual addresses to physical addresses, which hides physical addressing and allows a process to run anywhere in the real memory of the system. Processes typically begin with three logical segments, commonly called text, data, and stack. As Chapter 5 describes, dynamic linking creates more segments during execution, and a process can create additional segments for itself with system services.

Page Size

Memory is organized by pages, which are the smallest units of memory allocation in the system. Page size can vary from one system to another, depending on the processor, memory management unit, and system configuration. Processes can call sysconf(BA_OS) to determine the current page size.

Virtual Address Assignments

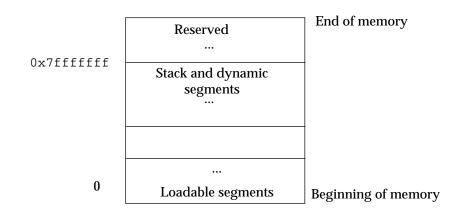
Although processes have the full 31-bit address space available, several factors limit the size of a process.

- The system reserves a configuration-dependent amount of virtual space.
- A tunable configuration parameter limits process size.
- A process that requires more memory than is available in system physical memory and secondary storage cannot run. Although some physical memory must be present to run any process, the system can execute processes that are bigger than physical memory, paging them to and from secondary storage. Nonetheless, both physical memory and secondary storage are shared resources. System load, which can vary from one program execution to the next, affects the available amount of memory.

Figure 3-23 shows virtual address configuration. The terms used in the figure are:

- The loadable segments of the processes can begin at 0. The exact addresses depend on the executable file format [see Chapters 4 and 5].
- The stack and dynamic segments reside below the reserved area. Processes can control the amount of virtual memory allotted for stack space, as described below.
- The reserved area resides at the top of virtual space.

Figure 3-23: Virtual Address Configuration



As Figure 3-23 shows, the system reserves the high end of virtual space, with the stack and dynamic segments of a process below that. Although the exact boundary between the reserved area and a process depends on system configuration, the reserved area will not consume more than 4 MBytes from the virtual address space. Thus the user virtual address range has a minimum upper bound of 0x7f-bfffff. Individual systems can reserve less space, increasing the processes virtual memory range. More information follows in 'Managing the Process Stack.'

Although applications can control their memory assignments, the typical arrangement follows the diagram in Figure 3-23. Loadable segments reside at low ad-

Coding Guidelines

Operating system facilities, such as mmap(KE_OS), allow a process to establish address mappings in two ways. First, the program can let the system choose an address. Second, the program can force the system to use an address the program supplies. This second alternative can cause application portability problems, because the requested address might not always be available. Differences in virtual address space between different architectures can be particularly troublesome, although the same problems can arise within a single architecture.

Process address spaces typically have three segment areas that can change size from one execution to the next: the stack [through setrlimit(BA_OS)], the data segment [through malloc(BA_OS)], and the dynamic segment area [through mmap(-KE_OS)]. Changes in one area can affect the virtual addresses available for another. Consequently, an address that is available in one process execution might not be available in the next. A program that uses mmap(KE_OS) to request a mapping at a specific address could work in some environments and fail in others. For this reason, programs that establish a mapping in their address space should use an address provided by the system.

Despite these warnings about requesting specific addresses, the facility can be used properly. For example, a multiprocess application can map several files into the address space of each process and build relative pointers among the data in the files. This is done by having each process specify a certain amount of memory at an address chosen by the system. After each process receives its own address from the system, it can map the desired files into memory, at specific addresses within the original area. This collection of mappings could be at different addresses in each process but their *relative* positions would be fixed. Without the ability to specify addresses, the application cannot build shared data structures, because the *relative* positions for files in each process would be unpredictable.

Exception Interface

In MIPS architecture, there are two execution modes: user and kernel. Processes run in user mode and the operating system kernel runs in kernel mode. The processor changes mode to handle precise or interrupting *exceptions*. Precise exceptions, which result from instruction execution, are explicitly generated by a process. This section, therefore, specifies those exception types with defined behavior.

An exception results in the operating system kernel taking some action. After handling the exception the kernel restarts the user process. It is not possible to determine that an exception took place, except by apparent slower execution. Some exceptions are considered errors, however, and cannot be handled by the operating system kernel. These exceptions cause either process termination or, if signal

catching is enabled, send a signal to the user process (see signal(BA_OS)).

Figure 3-24 lists the correspondence between exceptions and the signals specified by signal(BA_OS).

Figure 3-24: Hardware Exceptions and Signals

Exception	Signal
TLB modification	SIGBUS
Read TLB miss	SIGSEGV
Read TLB miss	SIGBUS
Write TLB miss	SIGSEGV
Read Address Error	SIGBUS
Write Address Error	SIGBUS
Instruction Bus Error	SIGBUS
Data Bus Error	SIGBUS
Syscall	SIGSYS
Breakpoint	SIGTRAP
Reserved Instruction	SIGILL
Coprocessor Unusable	SIGILL
Arithmetic Overflow	SIGFPE



A Read TLB miss generates a SIGSEGV signal when unmapped memory is NOTE accessed. A Read TLB miss generates a SIGBUS signal when mapped, but otherwise inaccessible memory is accessed. In other words, a SIGBUS is generated on a protection fault while a ${\tt SIGSEGV}$ is generated on a segmentation fault.

Floating-point instructions exist in the architecture, and can be implemented either in hardware or software. If the Coprocessor Unusable exception occurs because of a coprocessor 1 instruction, the process receives no signal. Instead, the system intercepts the exception, emulates the instruction, and returns control to the process. A process receives SIGILL for the Coprocessor Unusable exception only when the

accessed coprocessor is not present and when it is not coprocessor 1.

System calls, or requests for operating system services, use the Syscall exception for low level implementation. Normally, system calls do not generate a signal, but SIGSYS can occur in some error conditions.



The *ABI* does not define the implementation of individual system calls. Instead, programs should use the system libraries described in Chapter 6. Programs with embedded system call instructions do not conform to the *ABI*.

Stack Backtracing

There are standard called function rules for functions that allocate a stack frame and because the operating system kernel initializes the return address register \$31 to zero when starting a user program it is possible to trace back through any arbitrarily nested function calls. The following algorithm, which takes the set of general registers plus the program counter as input, produces the values the registers had at the most recent function call. Of course, only the saved registers plus gp, sp, ra, and pc can be reconstructed.

- Scan each instruction starting at the current program counter, going backwards. The compiler and linker must guarantee that a jump register to return address instruction will always precede each text section.
 - If the instruction is of the form "move Sr, Sp" or "addu Sr, SSp, SO, then the register Sr may be a frame pointer. The algorithm remembers the current instruction so it can continue its backward scan.

Then, it scans forward until it sees the "jr *ra*" instruction that marks the end of the current function.

Next, it scans backwards searching for an instruction of the form "move *sp*, *Sr*" or "addu *Ssp*, *Sr*, *S0*". This scan terminates when such an instruction is found or the branch or jump instruction that marks the beginning of the last basic block.

If a move or addu instruction of the kind described above was found, remember the register number of *Sr* as the frame pointer. Otherwise, *Sr* is not the frame pointer.

The algorithm should return to its original backwards scan starting with the instruction preceding the one remembered above.

■ If the instruction is a stack pointer decrement, exit the scan.

- If the instruction is a jump register to return address, exit the scan.
- If the last examined instruction is a jump register to the return address, it is the end of the previous function and no stack frame has yet been allocated for the current function. The address from which the current function was called is in the return address register minus eight. The other save registers had their current values when this function was called, so just return their current values.
- The stack decrement instruction must occur in the first basic block of the function. The amount of stack decrement is the size of the stack frame.
- Examine each instruction at increasing program addresses. If any instruction is a store of save registers *\$16-\$23*, *\$28*, *\$30*, or *\$31* through the frame pointer (or stack pointer if no frame pointer was used), then record its value by reading from the stack frame.
- Stop after examining the instruction in the first branch delay slot encountered. This marks the end of the first basic block.
- The frame pointer is the stack pointer value at the time the current function was called (or the stack pointer if no frame pointer was used) plus the size of the stack frame.
- The address from which the function is called is either the return address register value minus eight or, if the return address was saved on the stack, the saved value minus eight.

Process Initialization

This section describes the machine state that exec(BA_OS) creates for "infant" processes, including argument passing, register usage, stack frame layout, etc. Programming language systems use this initial program state to establish a standard environment for their application programs. For example, a C program begins ex-

ecution at a function named main, conventionally declared as follows:

```
extern int main(int argc, char *argv[], char *envp[]);
```

where argc is a non-negative argument count; argv is an array of argument strings, with argv[argc]==0; and envp is an array of environment strings, also terminated by a null pointer.

Although this section does not describe C program initialization, it does provide the information necessary to implement a call to main or to the entry point for a program in any other language.

Special Registers

As the architecture defines, two registers control and monitor the processor: the status register (SR) and the floating-point control and status register (csr). Applications cannot access the SR directly; they run in *user mode*. Instructions to read and write the SR are privileged. No fields in the SR affect user program behavior, except that the program can assume that coprocessor 1 instructions work as documented and that the user program executes in user mode with the possibility that interrupts are enabled. Nothing more should be inferred about the contents of the SR.

Figure 3-25 lists the initial values of the floating-point control and status register provided in the architecture

Figure 3-25: Floating-Point Control and Status Register Fields

led
•

The *ABI* specifies that coprocessor 1 always exists and that coprocessor 1 instructions (floating-point instructions) work as documented. Programs that directly ex-

ecute coprocessor 0, 2, or 3 instructions do not conform to the ABI. Individual system implementations may use one of these coprocessors under control of the system software, not the application.

Process Stack

When a process receives control, its stack holds the arguments and environment from exec(BA_OS). Figure 3-26 shows the initial process stack.

Figure 3-26: Initial Process Stack

\$sp+0

Unspecified	High addresses
Information block, including	8
argument strings	
environment strings	
auxiliary information	
duxinary information	
(size varies)	
Unspecified	
Null auxiliary vector entry	
Auxiliary vector	
(2-word entries)	
0 word	
Environment pointers	
(one word each)	
0 word	
Argument pointers	•
(Argument count words)	Low addresses

Argument strings, environment strings, and auxiliary information do not appear in a specific order with the information block. The system may leave an unspecified amount of memory between a null auxiliary vector entry and the beginning of an information block.

Except as shown below, general integer and floating-point register values are unspecified at process entry. Consequently, a program that requires specific register values must set them explicitly during process initialization. It should *not* rely on the operating system to set all registers to 0.

The registers listed below have the specified contents at process entry:

- \$2 A non-zero value specifies a function pointer the application should register with atexit(BA_OS). If \$2 contains zero, no action is required.
- Ssp The stack pointer holds the address of the bottom of the stack, which must be doubleword (8 byte) aligned.
- \$31 The return address register is set to zero so that programs that search backward through stack frames (stack backtracing) recognize the last stack frame, that is, a stack frame with a zero in the saved \$31 slot.

Every process has a stack, but the system does not define a fixed stack address. Furthermore, a program's stack address can change from one system to another even from one process invocation to another. Thus the process initialization code must use the stack address in *\$sp*. Data in the stack segment at addresses below the stack pointer contain undefined values.

Whereas the argument and environment vectors transmit information from one application program to another, the auxiliary vector conveys information from the operating system to the program. This vector is an array of the structures shown in Figure 3-27, interpreted according to the a_type member.

Figure 3-27: Auxillary Vector

```
typedef struct
{
    int a_type;
    union {
        long a_val;
        void *a_ptr;
        void (*a_fcn)();
    } a_un;
} auxv_t;
```

Figure 3-28: Auxillary Vector Types, a_type

Name	Value	a_un
AT_NULL	0	ignored
AT_IGNORE	1	ignored
AT_EXECFD	2	a_val
AT_PHDR	3	a_ptr
AT_PHENT	4	a_val
AT_PHNUM	5	a_val
AT_PAGESZ	6	a_val
AT_BASE	7	a_ptr
AT_FLAGS	8	a_val
AT_ENTRY	9	a_ptr
AT_NOTELF	10	a_val
AT_UID	11	a_val
AT_EUID	12	a_val
AT_GID	13	a_val
AT_EGID	14	a_val

The auxiliary vector types (a_type) shown in Figure 3-28 are explained in the paragraphs below:

AT_NULL	The auxiliary vector has no fixed length; instead the
	a type member of the last entry has this value.

AT_IGNORE This type indicates the entry has no meaning. The corresponding value of a_un is undefined.

AT_EXECFD As Chapter 5 describes, exec(BA_OS) can pass control to an interpreter program. When this happens, the system places either an entry of type AT_EXECFD or type AT_PHDR in the auxiliary vector. The entry for type AT_EXECFD uses the a_val member to contain a file descriptor open to read the application program object file.

AT_PHDR	Under some conditions, the system creates the memory image of the application program before passing control to the interpreter program. When this happens, the a_ptr member of the AT_PHDR entry tells the interpreter where to find the program header table in the memory image. If the AT_PHDR entry is present, entries of types AT_PHENT, AT_PHNUM, and AT_ENTRY are also present. See Chapter 5 in both the <i>System V ABI</i> and the processor supplement for more information about the program header table.
AT_PHENT	The <code>a_val</code> member of this entry holds the size, in bytes, of one entry in the program header table to which the <code>AT_PHDR</code> entry points.
AT_PHNUM	The a_val $$ member of this entry holds the number of entries in the program header table to which the AT_PHDR entry points.
AT_PAGESZ	If present, the a_val member of this entry gives the system page size, in bytes. The same information also is available through ${\tt sysconf}({\tt BA}_{\tt OS})$.
AT_BASE	The a_ptr member of this entry holds the base address at which the interpreter program was loaded into memory. See "Program Header" in the <i>System V ABI</i> for more information about the base address.
AT_FLAGS	If present, the a_val member of this entry holds one-bit flags. Bits with undefined semantics are set to zero.
AT_ENTRY	The a_ptr member of this entry holds the entry point of the application program to which the interpreter program should transfer control.
AT_NOTELF	The a_val member of this entry is zero if the executable is in ELF format as described in Chapter 4. It is non-zero if the executable is in MIPS XCOFF format.
AT_UID	If present, the a_val member of this entry holds the actual user id of the current user.
AT_EUID	If present, the <code>a_val</code> member of this entry holds the effective user id of the current user.
AT_GID	If present, the <code>a_val</code> member of this entry holds the actual

group id of the current user.

AT_EGID If present, the a_val member of this entry holds the effective group id of the current user.

Other auxiliary vector types are reserved. Currently, no flag definitions exist for AT_FLAGS. Nonetheless, bits under the 0xff000000 mask are reserved for system semantics.

In the following example, the stack resides below 0x7fc00000, growing toward lower addresses. The process receives three arguments:

- ср
- src
- dst

It also inherits two environment strings. (The example does not show a fully configured execution environment).

- HOME=/home/dir
- PATH=/home/dir/bin:/usr/bin:

Its auxiliary vector holds one non-null entry, a file descriptor for the executable file.

13

The initialization sequence preserves the stack pointer's doubleword (8 byte) alignment.

Figure 3-29: Example Process Stack

	n	:	\0	pad	High addresses
	r	/	b	i	
	:	/	u	s	
0x7fbffff0	/	b	i	n	
	/	d	i	r	
	h	0	m	е	
	Т	Н	=	/	
0x7fbfffe0	r	\0	Р	A	
	е	/	d	i	
	/	h	0	m	
	0	M	E	II	
0x7fbfffd0	s	t	\0	H	
	r	С	\0	d	
	С	р	\0	S	
			0		
0x7fbfffc0			0		
	13				
	2			Auxiliary vector	
	0				
0x7fbfffb0	0x7fbfffe2				
	0x7fbfffd3			Environment vector	
	0				
	0x7fbfffcf				
0x7fbfffa0		0x7fl	offfc	.b	
	0x7fbfffc8		:8	Argument vector	
\$sp+0 0x7fbfff90			3		Argument count
					Low addresses

Coding Examples

This section discusses example code sequences for basic operations such as calling functions, accessing static objects, and transferring control from one part of a program to another. Previous sections discuss how a program uses the machine or theoperating system, and specify what a program can or cannot assume about the execution environment. Unlike the previous material, the information here illustrates how operations *can* be done, not how they *must* be done.

As before, examples use the ANSI C language. Other programming languages may use the same conventions displayed below, but failure to do so does *not* prevent a program from conforming to the *ABI*. Two main object code models are available.

Absolute code

Instructions can hold absolute addresses under this model. To execute properly, the program must be loaded at a specific virtual address, making the program absolute addresses coincide with the process virtual addresses.

Position-independent code

Instructions under this model hold relative addresses, *not* absolute addresses. Consequently, the code is not tied to a specific load address, allowing it to execute properly at various positions in virtual memory.

The following sections describe the differences between absolute code and position-independent code. Code sequences for the models (when different) appear together, allowing easier comparison



The examples below show code fragments with various simplifications. They are intended to explain addressing modes, not to show optimal code sequences or to reproduce compiler output or actual assembler syntax.



When other sections of this document show assembly language code sequences, they typically show only the absolute versions. Information in this section explains how position–independent code would alter the examples.

Code Model Overview

When the system creates a process image, the executable file portion of the process has fixed addresses, and the system chooses shared object library virtual addresses to avoid conflicts with other segments in the process. To maximize text sharing, shared objects conventionally use position-independent code, in which instructions contain no absolute addresses. Shared object text segments can be loaded at various virtual addresses without changing the segment images. Thus multiple processes can share a single shared object text segment, even though the segment resides at a different virtual address in each process.

Position-independent code relies on two techniques:

- Control transfer instructions hold addresses relative to the program counter (PC). A PC-relative branch or function call computes its destination address in terms of the current program counter, *not* relative to any absolute address. If the target location exceeds the allowable offset for PCrelative addressing, the program requires an absolute address.
- When the program requires an absolute address, it computes the desired value. Instead of embedding absolute addresses in the the instructions, the compiler generates code to calculate an absolute address during execution.

Because the processor architecture provides PC-relative call and branch instructions, compilers can easily satisfy the first condition.

A *global offset table* provides information for address calculation. Position-independent object files (executable and shared object files) have a table in their data segment that holds addresses. When the system creates the memory image for an object file, the table entries are relocated to reflect the absolute virtual addresses assigned for an individual process. Because data segments are private for each process, the table entries can change - whereas text segments do not change because multiple processes share them.

Due to the 16-bit *offset* field of load and store instructions, the global offset table is limited to 16,384 entries (65,536 bytes).

The 16-bit offset fields of instructions require two instructions to load a 32-bit absolute value into a register. In the following code fragments wherever a 32-bit abso

lute value is loaded with a combination of lui and addiu instructions, the proper correction was made to the high 16 bits before setting the most significant (sign) bit of the low order 16 bits of the value.

Position-Independent Function Prologue

This section describes the function prologue for position-independent code. A function prologue first calculates the address of the global offset table, leaving the value in register *\$28*, hereafter referred to by its software name gp. This address is also known as the *context pointer*. This calculation is a constant offset between the text and data segments, known at the time the program is linked.

The offset between the start of a function and the global offset table (known because the global offset table is kept in the data segment) is added to the virtual address of the function to derive the virtual address of the global offset table. This value is maintained in the *gp* register throughout the function.

The virtual address of a called function is passed to the function in general register *\$25*, hereafter referred to by its software name t9. All callers of position independent functions must place the address of the called function in t9.



Although this section contains examples, an ABI compliant program must use register ± 9 for the context register. The interface to the system library routines described in Chapter 6 of the $System\ V\ ABI$ relies on the address of the called procedure being passed in ± 9 .

After calculating the *gp*, a function allocates the local stack space and saves the gp on the stack, so it can be restored after subsequent function calls. In other words, the gp is a *caller saved* register.

The code in the following figure illustrates a position-independent function prologue. _gp_disp represents the offset between the beginning of the function and the global offset table.

```
name:
la gp, _gp_disp
addu gp, gp, t9
addiu sp, sp, -64
sw gp, 32(sp)
```

Various optimizations are possible in this code example and the others that follow. For example, the calculation of *gp* need not be done for a position-independent function that is strictly local to an object module. However, the simplest, most general examples are used to keep the complexity to a minimum.

Data Objects

This section describes data objects with static storage duration. The discussion excludes stack-resident objects, because programs always compute their virtual addresses relative to the stack pointer.

In the MIPS architecture, only load and store instructions access memory. Because instructions cannot directly hold 32-bit addresses, a program normally computes an address into a register, using one instruction to load the high 16 bits of the address and another instruction to add the low 16 bits of the address.



In actual practice, most data references are performed by a single machine instruction using a gp relative address into the *global data area* (the global offset table and the global data area are both addressed by gp in position–independent code). However, those references are already position–independent and this section illustrates the differences between absolute addressing and position independent addressing.

Figure 3-30: Absolute Load and Store

C

Assembly

extern int src; extern int dst; extern int *ptr; ptr = &dst; *ptr = src;

```
.globl src, dst, ptr
lui
        t6, dst >> 16
addiu
        t6, t6, dst & 0xffff
lui
        t7, ptr >> 16
sw
        t6, ptr & 0xffff(t7)
        t6, src >> 16
lui
lw
        t6, src & 0xffff(t6)
lui
        t7, ptr >> 16
lw
        t7, ptr & 0xffff(t7)
        t6, 0(t7)
sw
```

Position-independent instructions cannot contain absolute addresses. Instead, instructions that reference symbols hold the symbols' offsets into the global offset table. Combining the offset with the global offset table address in gp gives the absolute address of the table entry holding the desired address.



The offset of data item name is represented as name_got_off in the global offset table. This is only a convention and there is no actual assembler support for these constructs.

Position-Independent Load and Store

	7
	٠

Assembly

	C
extern i extern i extern i	nt dst;
ptr = &d	
*ptr = s	rc;

globl	src, dst, ptr
lw	t7, dst_got_off(gp)
lw	<pre>t6, ptr_got_off(gp)</pre>
nop	
sw	t7, 0(t6)
lw	t7, src_got_off(gp)
nop	
lw	t7, 0(t7)
lw	<pre>t6, ptr_got_off(gp)</pre>
nop	
lw	t6, 0(t6)
nop	
sw	t7, 0(t6)

Function Calls

Programs use the jump and link instruction, jal, to make direct function calls. Since the jal instruction provides 28 bits of address and the program counter contributes the four most significant bits, direct function calls are limited to the current 256 MByte chunk of the address space as defined by the four most significant bits of pc.

Figure 3-31: Absolute Direct Function Call

C Assembly

extern void function(); jal function nop

Calls to functions outside the 256 MByte range and other indirect function calls are done by computing the address of the called function into a register and using the jump and link register, jalr, instruction.

Figure 3-32: Absolute Indirect Function Call

C Assembly

```
extern void (*ptr)();
extern void name();
ptr = name;

(*ptr)();
```



Normally, the data area for the variable ptr is kept in the global data area and is accessed relative to register gp. However, this example illustrates the difference between absolute data references and position–independent data references.

Calling position independent code functions is always done with the jalr instruction. The global offset table holds the absolute addresses of all position independent functions.

Figure 3-33: Position-Independent Function Calls

C

extern void (*ptr)();

extern void name();
name();

ptr = name;

Assembly

```
.global
         ptr, name
lw
         t9, name_got_off(gp)
nop
jalr
         t9
nop
         gp, 24(sp)
lw
nop
lw
         t7, name_got_off(gp)
lw
         t6, ptr_got_off(gp)
nop
         t7,0(t6) (*ptr)();
sw
lw
         t7, ptr_got_off(gp)
nop
lw
         t9, 0(t7)
nop
jalr
         t9
nop
lw
         gp, 24(sp)
nop
```

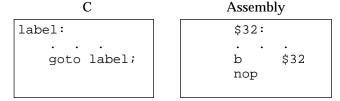


 ${\tt gp}$ must be restored on return because called position independent functions can change it. ${\tt gp}$ is saved in the stack frame in the prologue of position–independent code functions.

Branching

Programs use branch instructions to control execution flow. As defined by the architecture, branch instructions hold a PC-relative value with a 256 KByte range, allowing a jump to locations up to 128 KBytes away in either direction.

Figure 3-34: Branch Instruction, All Models



C switch statements provide multiway selection. When case labels of a switch statement satisfy grouping constraints, the compiler implements the selection with an address table. The address table is placed in a .rdata section; this so the linker can properly relocate the entries in the address table. Figures 3-36 and 3-37 use the following conventions to hide irrelevant details:

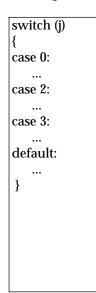
- \blacksquare The selection expression resides in register ± 7 ;
- case label constants begin at zero;
- case labels, default, and the address table use assembly names .Lcasei, .Ldef, and .Ltab, respectively.

Address table entries for absolute code contain virtual addresses; the selection code extracts the value of an entry and jumps to that address. Position-independent table entries hold offsets; the selection code compute the absolute address of a destination.

Figure 3-35: Absolute switch Code

C

Assembly



```
at, t7, 4
            sltiu
            beq
                         at, zero, .Ldef
            sll
                         t7, t7, 2
                         t6, .Ltab >> 16
            lui
            addiu
                         t6, .Ltab & 0xffff
                         t6, t6, t7
            addu
            lw
                         t7, 0(t6)
            nop
                         t7
            jr
            nop
.Ltab:
            .word
                         .Lcase0
                         .Ldef
            .word
                         .Lcase2
            .word
            .word
                         . Lcase 3 \\
```

Figure 3-36: Position-independent switch Code

C

Assembly

switch (j)
{
case 0:
case 2:
case 3:
default:
}

```
sltiu
                          at, t7, 4
                          at, zero, .Ldef
             beq
                          t7, t7, 2
             sll
                          at, .Ltab_got_off(gp)
             lw
             nop
             addu
                          at, at, t7
                          t6, 0(at)
             lw
             nop
                          t6, t6, gp
             addu
             jr
             nop
             .rdata
                          . L case 0\_gp\_off
.Ltab:
             . word \\
                          .Ldef_gp_off
             .word
                          .Lcase2_gp_off
.Lcase3_gp_off
             .word
             .word
```

C Stack Frame

Figure 3-37 shows the C stack frame organization. It conforms to the standard stack frame with designated roles for unspecified areas in standard frame.

Figure 3-37: C Stack Frame

Base	Offset	Contents	Frame
		local space: automatic variables	High addresses
		 compiler scratch space: temporaries register save area	Current
\$sp	16	outgoing arguments 5	
		outgoing argument 4	
\$sp	0	outgoing argument 1	Low addresses

A C stack frame does not normally change size during execution. The exception is dynamically allocated stack memory, discussed below. By convention, a function allocates automatic (local) variables in the top of its frame and references them as positive offsets from *sp*. Its incoming arguments reside in the previous frame, referenced as positive offsets from *sp* plus the size of the stack frame.

Variable Argument List

Previous sections describe the rules for passing arguments. Unfortunately, some otherwise portable C programs depend on other argument passing schemes, implicitly assuming that 1) all arguments reside on the stack, and 2) arguments appear in increasing order on the stack. Programs that make these assumptions never have been portable, but they have worked on many machines. They do *not* work on MIPS based systems because some arguments can reside in registers. Portable C programs should use the facilities defined in the header files <stdarg.h> or <varargs.h> to deal with variable argument lists (on MIPS and other machines as well). A program implicitly uses <stdarg.h> when it specifies a prototype declaration with an ellipsis ("...") in the argument list. No prototype or a prototype with no ellipsis causes <varargs.h> to be used.

When a function uses <stdarg.h>, the compiler modifies the argument passing

rules described above. In the calling function, the compiler passes the first 4 32-bit words of arguments in registers \$4, \$5, \$6, and \$7, regardless of data type. In particular, this means that floats and doubles are passed in the integer register. In the called function, the compiler arranges that the argument registers are saved on the stack in the locations reserved for incoming arguments. This allows the called function to reference all incoming arguments from consecutive locations on the stack.

When a function uses <varargs.h>, the situation is somewhat different. The calling function uses the argument passing rules exactly as described in the the section on argument passing rules. However, the called function allocates 32 bytes immediately adjacent to the space for incoming arguments in which to save incoming floating-point argument values.

If va_list appears as the first argument, it spills the \$f12/\$f13, and \$f14/\$f15 register pairs at -24 and -32 bytes respectively, relative to the increasing argument area. If va_alist appears as the second argument, it spills the \$f14/\$f15 register pair at -24 bytes relative to the incoming argument area.

Figure 3-38: Called Function Stack Frame

Base	Offset	Contents	Frame		
		unspecified	High addresses		
		variable size			
		(if present)			
		incoming arguments	Previous		
	+16	passed in stack frame			
		space for incoming			
old \$s	p +0	arguments 1-4			
		16 bytes reserved			
		8 bytes to spill <i>\$f12/\$f13</i>			
		8 bytes to spill <i>\$f14/\$f15</i>			
		locals and			
		temporaries			
		general register			
		save area	Current		
		floating-point			
		register save area			
		argument			
\$sp	+0	build area	Low addresses		

The 30 most-significant bits of the va_list type locate the next address in the incoming arguments to process with the va_arg macro. This address is calculated by the rules given below. The two least significant bits encode whether the va_arg macro will read floating-point values from the incoming argument area or from the floating-point save area described in the previous paragraph.

The va_start() macro in <varargs.h> encodes the following states in the two least significant bits of the va_list type:

- If the va_list pointer points to the first argument, va_start subtracts 1 from the va_list pointer, leaving it completely misaligned.
- If the va_list pointer points to the second argument, and the first argument was type double, va_start subtracts 2 from the va_list pointer, leaving it 2-byte aligned.
- For all other cases, va_start leaves the low-order bits of the va_list pointer set to zero (leaving it 4-byte aligned).

The va_start() macro in <varargs.h> requires built-in compiler support to determine which position in the argument list the va_alist parameter appears.

The va_start()macro in <stdarg.h> always sets the two least significant bits of the va_list type to zero.

If the second argument of the va_arg() macro is not the type double or the va_list pointer is 4-byte aligned, it zeroes the two least significant bits of the va_list pointer in calculating the next argument to return. It advances the value of the va_list pointer by the size of the type passed to va_arg. This leaves the va_list pointer 4-byte aligned.

If the second argument to va_arg() is type double and the va_list pointer's least significant bit is 1, it returns the value of the *Sf12/Sf13* register pair saved 32 bytes below the incoming argument. The address of the save area must be calculated by subtracting 31 from the value of the va_list pointer. The va_arg macro advances va_list pointer by 7 leaving it 2-byte aligned.

If the second argument to va_arg() is type double and the va_list pointer's value is 2-byte aligned, it returns the value of the *\$f14/\$f15* register pair saved 16 bytes below the incoming argument area. The address of the save area must be calculated by subtracting -30 from the value of the va_list pointer. The va_arg macro advances va_list pointer by 10 leaving it 4-byte aligned.

Dynamic Allocation of Stack Space

The C language does not require dynamic stack allocation *within* a stack frame. Frames are allocated dynamically on the program stack, depending on program execution. The architecture, standard calling sequence, and stack frame support dynamic allocation for programming languages that require it. Thus languages that need dynamic stack frame sizes can call C functions and vice versa.

When a function requires dynamically allocated stack space it manifests a *frame pointer* on entry to the function. The frame pointer is kept in a callee-saved register so that it is not changed across subsequent function calls. Dynamic stack allocation requires the following steps.

1. On function entry, the function adjusts the stack pointer by the size of the static stack frame. The frame pointer is then set to this initial *sp* value and is used for referencing the static elements within the stack frame, performing the normal function of the stack pointer.

- 2. Stack frames are doubleword (8 byte) aligned; dynamic allocation preserves this property. Thus, the program rounds (up) the desired byte count to a multiple of 8.
- 3. To allocate dynamic stack space, the program decreases the stack pointer by the rounded byte count, increasing its frame size. At this point, the new space resides between the register save area and the argument build area and the argument build area effectively moves down.



Standard calling sequence rules require that any frame pointer manifest within a function be initialized within the first basic block of the function. In other words, it must be set before any branches or calls.

Even in the presence of signals, dynamic allocation is "safe." If a signal interrupts allocation, one of three things can happen.

- The signal handler can return. The process resumes the dynamic allocation from the point of interruption.
- The signal handler can execute a non-local goto, or longjmp [see set-jmp(BA_LIB)]. This resets the process to a new context in a previous stack frame, automatically discarding the dynamic allocation.
- The process can terminate.

Regardless of when the signal arrives during dynamic allocation, the result is a consistent (though possibly dead) process.

Existing stack objects reside at fixed offsets from the frame pointer; stack heap allocation does not move them. No special code is needed to free dynamically allocated stack memory. The function epilogue resets the stack pointer and removes the entire stack frame, including the heap, from the stack. Naturally, a program should not reference heap objects after they have gone out of scope.

ELF Header

Machine Information

For file identification in <code>e_ident[]</code> , MIPS requires the values listed in Figure 4-1.

Figure 4–1: MIPS Identification, e_ident[]

Value
ELFCLASS32
ELFDATA2MSB

Processor identification resides in the ELF header e_machine member and must have the value 8, defined as the name ${\tt EM_MIPS}$.

The ELF header e_flags member holds bit flags associated with the file, as listed in Figure 4-2.

Figure 4–2: Processor–Specific Flags, e_flags

Name	Value
EF_MIPS_NOREORDER EF_MIPS_PIC EF_MIPS_CPIC EF_MIPS_ARCH	0x00000001 0x00000002 0x00000004 0xf0000000

EF_MIPS_NOREORDER	This bit is asserted when at least one .noreorder directive in an assembly language source contributes to the object module.
EF_MIPS_PIC	This bit is asserted when the file contains position-independent code that can be relocated in memory.
EF_MIPS_CPIC	This bit is asserted when the file contains code that follows standard calling sequence rules for calling position-independent code. The code in this file is not necessarily position independent. The EF_MIPS_PIC and EF_MIPS_CPIC flags

OBJECT FILES 4-1

must be mutually exclusive.

EF_MIPS_ARCH

The integer value formed by these four bits identify extensions to the basic MIPS I architecture. An *ABI* compliant file must have the value zero in these four bits. Non-zero values indicate the object file or executable contains program text that uses architectural extensions to the MIPS I architecture.

Sections

Figure 4-3 lists the MIPS-defined special section index which is provided in addition to the standard special section indexes.

Figure 4–3: Special Section Indexes

Name		Value				
SHN_MIPS_ACOMMON SHN_MIPS_TEXT SHN_MIPS_DATA SHN_MIPS_ SCOMMON SHN_MIPS_SUNDEFINED		0xff02 0xff03	or or or	(SHN_LOPROC (SHN_LOPROC (SHN_LOPROC (SHN_LOPROC (SHN_LOPROC	T; + + T; +	1) 2) 3)
mon s st_vali tual ac reloca addres found SHN_ ject. T SHN_		ymbols who we member ddress for ted, the aling selection in shared COMMON The dyname.	nich r of s that gnn ved ved obje V are	ative to this sec are defined and such a symbol of symbol. If the s nent indicated l , up to modulo of ects with section e not allocated inker must alloc mbols that do r	d all cont secti by tl 35,5 n inc n th	ocated. The tains the virion must be he virtual 36. Symbols dex e shared obspace for
SHN_MIPS_TEXT						
only p the pi progra relativ		ols defined relative to these two sections are present after a program has been rewritten by ixie code profiling program. Such rewritten ams are not ABI-compliant. Symbols defined we to these two sections will never occur in an compliant program.				
mon s data a Area"		symbols w rea (are <i>gp</i>	hich -ado pter	ative to this sec can be placed i dressable). See . This section o	n th "Gl	ne global obal Data

OBJECT FILES 4-3

SHN MIPS SUNDEFINED

Undefined symbols with this special section index in the *st_shndx* field can be placed in the global data area (gp-addressable). See "Global Data Area" in this chapter. This section only occurs in relocatable object files.

Figure 4-4 lists the MIPS-defined section types in addition to the standard section types.

Figure 4–4: Section Types, sh_type

Name	Value
SHT_MIPS_LIBLIST	0x70000000 or (SHT_LOPROC + 0)
SHT_MIPS_CONFLICT	0x70000002 or (SHT_LOPROC + 2)
SHT_MIPS_GPTAB	0x70000003 or (SHT_LOPROC + 3)
SHT_MIPS_UCODE	0x70000004 or (SHT_LOPROC + 4)
SHT_MIPS_DEBUG	0x70000005 or (SHT_LOPROC + 5)
SHT_MIPS_REGINFO	0x70000006 or (SHT_LOPROC + 6)

SHT MIPS LIBLIST

The section contains information about the set of dynamic shared object libraries used when statically linking a program. Each entry contains information such as the library name, timestamp, and version. See "Quickstart" in Chapter 5 for details.

SHT_MIPS_CONFLICT The section contains a list of symbols in an executable whose definitions conflict with shared-object defined symbols. See "Quickstart" in Chapter 5 for details.

SHT_MIPS_GPTAB

The section contains the *global pointer table*. The global pointer table includes a list of possible global data area sizes. The list allows the linker to provide the user with information on the optimal size criteria to use for gp register relative addressing. See "Global Data Area" below for details.

4-4

SHT_MIPS_UCODE	This section type is reserved and the contents are unspecified. The section contents can be ignored.
SHT_MIPS_DEBUG	The section contains debug information specific to MIPS. An ABI-compliant application does not need to have a section of this type.
SHT_MIPS_REGINFO	The section contains information regarding register usage information for the object file. See Register Information for details.

A section header sh_flags member holds 1-bit flags that describe the attributes of the section. In addition to the values defined in the *System VABI*, Figure 4-5 lists the MIPS-defined flag.

Figure 4-5: Section Attribute Flags, sh_flags

Name	Value	
SHF_MIPS_GPREL	0x10000000	

SHF_MIPS_GPREL

The section contains data that must be part of the global data area during program execution. Data in this area is addressable with a gp relative address. Any section with the SHF_MIPS_GPREL attribute must have a section header index of one of the .gptab special sections in the sh_link member of its section header table entry. See "Global Data Area" below for details.

The static linker does not guarantee that a section with the SHF_MIPS_GPREL attribute will remain in the global data area after static linking.

Figure 4-6 lists the MIPS-defined section header sh_link and sh_info members interpretation for the MIPS-specific section types.

OBJECT FILES 4-5

Figure 4–6: sh_link and sh_info interpretation

sh_type	sh_link	sh_info	
SHT_MIPS_LIBLIST	The section header index of the string table used by en- tries in this section.	The number of entries in this section.	
SHT_MIPS_GPTAB	not used	The section header index of the SHF_ALLOC + SHF_WRITE section. See " Global Data Area" in this chapter.	

Special Sections

MIPS defines several additional special sections. Figure 4-7 lists their types and corresponding attributes.

Figure 4-7: Special Sections

Name	Туре	Attributes
.text	SHT_PROGBITS	SHF_ALLOC + SHF_EXECINSTR
.sdata	SHT_PROGBITS	SHF_ALLOC + SHF_WRITE + \
		SHF_MIPS_GPREL
.sbss	SHT_NOBITS	SHF_ALLOC + SHF_WRITE + \
		SHF_MIPS_GPREL
.lit4	SHT_PROGBITS	SHF_ALLOC + SHF_WRITE + \
		SHF_MIPS_GPREL
.lit8	SHT_PROGBITS	SHF_ALLOC + SHF_WRITE + \
		SHF_MIPS_GPREL
.reginfo	SHT_MIPS_REGINFO	SHF_ALLOC
.liblist	SHT_MIPS_LIBLIST	SHF_ALLOC
.conflict	SHT_CONFLICT	SHF_ALLOC
.gptab	SHT_MIPS_GPTAB	none
.got	SHT_PROGBITS	SHF_ALLOC + SHF_WRITE + \
		SHF_MIPS_GPREL
.ucode	SHT_MIPS_UCODE	none
.mdebug	SHT_MIPS_DEBUG	none
.dynamic	SHT_DYNAMIC	SHF_ALLOC
.rel.dyn	SHT_REL	SHF_ALLOC



A $MIPS\ ABI$ compliant system must support the .sdata, .sbss, .lit4, .lit8, .reginfo, and .gptab sections. A $MIPS\ ABI$ compliant system must recognize, but may choose to ignore the contents of the .liblist or .conflict sections. However, if either of these optional sections is supported, both must be supported.

.text

This section contains only executable instructions. The first two instructions immediately preceding the first function in the section must be a jump to return address instruction followed by a nop. The stack traceback algorithm, described in Chapter 3, depends on this.

.sdata

This section holds initialized short data that contribute to the program memory image. See "Global Data Area" below for details.

OBJECT FILES 4-7

. sbss This section holds uninitialized short data that contribute to the program memory image. By definition, the system initializes the data with zeros when the program begins to run. See "Global Data Area" below for details.

.lit4

.reginfo

.conflict

.gptab

.ucode

This section holds 4 byte read-only literals that contribute to the program memory image. Its purpose is to provide a list of unique 4-byte literals used by a program. See "Global Data Area" below for details. Although this section has the SHF_WRITE attribute, it is not expected to be written. Placing this section in the data segment mandates the SHF_WRITE attribute.

This section holds 8 byte read-only literals that contribute to the program memory image. Its purpose is to provide a list of unique 8-byte literals used by a program. See "Global Data Area" below for details. Although this section has the SHF_WRITE attribute, it is not expected to be written. Placing this section in the data segment mandates the SHF_WRITE attribute.

This section provides information on the program register usage to the system. See "Register Information" below for details.

.liblist This section contains information on each of the libraries used at static link time as described in "Quickstart" in Chapter 5.

This section provides additional dynamic linking information about symbols in an executable file that conflict with symbols defined in the dynamic shared libraries with which the file is linked. See "Quickstart" in Chapter 5 for details.

This section contains a global pointer table. The global pointer table is described in "Global Data Area" in this chapter. The section is named .gptab.sbss,.gptab.sdata, gptab.bss, or .gptab.data depending on which data section the particular .gptab refers.

This section name is reserved and the contents of this type of section are unspecified. The section contents can be ignored .mdebug

This section contains symbol table information as emitted by the MIPS compilers. Its content is described in Chapter 10 of the MIPS Assembly Language Programmer's Guide, order number ASM-01-DOC, (Copyright © 1989, MIPS Computer Systems, Inc.). The information in this section is dependent on the location of other sections in the file; if an object is relocated, the section must be updated. Discard this section if an object file is relocated and the ABI compliant system does not update the section.

.got

This section holds the global offset table. See "Coding Examples" in Chapter 3 and " Global Offset Table" in Chapter 5 for more information.

.dynamic

This is the same as the generic ABI section of the same type, but the MIPS-specific version does not include the SHF_WRITE attribute.

.rel.dyn

This relocation section contains run-time entries for the <code>.data</code> and <code>.sdata</code> sections. See "Relocations" in Chapter 5 for more information.



Sections that contribute to a loadable program segment must not contain overlapping virtual addresses.

OBJECT FILES 4-9

Symbol Table

Symbol Values

If an executable or shared object contains a reference to a function defined in one of its associated shared objects, the symbol table section for that file will contain an entry for that symbol. The st_shndx member of that symbol table entry contains SHN_UNDEF. This signals to the dynamic linker that the symbol definition for that function is not contained in the executable file. If there is a stub for that symbol in the executable file and the st_value member for the symbol table entry is non-zero, the value will contain the virtual address of the first instruction of that procedure's stub. Otherwise, the st_value member contains zero. This stub calls the dynamic linker at runtime for lazy text evaluation. See "Function Addresses" in Chapter 5 for details.

Global Data Area

The global data area is part of the data segment of an executable program. It con-

OBJECT FILES 4-11

Figure 4–8: Global Pointer Table

```
typedef union {
    struct {
        Elf32_Word gt_current_g_value;
        Elf32_Word gt_unused;
    } gt_header;
    struct {
        Elf32_Word gt_g_value;
        Elf32_Word gt_bytes;
    } gt_entry;
} Elf32_gptab;
```

```
gt_header.gt_current_g_value
```

This member is the size criterion actually used for this object file. Data items of this size or smaller are referenced with gp relative addressing and reside in a SHF_MIPS_GPREL section.

gt_header.gt_unused

This member is not used in the first entry of the Elf32_gptab array.

gt_entry.gt_g_value

This member is a hypothetical size criterion value.

gt_entry.gt_bytes

This member indicates the length of the global data area if the corresponding gt_entry.gt_g_value were used.

The first element of the ELF_32_gptab array is alway of type gt_header; this entry must always exist. Additional elements of the array are of type gt_entry.

Each of the $gt_entry.gt_g_value$ fields is the size of an actual data item encountered during compilation or assembly, including zero. Each separate size criteria results in a overall size for the global data area. The various entries are

sorted and duplicates are removed. The resulting set of entries, including the actual size criterion used, yields the <code>.gptab</code> section.

There are always at least two .gptab

OBJECT FILES 4-13

Register Information

The compilers and assembler collect information on the registers used by the code in the object file. This information is communicated to the operating system kernel using a <code>.reginfo</code> section. The operating system kernel can use this information to decide what registers it does not need to save or which coprocessors the program uses. The section also contains a field which specifies the initial value for the <code>gp register</code>, based on the final location of the global data area in memory.

Figure 4-9: Register Information Structure

```
typedef struct {
   Elf32_Word ri_gprmask;;
   Elf32_Word ri_cprmask[4];
   Elf32_SWord ri_gp_value;
} ELF_RegInfo;
```

ri_gprmask

This member contains a bit-mask of general registers used by the program. Each set bit indicates a general integer register used by the program. Each clear bit indicates a general integer register not used by the program. For instance, bit 31 set indicates register \$31 is used by the program; bit 27 clear indicates register \$27 is not used by the program.

ri_cprmask

This member contains the bit-mask of co-processor registers used by the program. The MIPS RISC architecture supports up to four co-processors, each with 32 registers. Each array element corresponds to one set of co-processor registers. Each of the bits within the element corresponds to individual register in the co-processor register set. The 32 bits of the words correspond to the 32 registers, with bit number 31 corresponding to register 31, bit number 30 to register 30, etc. Set bits indicate the corresponding register is used by the program; clear bits indicate the program does not use the corresponding register.

ri_gp_value

This member contains the gp register value. In relocatable object files it is used for relocation of the R_MIPS_GPREL and $R_MIPS_LITERAL$ relocation types.



Only co-processor 1 can be used by ABI-compliant programs. This means that only the ri_cprmask[1] array element can have a non-zero value.

ri_cpr-mask[0], ri_cprmask[2], and ri_cprmask[3] must all be zero in an ABI-compliant program.

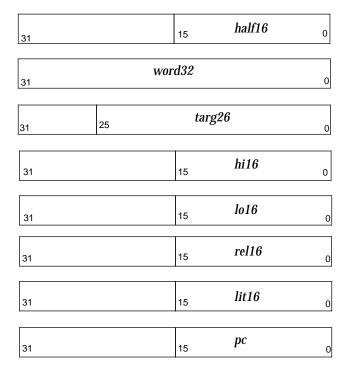
OBJECT FILES 4-15

Relocation

Relocation Types

Relocation entries describe how to alter the following instruction and data fields shown in Figure 4-10; bit numbers appear in the lower box corners.

Figure 4–10: Relocatable Fields



Calculations below assume the actions are transforming a relocatable file into either an executable or a shared object file. Conceptually, the linker merges one or more relocatable files to form the output. It first determines how to combine and locate the input files; then it updates the symbol values, and finally it performs the relocation.

Relocations applied to executable or shared object files are similar and accomplish the same result. Descriptions below use the following notation.

- A Represents the addend used to compute the value of the relocatable field.
- AHL Identifies another type of addend used to compute the value of the relocatable field. See the note below for more detail.
- P Represents the place (section offset or address) of the storage unit being relocated (computed using r_offset).
- S Represents the value of the symbol whose index resides in the relocation entry, unless the the symbol is STB_LOCAL and is of type STT_SECTION in which case S represents the original sh_addr minus the final sh_addr.
- Represents the offset into the global offset table at which the address of the relocation entry symbol resides during execution. See "Coding Examples" in Chapter 3 and "Global Offset Table" in Chapter 5 for more information.
- Represents the final gp value to be used for the relocatable, executable, or shared object file being produced.
- GPO Represents the gp value used to create the relocatable object.
- EA Represents the effective address of the symbol prior to relocation.
- Represents the .lit4 or .lit8 literal table offset. Prior to relocation the addend field of a literal reference contains the offset into the global data area. During relocation, each literal section from each contributing file is merged and sorted, after which duplicate entries are removed and the section compressed, leaving only unique entries. The relocation factor L is the mapping from the old offset of the original gp to the value of gp used in the final file.

A relocation entry r_{offset} value designates the offset or virtual address of the first byte of the affected storage unit. The relocation type specifies which bits to change and how to calculate their values. Because MIPS uses only Elf32_Rel relocation entries, the relocated field holds the addend.

The AHL addend is a composite computed from the addends of two consecutive relocation entries. Each relocation type of R_{MIPS_HI16} must have an associated R_{MIPS_LO16} entry immediately following it in the list of relocations.

OBJECT FILES 4-17

These relocation entries are always processed as a pair and both addend fields contribute to the AHL addend. If AHI and ALO are the addends from the paired R_MIPS_HI16 and R_MIPS_LO16 entries, then the addend AHL is computed as (AHI << 16) + (short)ALO. R_MIPS_LO16 entries without an R_MIPS_HI16 entry immediately preceding are orphaned and the previously defined R_MIPS_HI16 is used for computing the addend.



The field names in Table 4–11 tell whether the relocation type checks for overflow. A calculated relocation value can be larger than the intended field, and a relocation type can verify (V) the value fits or truncate (T) the result. As an example, V–half16 means the computed value cannot have significant non–zero bits outside the half16 field.

Figure 4-11: Relocation Types

Name	Valu	e Field	Symbol	Calculation
R_MIPS_NONE	0	none	local	none
R_MIPS_16	1	V-half16	external	S + sign-extend(A)
	1	V-half16	local	S + sign-extend(A)
R_MIPS_32	2	T-word32	external	S + A
	2	T-word32	local	S + A
R_MIPS_REL32	3	T-word32	external	A - EA + S
R_MIPS_REL32	3	T-word32	local	A - EA + S
R_MIPS_26	4	T-targ26	local	(((A << 2) \
				(P & 0xf0000000) + S) >> 2
	4	T-targ26	external	(sign-extend(A < 2) + S) >> 2
R_MIPS_HI16	5	T-hi16	external	((AHL + S) - \
				(short)(AHL + S)) >> 16
	5	T-hi16	local	((AHL + S) - \
				(short)(AHL + S)) >> 16
	5	V-hi16	_gp_disp	(AHL + GP - P) - (short) \
				(AHL + GP - P)) >> 16
R_MIPS_L016	6	T-1o16	external	AHL + S
	6	T-1o16	local	AHL + S
	6	V-lo16	_gp_disp	AHL + GP - P + 4
R_MIPS_GPREL16	7	V-rel16	external	sign-extend(A) + S + GP
	7	V-rel16	local	sign-extend(A) + S + GP0 - GP
R_MIPS_LITERAL	8	V-lit16	local	sign-extend(A) + L
R_MIPS_GOT16	9	V-rel16	external	G
	9	V-rel16	local	see below
R_MIPS_PC16	10	V-pc16	external	sign-extend(A) + S - P
R_MIPS_CALL16	11	V-rel16	external	G
R_MIPS_GPREL32	12	T-word32	local	A + S + GPO - GP
R_MIPS_GOTHI16	21	T-hi16	external	(G - (short)G) >> 16 + A
R_MIPS_GOTLO16	22	T-1o16	external	
R_MIPS_CALLHI16	1	T-hi16	external	(- (
R_MIPS_CALLLO16	31	T-lo16	external	G & 0xffff

In the Symbol column in the table above, local refers to a symbol referenced by the symbol table index in the relocation entry $STB_LOCAL/STT_SECTION$. Otherwise, the relocation is considered an *external* relocation. See below for $_gp_disp$ relocations.

The R_MIPS_Rel32 relocation type is the only relocation performed by the dynamic linker. The value EA $\,$ used by the dynamic linker to relocate an

OBJECT FILES 4-19

R_MIPS_REL32 relocation depends on its r_symndx value. If the relocation entry r_symndx is less than DT_MIPS_GOTSYM, the value of EA is the symbol st_value plus displacement. Otherwise, the value of EA is the value in the GOT entry corresponding to the relocation entry r_symndx. The correspondence between the GOT and the dynamic symbol table is described in the "Global Offset Table" section in Chapter 5.

If an R_MIPS_GOT16 refers to a locally defined symbol, then the relocation is done differently than if it refers to an external symbol. In the local case, the R_MIPS_GOT16 must be followed immediately with a R_MIPS_LO16 relocation. The AHL addend is extracted and the section in which the referenced data item resides is determined (requiring that all sections in an object module have unique addresses and not overlap). From this address the final address of the data item is calculated. If necessary, a global offset table entry is created to hold the high 16 bits of this address (an existing entry is used when possible). The *rel16* field is replaced by the offset of this entry in the global offset table. The *lo16* field in the following R_MIPS_LO16 relocation is replaced by the low 16 bits of the actual destination address. This is meant for local data references in position-independent code so that only one global offset table entry is necessary for every 64 KBytes of local data.

The first instance of R_MIPS_GOT16 , R_MIPS_CALL16 , $R_MIPS_GOT_H116$, $R_MIPS_CALL_H116$, $R_MIPS_GOT_LO16$, or $R_MIPS_CALL_LO16$. Relocations cause the link editor to build a global offset table if one has not already been built.

The symbol name <code>_gp_disp</code> is reserved. Only <code>R_MIPS_HI16</code> and <code>R_MIPS_LO16</code> relocations are permitted with <code>_gp_disp</code>. These relocation entries must appear consecutively in the relocation section and they must reference consecutive relocation area addresses.

R_MIPS_CALL16, R_MIPS_CALL_HI16, and R_MIPS_CALL_LO16 relocation entries load function addresses from the global offset table and indicate that the dynamic linker can perform lazy binding. See "Global Offset Table" in Chapter 5.

Program Loading

As the system creates or augments a process image, it logically copies a file segment to a virtual memory segment. When and if the system physically reads the file depends on the program's execution behavior, system load, etc. A process does not require a physical page unless it references a logical page during execution. Processes commonly leave many pages unreferenced; therefore delaying physical reads frequently obviates them, improving system performance. To obtain this efficiency in practice, executable and shared object files must have segment images whose virtual addresses are zero, modulo the file system block size.

Virtual addresses and file offsets for MIPS segments are congruent modulo 64 KByte (0x10000) or larger powers of 2. Because 64 KBytes is the maximum page size, the files are suitable for paging regardless of physical page size.

Figure 5-1: Example Executable File

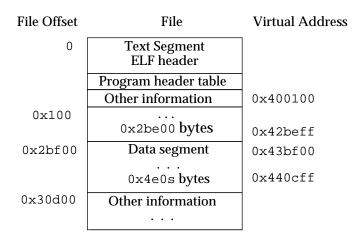


Figure 5-2: Program Header Segments

Member	Text	Data
p_type	PT_LOAD	PT_LOAD
p_offset	0	0x2bf00
p_vaddr	400100	0x43bf00
p_paddr	unspecified	unspecified
p_filesz	0x2bf00	0x4e00
p_memsz	0x2bf00	0x5e24
p_flags	PF_R+PF_X	PF_R+PF_W+PF_X
p_align	0x10000	0x10000

Because the page size can be larger than the alignment restriction of a segment file offset, up to four file pages can hold impure text or data (depending on page size and file system block size).

- The first text page contains the ELF header, the program header table, and other information.
- The last text page can hold a copy of the beginning of data.
- The first data page can have a copy of the end of text.
- The last data page can contain file information not relevant to the running process.

Logically, the system enforces the memory permissions as if each segment were complete and separate; segment addresses are adjusted to ensure each logical page in the address space has a single set of permissions. In the example in Figure 5-1, the file region holding the end of text and the beginning of data is mapped twice: once at one virtual address for text and once at a different virtual address for data.

The end of the data segment requires special handling for uninitialized data which the system defines to begin with zero values. Thus if the last data page of a file includes information not in the logical memory page, the extraneous data must be set to zero, rather than the unknown contents of the executable file. "Impurities" in the other three pages are not logically part of the process image; whether the system expunges them is unspecified.

There is one aspect of segment loading that differs between executable files and shared objects. Executable file segments typically contain absolute code [see "Coding Examples" in Chapter 3]. To let the process execute correctly, the segments must reside at the virtual addresses used to build the executable file, with

the system using the p_vaddr values unchanged as virtual addresses.

Shared object segments typically contain position-independent code, allowing a segment virtual address to change from one process to another without invalidating execution behavior. Though the system chooses virtual addresses for individual processes, it maintains the *relative positions* of the segments. Because position-independent code uses relative addressing between segments, the difference between virtual addresses in memory must match the difference between virtual addresses in the file. The following table shows possible shared object virtual address assignments for several processes, illustrating constant relative positioning. The table also illustrates the base address computations.

Figure 5-3: Example Shared Object Segment Addresses

Source	Text	Data	Base Address
File	0x200	0x2a400	0x0
Process 1	0x50000200	0x5002a400	0x50000000
Process 2	0x50010200	0x5003a400	0x50010000
Process 3	0x60020200	0x6004a400	0x60020000
Process 4	0x60030200	0x6005a400	0x60030000



In addition to maintaining the relative positionsof the segments, the system must also ensure that relocations occur in 64 KByte increments; position–independent code relies on this property.



By convention, no more than one segment will occupy addresses in the same chunk of memory, modulo 256 KBytes.

Program Header

There is one program header type specific to this supplement.

Figure 5-4: MIPS Specific Segment Types, p_type

Name	Value
PT_MIPS_REGINFO	0x70000000

PT_MIPS_REGINFO

Specifies register usage information for the executable or shared object; it cannot occur more than once in a file. Its presence is mandatory and it must precede any loadable segment entry. It identifies one <code>.reginfo</code> type section. See Register Information" in Chapter 4 for more information

Segment Contents

Figures 5-5 and 5-6 below illustrate typical segment contents for a MIPS executable or shared object. The actual order and membership of sections within a segment may alter the examples below.

Figure 5-5: Text Segment

.reginfo
.dynami
.liblist
.rel.dyn
.conflict
.dynstr
.dynsym
.hash
.rodata
.text

Figure 5-6: Data Segment

.got
.lit4
.lit8
.sdata
.data
.sbss
.bss

Dynamic Linking

Dynamic Section

Dynamic section entries give information to the dynamic linker. Some of this information is processor-specific, including the interpretation of some entries in the dynamic structure.

Figure 5-7: Dynamic Array Tags d_tag

DT MIPS RLD VERSION	0x70000001	d val	mandatory	mandatory
		_		U
DT_MIPS_TIME_STAMP	0x70000002	d_val	optional	optional
DT_MIPS_ICHECKSUM	0x70000003	d_val	optional	optional
DT_MIPS_IVERSION	0x70000004	d_val	optional	optional
DT_MIPS_FLAGS	0x70000005	d_val	mandatory	mandatory
DT_MIPS_BASE_ADDRESS	0x70000006	d_ptr	mandatory	mandatory
DT_MIPS_CONFLICT	0x70000008	d_ptr	optional	optional
DT_MIPS_LIBLIST	0x70000009	d_ptr	optional	optional
DT_MIPS_LOCAL_GOTNO	0x7000000A	d_val	mandatory	mandatory
DT_MIPS_CONFLICTNO	0x7000000B	d_val	optional	optional
DT_MIPS_LIBLISTNO	0x70000010	d_val	optional	optional
DT_MIPS_SYMTABNO	0x70000011	d_val	mandatory	mandatory
DT_MIPS_UNREFEXTNO	0x70000012	d_val	optional	optional
DT_MIPS_GOTSYM	0x70000013	d_val	mandatory	mandatory
DT_MIPS_HIPAGENO	0x70000014	d_val	optional	optional
DT_MIPS_RLD_MAP	0x70000016	d_ptr	mandatory	ignored
DT_PLTGOT	3	d_ptr	mandatory	mandatory
DT_RPATH	15	d_val	optional	optional

DT_MIPS_RLD_VERSION

This element holds a 32-bit version id for the *Runtime Linker Interface*. This will start at integer value 1.

DT_MIPS_TIME_STAMP

This element holds a 32-bit time stamp.

DT_MIPS_ICHECKSUM

This element holds the sum of all external strings and common sizes.

DT_MIPS_IVERSION

This element holds an index into the object file string table. The version string is a series of version strings separated by colons (:). An index value of zero means no version string was specified.

DT_MIPS_FLAGS

This element holds a set of 1-bit flags. Flag definitions appear below.

DT_MIPS_BASE_ADDRESS

This member holds the *base address* of the segment. That is, it holds the virtual address of the segment as if the the segment were actually loaded at the addressed specified at static link time. It can be adjusted when the operating system kernel actually maps segments. It is used to adjust pointers based on the difference between the static link time value and the actual address.

DT_MIPS_CONFLICT

This member holds the address of the $\mbox{.}\mbox{conflict}$ section.

DT_MIPS_LIBLIST

This member holds address of the .liblist section.

DT_MIPS_LOCAL_GOTNO

This member holds the number of local global offset table entries.

DT_MIPS_CONFLICTNO

This member holds the number of entries in the .conflict section. This field is mandatory if there is a .conflict section.

DT_PLTGOT

This member holds the address of the .got section.

DT_MIPS_SYMTABNO

This member holds the number of entries in the .dynsym section.

DT_MIPS_LIBLISTNO

This member holds the number of entries in the .liblist section.

DT_MIPS_UNREFEXTNO

This member holds the index into the dynamic symbol table which is the entry of the first external symbol that is not referenced within the same object.

DT_MIPS_GOTSYM	This member holds the index of the first dynamic symbol table entry that corresponds to an entry in the global offset table. See "Global Offset Table" in this chapter.
DT_MIPS_HIPAGENO	This member holds the number of page table entries in the global offset table. A page table entry here refers to a 64 Kb chunk of data space. This member is used by profiling tools and is optional.
DT_RPATH	This member optionally appears in a shared object. If it is present in a shared object at static link time, it is propagated to the final executable's $\mathtt{DT}_\mathtt{RPATH}$.
DT_DEBUG	This member is specifically disallowed.
DT_MIPS_RLD_MAP	This member is used by debugging. It contains the address of a 32-bit word in the .data section which is supplied by the compilation environment. The word's contents are not specified and programs using this value are not <i>ABI</i> - compliant.

Figure 5-8: Dynamic section, DT_MIPS_FLAGS

Name	Value	Meaning
RHF_NONE	0x0000000	none
RHF_QUICKSTART	0x00000001	use shortcut pointers
RHF_NOTPOT	0x00000002	hash size not power of two
RHF_NO_LIBRARY_REPLACEMENT	0x00000004	ignore LD_LIBRARY_PATH

The RHF_NO_LIBRARY_REPLACEMENT flag directs the dynamic linker to ignore the LD_LIBRARY_PATH environment variable when searching for shared objects.

Shared Object Dependencies

The $System\ V\ ABI$ defines the default library search path to be /usr/lib; MIPS defines the default library search path to be /lib:/usr/lib:/usr/lib/cmplrs/cc.

Global Offset Table

In general, position-independent code cannot contain absolute virtual addresses.

Global offset tables (or GOTs) hold absolute addresses in private data, making the addresses available without compromising position-independence and sharability of a program text. A program references its global offset table using position-independent addressing and extracts absolute values, thus redirecting position-independent references to absolute locations.

The global offset table is split into two logically separate subtables: locals and externals. Local entries reside in the first part of the global offset table. The value of the dynamic tag DT_MIPS_LOCAL_GOTNO holds the number of local global offset table entries. These entries only require relocation if they occur in a shared object and the shared object memory load address differs from the virtual address of the loadable segments of the shared object. As with defined external entries in the global offset table, these local entries contain actual addresses.

External entries reside in the second part of the global offset table. Each entry in the external section corresponds to an entry in the global offset table mapped part of the .dynsym section (see "Symbols" below for a definition). The first symbol in the .dynsym section corresponds to the first word of the global offset table; the second symbol corresponds to the second word, and so on. Each word in the external entry part of the global offset table contains the *actual address* for its corresponding symbol. The external entries for defined symbols must contain actual addresses. If an entry corresponds to an undefined symbol and the global offset table entry contains a zero, the entry must be resolved by the dynamic linker, even if the dynamic linker is performing a *quickstart*. See "Quickstart" below for more information.

The following table details the various possibilities for the initial state of the global offset table mapped dynamic symbol table section and the global part of the global offset table.

Figure 5-9: Initial State, global GOT and .dynsym

Section	Type	st_value	GOT Entry	Comments
SHN_UNDEF	STT_FUNC	0	0/QS	1
SHN_UNDEF	STT_FUNC	stub addr	stub address/ QS	2
SHN_UNDEF SHN_COMMON	any	0/alignment	0/QS	
all others	STT_FUNC	address	stub address/ address	2
all others	any	address	address	3

QS stands for the *Quickstart* value of the symbol.

Comments:

- 1: had relocations related to taking the function's address
- 2: only had call related relocations defined STT_FUNC
- 3: non-STT_FUNC defined globals

After the system creates memory segments for a loadable object file, the dynamic linker can process the relocation entries. The only relocation entries remaining are type R_MIPS_REL32 referring to data containing addresses. The dynamic linker determines the associated symbol (or section) values, calculates their absolute addresses, and sets the proper values. Although the absolute addresses may be unknown when the link editor builds an object file, the dynamic linker knows the addresses of all memory segments and can find the correct symbols, thus calculating the absolute addresses contained therein.

The dynamic linker relocates the global offset table by first adding the difference between the base where the shared object is loaded and the value of the dynamic tag DT_MIPS_BASE_ADDRESS to all local global offset table entries. Next, the global GOT entries are relocated. For each global GOT entry the following relocation is performed:

Figure 5-10: Global Offset Table Relocation Algorithm

Section	Type	st_value	GOT Entry	Relocation
SHN_UNDEF	STT_FUNC	0	0/QS	1
SHN_UNDEF	STT_FUNC	stub addr	stub addr	2
SHN_UNDEF	STT_FUNC	stub addr	!= stub addr	3
SHN_UNDEF SHN_COMMON	all others	any	0/QS	1
all others	STT_FUNC	address	stub address != address*	2
all others	all others	address	address	1

^{*} Stub address must be in this executable and can only be applied the first time the GOT is modified.

Relocation:

- 1: resolve immediately or use Quickstart value
- 2: add run-time displacement to GOT entry
- 3: set GOT entry to stub address plus run-time displacement

Certain optimizations are possible with information from Quickstart. An ABI-compliant system performing such optimizations guarantees that the values of the GOT entries are the same as if the dynamic linker performed the relocation algorithm described in Figure 5-10.

If a program requires direct access to the absolute address of a symbol, it uses the appropriate global offset table entry. Because the executable file and shared objects have separate global offset tables, the address of a symbol can appear in several tables. The dynamic linker processes all necessary relocations before giving control to any code in the process image, thus ensuring the absolute addresses are available during execution.

The zero entry in the global offset table is reserved to hold the address of the entry point in the dynamic linker to call when lazy resolving text symbols. The dynamic linker must always initialize this entry regardless of whether lazy binding is or is not enabled.

The system can choose different memory segment addresses for the same shared object in different programs; it can even choose different library addresses for dif-

ferent executions of the same program. Nonetheless, memory segments do not change addresses once the process image is established. As long as a process exists, its memory segments reside at fixed virtual addresses.

Calling Position–Independent Functions

The global offset table is used to hold addresses of position-independent functions as well as data addresses. It is not possible to resolve function calls from one executable file or shared object to another at static link time, so all of the function address entries in the global offset table are normally resolved at execution time. The dynamic linker then resolves all of these undefined relocation entries at run-time. Through the use of specially constructed pieces of code known as stubs, this runtime resolution can be be deferred through a technique known as "binding, lazy binding".

Using this technique, the link editor (or a combination of the compiler, assembler, and link editor) builds a stub for each called function, and allocates a global offset table entry that initially points to the stub. Because of the normal calling sequence for position-independent code, the call ends up invoking the stub the first time the call is made.

Figure 5-11: Sample Stub Code

```
stub_xyz: .
  lw    t9, 0(gp)
  move    t7, ra
  jal    t9
  li    t8, .dynsym_index  # branch delay slot
```

In the example in Figure 5-11, the stub code loads register t9 with an entry from the global offset table which contains a well-known entry point in the dynamic linker; it also loads register t8 with the index into the .dynsym section of the referenced external. The code must save register ra in register t7 and transfer control to the dynamic linker.

The dynamic linker determines the correct address for the actual called function and replaces the address of the stub in the global offset table with the address of the function.

Most undefined text references can be handled by lazy text evaluation except when the address of a function is relocated using relocations of type

```
R_MIPS_CALL16 or R_MIPS_26.
```

The LD_BIND_NOW environment variable can also change dynamic linking behavior. If its value is non-null, the dynamic linker evaluates all symbol table entries of type STT_FUNC, replacing their stub addresses in the global offset table with the actual address of the referenced function.



Lazy binding generally improves overall application performance because NOTE unused symbols do not incur the dynamic linking overhead. Nevertheless, two situations make lazy binding undesirable for some applications. First, the initial reference to a shared object function takes longer than subsequent calls, because the dynamic linker intercepts the call to resolve the symbol. Some applications cannot tolerate this unpredictability. Second, if an error occurs and the dynamic linker cannot resolve the symbol, the dynamic linker terminates the program. Under lazy binding, this might occur at arbitrary times. Once again, some applications cannot tolerate this unpredictability. By turning off lazy binding, the dynamic linker forces the failure to occur during process initialization, before the application receives control.

Symbols

All externally visible symbols, both defined and undefined, must be hashed into the hash table.

Undefined symbols of type STT_FUNC, which have been referenced only by R MIPS CALL16 and R MIPS 26 relocations, can contain non-zero values in the their st_value field, denoting the stub address used for lazy evaluation for this symbol. The run-time linker uses this to reset the global offset table entry for this external to its stub address when unlinking a shared object. All other undefined symbols must contain zero in their st value fields.

The dynamic symbol table, like all ELF symbol tables, is divided into local and global parts. The global part of the dynamic symbol table is further divided into two parts: symbols that do not have GOT entries associated with them and symbols that do have GOT entries associated with them. The part of the dynamic symbol table with GOT entries is called the "global offset table mapped" part or "GOT mapped" part. Symbols with GOT entries have a one-to-one mapping with the global part of the GOT.

The value of the dynamic tag DT MIPS GOTSYM is the index of the first symbol with a global offset table entry in the dynamic symbol table.

Relocations

There may be only one dynamic relocation section to resolve addresses in data. It must be called .rel.dyn. Executables can contain normal relocation sections in

addition to a dynamic relocation section. The normal relocation sections may contain resolutions for any absolute values in the main program. The dynamic linker does not resolve these or relocate the main program.

As noted previously, only R_MIPS_REL32 relocation entries are supported in the dynamic relocation section.

Because sufficient information is available in the .dynamic section, the GOT has no relocation information. The relocation algorithm for the GOT is described above.

The entries in the dynamic relocation section must be ordered by increasing r_{symndx} value.

Ordering

To take advantage of *Quickstart* functionality, the .dynsym and .rel.dyn sections must obey ordering constraints. The GOT-mapped portion of the .dynsym section must be ordered on increasing values in the st_value field. This requires that the .got section have the same order, since it must correspond to the .dynsym section.

The .rel.dyn section must have all local entries first, followed by the external entries. Within these sub-sections, the entries must be ordered by symbol index.

Quickstart

The MIPS supplement to the *ABI* defines two sections which are useful for faster start-up of programs when the programs have been linked with dynamic shared objects. The group of structures defined in these sections allow the dynamic linker to operate more efficiently than when these sections are not present. These additional sections are also used for more complete dynamic shared object version control.



An *ABI* compliant system can ignore the sections defined here, but if it supports one of these sections, it must support both of them. If you relink or relocate the object file on secondary storage and cannot process these sections, you must delete them.

Shared Object List

A shared object list section is an array of structures that contain information about the various dynamic shared objects used to statically link this object file. Each separate shared object used generates one Elf32_Lib array element. The shared object list is used for more complete shared object version control.

Figure 5-12: Shared Object Information Structure

```
typedef struct {
   Elf32_Word l_name;
   Elf32_Word l_time_stamp;
   Elf32_Word l_checksum;
   Elf32_Word l_version;
   Elf32_Word l_flags;
} Elf32_Lib;
```

1_name

This member specifies the name of a shared object. Its value is a string table index. This name can be a trailing component of the path to be used with RPATH + LD_LIBPATH or a name containing '/'s, which is relative to '.', or it can be a full pathname.

1_time_stamp This member's value is a 32 bit time stamp. It can be com-

bined with the <code>l_checksum</code> value and the <code>l_version</code> string to form an unique id for this shared object.

1_checksum This member's value is the sum of all externally visible sym-

bol's string names and common sizes.

1_version This member specifies the interface version. Its value is a

string table index. The interface version is a single string containing no colons (:). It is compared against a colon separated string of versions pointed to by a dynamic section entry of the shared object. Shared objects with matching names are considered incompatible if the interface version strings are deemed incompatible. An index value of zero

means no version string is specified.

flags This is a set of 1 bit flags. Flag definitions appear below.

Figure 5-13: Library Flags, l_flags

Name Value	Meaning	
LL_EXACT_MATCH LL_IGNORE_INT_VER		require exact match ignore interface version

LL_EXACT_MATCH At run-time use a unique id composed of the

1_time_stamp, 1_checksum, and 1_version fields to demand that the run-time dynamic shared library match exactly the shared library used at static link time.

LL_IGNORE_INT_VER

At run-time, ignore any version incompatibilities between the dynamic shared library and the library used at static link time.

Normally, if neither LL_EXACT_MATCH nor LL_IGNORE_INT_VER bits are set, the dynamic linker requires that the version of the dynamic shared library match at least one of the colon separated version strings indexed by the l_version string table index.

Conflict Section

The .conflict section is an array of indexes into the .dynsym section. Each index identifies a symbol whose attributes conflict with a shared object on which it depends, either in type or size such that this definition will preempt the shared object's definition. The dependent shared object is identified at static link time.

Figure 5-14: Conflict Section

typedef Elf32_Addr Elf32_Conflict;

System Library

Additional Entry Points

The following routines are included in the **libsys** library to provide entry points for the required source-level interfaces listed in the *System V ABI*. A description and syntax summary for each function follows the table.

Figure 6-1: libsys Additional Required Entry Points

_fxstat _lxstat _xmknod _xstat nuname _nuname

int _ fxstat (int, int, struct stat *);

The semantics of this function are identical to those of the fstat (BA_OS) function described in the *System V Interface Definition, Third Edition*. Its only difference is that it requires an extra first argument whose value must be 2.

int _lxstat (int, char *, struct stat *);

The semantics of this function are identical to those of the lstat (BA_OS) function described in the *System V Interface Definition*, *Third Edition*. Its only difference is that it requires an extra first argument whose value must be 2.

int _nuname (struct utsname *);

The semantics and syntax of this function are identical to those of the uname(BA_OS) function described in the *System V Interface Definition,*-*Third Edition*. The symbol _nuname is also available with the same semantics.

int _xmknod(int, char *, mode_t, dev_t);

The semantics and syntax of this function are identical to those of the mknod(BA_OS) function described in the *System V Interface Definition*,-*Third Edition*. Its only difference is that it requires an extra first argument whose value must be 2.

int _xstat(int, char *, struct stat *);

The semantics of this function are identical to those of the stat(BA_OS) function described in the *System V Interface Definition*,

6-1 LIBRARIES

Third Edition. Its only difference is that it requires an extra first argument whose value must be 2.

Support Routines

Besides operating system services, **libsys** contains the following processor-specific support routines.

Figure 6-2: libsys Support Routines

```
sbrk _sbrk _sqrt_s _sqrt_d _test_and_set _flush_cache
```

The routines listed below employ the standard calling sequence described in Chapter 3, "Function Calling Sequence."

char *sbrk(int incr);

This function adds <code>incr</code> bytes to the <code>break value</code> and changes the allocated space accordingly. <code>Incr</code> can be negative, in which case the amount of allocated space is decreased. The break value is the address of the first allocation beyond the end of the data segment. The amount of allocated space increases as the break value increases. Newly allocated space is set to zero. If, however, the same memory is reallocated to the same process, its contents are undefined. Upon successful completion, <code>sbrk</code> returns the old break value. Otherwise, it returns -1 and sets <code>errno</code> to indicate the error. The symbol <code>_sbrk</code> is also available with the same semantics. NOTE: mixing <code>sbrk</code> & malloc is hazardous to your program's health.

float _sqrt_s(float v)

This function computes $\sqrt{\nu}$ using single-precision floating point arithmetic and returns the resulting value. The result is rounded as if calculated to infinite precision and then rounded to single-precision according to the current rounding modes specified by the floating point control/status register. If the value is -0, the result is -0. $_{sqrt_s}$ can trigger the floating point exceptions *Invalid Operation* when v is less than 0 or *Inexact*.

double _sqrt_d(double v)

This function computes \sqrt{v} using double-precision floating point arithmetic and returns the resulting value. The result is rounded as if calculated to infinite precision and then rounded to double-precision according to the current rounding modes specified by the floating point con-

trol/status register. If the value is -0, the result is -0. _sqrt_d can trigger the floating point exceptions *Invalid Operation* when v is less than 0 or *Inexact*.

```
int _test_and_set(int *p, int v)
```

This function performs an atomic *test and set* operation on the integer pointed to by p. It effectively performs the following operations, but with a guarantee that no other process executing on the system can interrupt the operation.

```
temp = *p;
*p = v;
return(temp);
```

ICACHE - Flush only the instruction cache.

DCACHE - Flush only the data cache.

BCACHE - Flush both instruction and data cache.

These definitions are in the include file <sys/cachectl.h>. The function returns zero when no errors are detected and returns -1 otherwise, with the error cause indicated in errno. On error, the two possible errno values are either EINVAL, indicating an invalid value for the cache parameter, or EFAULT, indicating some part or all of the address range specified is not accessable.

Global Data Symbols

The **libsys** library requires that some global external data objects be defined for the routines to work properly. In addition to the corresponding data symbols listed in the *System V ABI*, the following symbols must be provided in the system library on all *ABI*-conforming systems implemented with the MIPS processor architecture. Declarations for the data objects listed below can be found in the "Data Definitions" section.

LIBRARIES 6-3

Figure 6-3: libsys, Global External Data Symbols

__huge_val

Application Constraints

As described above, *libsys* provides symbols for applications. In a few cases, however, an application must provide symbols for the library. In addition to the application-provided symbols listed in this section of the *System V ABI*, conforming applications on the MIPS processor architecture are also required to provide the following symbols.

extern _end;

This symbol refers neither to a routine nor to a location with interesting contents. Instead, its address must correspond to the beginning of the dynamic allocation area of a program, called the *heap*. Typically, the heap begins immediately after the data segment of the program executable file.

extern _gp;

This symbol is defined by the link editor and provides the value used for the gp register for this executable or shared object file.

extern const int _lib_version;

This variable's value specifies the compilation and execution mode for the program. If the value is zero, the program preserves the semantics of older (pre-ANSI) C, where conflicts exist with ANSI. Otherwise, the value is non-zero, and the program requires ANSI C semantics.

extern DYNAMIC LINKING;

This variable is a flag that the static linker sets to non-zero if the object is dynamically linked and is capable of linking with other dynamic shared objects at run time. The value is set to zero otherwise.

System Data Interfaces

Data Definitions

This section contains standard header files that describe system data. These header files are referred to by their names in angle brackets: < name.h> and < sys/name.h>. Included in these header files are macro and data definitions.

The data objects described in this section are part of the interface between an ABI-conforming application and the underlying *ABI*-conforming system where it runs. While an *ABI*-conforming system must provide these interfaces, it is not required to contain the actual header files referenced here.

ANSI C serves as the *ABI* reference programming language, and data definitions are specificed in ANSI C format. The C language is used here as a convenient notation. Using a C language description of these data objects does *not* preclude their use by other programming languages.

Figure 6-4: <assert.h>

```
extern void __assert(const char *, const char *, int);
#define assert(EX) (void)((EX)||(__assert(#EX, __FILE__, __LINE__), 0))
```

Figure 6-5: <sys/cachectl.h>

```
#define ICACHE 0x1
#define DCACHE 0x2
#define BCACHE (ICACHE | DCACHE)
```

LIBRARIES 6-5

Figure 6-6: <ctype.h>

```
#define U
                       01
#define _L
                       02
#define _N
                       04
#define _S
                       010
#define _P
                       020
#define _C
                       040
#define _B
                       0100
#define _X
                       0200
extern unsigned char
                      __ctype[];
#define isalpha(c)
                       ((__ctype+1)[c]&(_U|_L))
#define isupper(c)
                       ((__ctype+1)[c]&_U)
#define islower(c)
                       ((__ctype+1)[c]&_L)
#define isdigit(c)
                       ((__ctype+1)[c]&_N)
#define isxdigit(c)
                       ((__ctype+1)[c]&_X)
#define isalnum(c)
                       ((\_ctype+1)[c]&(\_U|\_L|\_N))
#define isspace(c)
                       ((__ctype+1)[c]&_S)
#define ispunct(c)
                       ((__ctype+1)[c]&_P)
#define isprint(c)
                       ((\_ctype+1)[c]&(\_P|\_U|\_L|\_N|\_B))
#define isgraph(c)
                       ((__ctype+1)[c]&(_P|_U|_L|_N))
#define iscntrl(c)
                       ((__ctype+1)[c]&_C)
#define isascii(c)
                       (!((c)&\sim0177))
#define _toupper(c)
                       ((__ctype+258)[c])
#define _tolower(c)
                       ((__ctype+258)[c])
#define toascii(c)
                       ((c) \& 0177)
```

Figure 6-7: <dirent.h>

```
typedef struct {
                               dd_fd;
            int
                               dd_loc;
            int
            int
                               dd_size;
                               *dd_buf;
            char
} DIR;
struct dirent {
            ino_t
                               d_ino;
            off_t
                               d_off;
            unsigned short
                              d_reclen;
            char
                               d_name[1];
};
```

LIBRARIES 6-7

Figure 6-8: <errno.h>

	.nt errno;		
#define	EPERM	1	
#define	ENOENT	2	
#define	ESRCH	3	
#define	EINTR	4	
#define	EIO	5	
#define	ENXIO	6	
#define	E2BIG	7	
#define	ENOEXEC	8	
#define	EBADF	9	
#define	ECHILD	10	
#define	EAGAIN	11	
#define	ENOMEM	12	
#define	EACCES	13	
#define	EFAULT	14	
#define	ENOTBLK	15	
#define	EBUSY	16	
#define	EEXIST	17	
#define	EXDEV	18	
#define	ENODEV	19	
#define	ENOTDIR	20	
#define	EISDIR	21	
#define	EINVAL	22	
#define	ENFILE	23	
#define	EMFILE	24	
#define	ENOTTY	25	
#define	ETXTBSY	26	
#define	EFBIG	27	
#define	ENOSPC	28	
#define	ESPIPE	29	

Figure 6-8: <errno.h> (continued)

#define	EROFS	30	
#define	EMLINK	31	
#define	EPIPE	32	
#define	EDOM	33	
#define	ERANGE	34	
#define	ENOMSG	35	
#define	EIDRM	36	
#define	ECHRNG	37	
#define	EL2NSYNC	38	
#define	EL3HLT	39	
#define	EL3RST	40	
#define	ELNRNG	41	
#define	EUNATCH	42	
#define	ENOCSI	43	
#define	EL2HLT	44	
#define	EDEADLK	45	
#define	ENOLCK	46	
#define	ENOSTR	60	
#define	ENODATA	61	
#define	ETIME	62	
#define	ENOSR	63	
#define	ENONET	64	
#define	ENOPKG	65	
#define	EREMOTE	66	
#define	ENOLINK	67	
#define	EADV	68	
#define	ESRMNT	69	

Figure 6-8: <errno.h> (continued)

#define ECOMM	70
#define EPROTO	71
#define EMULTIHOP	74
#define EBADMSG	77
#define ENAMETOOLONG	78
#define EOVERFLOW	79
#define ENOTUNIQ	80
#define EBADFD	81
#define EREMCHG	82
#define ENOSYS	89
#define ELOOP	90
#define ERESTART	91
#define ESTRPIPE	92
#define ENOTEMPTY	93
#define EUSERS	94
#define ECONNABORTED	130
#define ECONNRESET	131
#define ECONNREFUSED	146
#define ESTALE	151

```
Figure 6-9: <fcntl.h>
```

```
#define O_RDONLY
                        0
#define O_WRONLY
                        1
#define O_RDWR
                        2
#define O_APPEND
                        0x08
#define O_SYNC
                        0x10
#define O_NONBLOCK
                        0x80
#define O_CREAT
                        0x100
#define O_TRUNC
                        0x200
#define O_EXCL
                        0x400
#define O_NOCTTY
                        0x800
#define F_DUPFD
                        0
#define F_GETFD
                       1
#define F_SETFD
                        2
#define F_GETFL
#define F_SETFL
                        4
#define F_GETLK
                       14
#define F_SETLK
                        6
#define F_SETLKW
                        7
#define FD_CLOEXEC
#define O_ACCMODE
typedef struct flock {
      short
                        l_type;
      short
                       l_whence;
                       l_start;
      off_t
      off_t
                       l_len;
      long
                       l_sysid;
      pid_t
                       l_pid;
      long
                       pad[4];
} flock_t;
                        01
#define F_RDLCK
#define F_WRLCK
                        02
#define F_UNLCK
                        03
```

Figure 6-10: <float.h>

```
extern int __flt_rounds;
#define FLT_ROUNDS __flt_rounds
```

```
Figure 6-11: <fmtmsg.h>
```

```
#define MM_NULL
                        0L
#define MM_HARD
                        0x0000001L
#define MM_SOFT
                        0x00000002L
#define MM_FIRM
                         0x00000004L
#define MM_RECOVER
                        0x00000100L
#define MM_NRECOV
                        0x00000200L
#define MM APPL
                        0x0000008L
#define MM_UTIL
                        0x0000010L
#define MM_OPSYS
                        0x00000020L
#define MM PRINT
                        0x00000040L
#define MM_CONSOLE
                        0x00000080L
#define MM_NOSEV
                        0
#define MM_HALT
                        1
#define MM_ERROR
#define MM_WARNING
                        3
#define MM_INFO
#define MM_NULLLBL
                         ((char *) NULL)
#define MM_NULLSEV
                        MM NOSEV
#define MM_NULLMC
                        MM_NULL
#define MM_NULLTXT
                        ((char *) NULL)
#define MM NULLACT
                         ((char *) NULL)
#define MM_NULLTAG
                        ((char *) NULL)
#define MM_NOTOK
                        -1
#define MM_OK
                         0 \times 0
#define MM_NOMSG
                         0x01
#define MM_NOCON
                         0x04
```

Figure 6-12: <ftw.h>

```
#define FTW_PHYS
                                    01
#define FTW_MOUNT
                                    02
#define FTW_CHDIR
                                    04
#define FTW_DEPTH
                                 0 10
#define FTW_F
                                    0
#define FTW_D
                                    1
                                    2
#define FTW_DNR
#define FTW_NS
                                    3
#define FTW_SL
                                    4
#define FTW_DP
                                    6
struct FTW
                         int
                                 quit;
                         int
                                 base;
                                  level;
                         int
};
```

Figure 6-13: <grp.h>

Figure 6-14: <sys/ipc.h>

```
struct ipc_perm {
                         uid;
         uid_t
         gid_t
                         gid;
                         cuid;
         uid_t
         gid_t
                         cgid;
                         mode;
         mode_t
         unsigned long
                         seq;
                         key;
         key_t
         long
                         pad[4];
};
#define IPC_CREAT
                         0001000
#define IPC_EXCL
                         0002000
#define IPC_NOWAIT
                         0004000
#define IPC_PRIVATE
                         (key_t)0
#define IPC_RMID
                         10
#define IPC_SET
                         11
#define IPC_STAT
                         12
```

Figure 6-15: <langinfo.h>

#define DAY_1 1 #define DAY_2 2 #define DAY_3 3	
#define DAY_3 3	
<u> </u>	
U 3 C ' D 3 T 4	
#define DAY_4 4	
#define DAY_5 5	
#define DAY_6 6	
#define DAY_7 7	
#define ABDAY_1 8	
#define ABDAY_2 9	
#define ABDAY_3 10	
#define ABDAY_4 11	
#define ABDAY_5 12	
#define ABDAY_6 13	
#define ABDAY_7 14	
#define MON_1 15	
#define MON_2 16	
#define MON_3 17	
#define MON_4 18	
#define MON_5 19	
#define MON_6 20	
#define MON_7 21	
#define MON_8 22	
#define MON_9 23	
#define MON_10 24	
#define MON_11 25	
#define MON_12 26	

Figure 6-15: <langinfo.h> (continued)

#define	ABMON_1	27
#define	ABMON_2	28
#define	ABMON_3	29
#define	ABMON_4	30
#define	ABMON_5	31
#define	ABMON_6	32
#define	ABMON_7	33
#define	ABMON_8	34
#define	ABMON_9	35
#define	ABMON_10	36
#define	ABMON_11	37
#define	ABMON_12	38
#define	RADIXCHAR	39
#define	THOUSEP	40
#define	YESSTR	41
#define	NOSTR	42
#define	CRNCYSTR	43
#define	D_T_FMT	44
#define	D_FMT	45
#define	T_FMT	46
#define	AM_STR	47
#define	PM_STR	48

Figure 6-16: <limits.h>

```
#define MB_LEN_MAX
#define ARG_MAX
#define CHILD_MAX
#define MAX_CANON
#define NGROUPS_MAX
#define LINK_MAX
#define NAME_MAX
#define OPEN_MAX
#define PASS_MAX
#define PATH_MAX
#define PIPE_MAX
#define PIPE_BUF
#define MAX_INPUT
 /* starred values vary and should be
    retrieved using sysconf() or pathconf()
#define NL_ARGMAX
#define NL_LANGMAX
                        14
#define NL_MSGMAX
                        32767
#define NL_NMAX
#define NL_SETMAX
                        255
#define NL_TEXTMAX
                        255
#define NZERO
                        20
#define TMP_MAX
                        17576
#define FCHR_MAX
                        2147483647
```

Figure 6-17: <locale.h>

```
struct lconv {
    char *decimal_point;
    char *thousands_sep;
    char *grouping;
    char *int_curr_symbol;
    char *currency_symbol;
    char *mon_decimal_point;
    char *mon_thousands_sep;
    char *mon_grouping;
    char *positive_sign;
    char *negative_sign;
    char int_frac_digits;
    char frac_digits;
    char p_cs_precedes;
    char p_sep_by_space;
    char n_cs_precedes;
    char n_sep_by_space;
    char p_sign_posn;
    char n_sign_posn;
};
#define LC_CTYPE
                           0
#define LC_NUMERIC
                           1
#define LC_TIME
                           2
#define LC_COLLATE
                           3
#define LC MONETARY
                           4
#define LC_MESSAGES
                           5
#define LC_ALL
                           6
#define NULL
                           0
```

Figure 6-18: <math.h>

```
typedef union _h_val {
    unsigned long i[2];
    double d;
} _h_val;

extern const _h_val __huge_val;
#define HUGE_VAL __huge_val.d
```

Figure 6-19: <sys/mman.h>

```
#define PROT_READ
                        0x1
#define PROT_WRITE
                        0x2
#define PROT_EXEC
                        0x4
#define PROT_NONE
                        0x0
#define MAP_SHARED
                        1
#define MAP_PRIVATE
                        2
#define MAP_FIXED
                        0x10
#define MS_SYNC
                        0x0
#define MS_ASYNC
                        0x1
#define MS_INVALIDATE
                        0x2
```

Figure 6-20: <sys/mount.h>

```
#define MS_RDONLY 0x01
#define MS_DATA 0x04
#define MS_NOSUID 0x10
#define MS_REMOUNT 0x20
```

Figure 6-21: <sys/msg.h>

```
struct msqid_ds {
         struct ipc_perm
                            msg_perm;
                            *msg_first;
         struct msg
         struct msg
                            *msg_last;
         unsigned long
                            msg_cbytes;
         unsigned long
                            msg_qnum;
         unsigned long
                            msg_qbytes;
                            msg_lspid;
         pid_t
         pid_t
                            msg_lrpid;
         time_t
                            msg_stime;
                            msg_pad1;
         long
         time_t
                            msg_rtime;
         long
                            msg_pad2;
                            msg_ctime;
         time_t
         long
                            msg_pad3;
         long
                            msg_pad4[4];
};
#define MSG_NOERROR
                            010000
```

Figure 6-22: <netconfig.h>

```
struct netconfig{
     char
                         *nc_netid;
     unsigned long
                         nc_semantics;
     unsigned long
                         nc_flag;
     char
                         *nc_protofmly;
     char
                         *nc_proto;
     char
                         *nc_device;
     unsigned long
                         nc_nlookups;
                         **nc_lookups;
     char
     unsigned long
                         nc_unused[8];
};
#define NC_TPI_CLTS
                         1
#define NC_TPI_COTS
                         2
#define NC_TPI_COTS_ORD 3
                         4
#define NC_TPI_RAW
#define NC_NOFLAG
                         00
#define NC_VISIBLE
                         01
```

Figure 6-22: <netconfig.h> (continued)

```
#define NC NOPROTOFMLY
                         "-"
#define NC_LOOPBACK
                         "loopback"
#define NC_INET
                         "inet"
#define NC_IMPLINK
                         "implink"
#define NC_PUP
                         "pup"
#define NC_CHAOS
                         "chaos"
#define NC NS
                         "ns"
#define NC_NBS
                         "nbs"
#define NC_ECMA
                         "ecma"
#define NC DATAKIT
                         "datakit"
#define NC_CCITT
                         "ccitt"
#define NC_SNA
                         "sna"
#define NC_DECNET
                         "decnet"
#define NC_DLI
                         "dli"
#define NC_LAT
                         "lat"
#define NC_HYLINK
                         "hylink"
#define NC_APPLETALK
                         "appletalk"
#define NC_NIT
                         "nit"
#define NC_IEEE802
                         "ieee802"
#define NC_OSI
                         "osi"
#define NC_X25
                         "x25"
#define NC_OSINET
                         "osinet"
#define NC GOSIP
                         "qosip"
#define NC_NOPROTO
                         " – "
#define NC_TCP
                         "tcp"
#define NC_UDP
                         "udp"
#define NC_ICMP
                         "icmp"
```

Figure 6-23: <netdir.h>

```
struct nd_addrlist{
     int
                          n_cnt;
     struct netbuf
                          *n_addrs;
};
 struct nd_hostservlist {
     int
                          h_cnt;
     struct nd_hostserv
                          *h_hostservs;
};
 struct nd_hostserv {
     char *h_host;
     char *h_serv;
};
#define ND_BADARG
                          -2
#define ND_NOMEM
                          -1
#define ND_OK
                          0
#define ND_NOHOST
                          1
#define ND_NOSERV
                          2
#define ND_NOSYM
                          3
#define ND_OPEN
#define ND_ACCESS
                          5
#define ND_UKNWN
                          6
                          7
#define ND_NOCTRL
#define ND_FAILCTRL
                          8
#define ND_SYSTEM
```

Figure 6-23: <netdir.h> (continued)

```
#define ND_HOSTSERV
                                 0
#define ND_HOSTSERVLIST
                                1
#define ND_ADDR
#define ND_ADDRLIST
                                3
#define HOST_SELF
                                 "\\1"
                                 "\\2"
#define HOST_ANY
#define HOST_BROADCAST
                                 "\\3"
#define ND_SET_BROADCAST
                                1
#define ND_SET_RESERVEDPORT
                                2
#define ND_CHECK_RESERVEDPORT
                                3
#define ND_MERGEADDR
```

Figure 6-24: <nl_types.h>

```
#define NL_SETD 1

typedef int nl_item ;
typedef void *nl_catd;
```

Figure 6-25: <sys/param.h>

#define	CANBSIZ	256
#define	HZ	100
#define	NGROUPS_UMIN	0
#define	MAXPATHLEN	1024
#define	MAXSYMLINKS	30
#define	MAXNAMELEN	256
#define	NADDR	13
#define	NBBY	8
#define	NBPSCTR	512

Figure 6-26: <poll.h>

```
struct pollfd {
                         int fd;
                         short events;
                         short revents;
};
#define POLLIN
                         0x0001
#define POLLPRI
                         0x0002
#define POLLOUT
                         0x0004
#define POLLRDNORM
                         0 \times 0040
#define POLLWRNORM
                         POLLOUT
#define POLLRDBAND
                         0x0080
#define POLLWRBAND
                         0 \times 0100
#define POLLNORM
                         POLLRDNORM
#define POLLERR
                         8000x0
#define POLLHUP
                         0x0010
#define POLLNVAL
                         0x0020
```

Figure 6-27: <sys/procset.h>

```
#define P_INITPID
                                 1
#define P_INITUID
                                 0
#define P_INITPGID
                                  0
typedef long id_t;
typedef enum idtype{
                        P_PID,
                         P_PPID,
                        P_PGID,
                         P_SID,
                         P_CID,
                        P_UID,
                         P_GID,
                         P_ALL
} idtype_t;
typedef enum idop {
                         POP_DIFF,
                         POP_AND,
                        POP_OR,
                         POP_XOR
} idop_t;
```

Figure 6-27: <sys/procset.h> (continued)

Figure 6-28: <pwd.h>

```
struct passwd {
         char
                    *pw_name;
         char
                    *pw_passwd;
         uid_t
                    pw_uid;
         gid_t
                    pw_gid;
         char
                    *pw_age;
         char
                    *pw_comment;
         char
                    *pw_gecos;
         char
                    *pw_dir;
         char
                    *pw_shell;
```

Figure 6-29: <sys/resource.h>

```
#define RLIMIT_CPU
                               0
#define RLIMIT_FSIZE
                               1
                               2
#define RLIMIT_DATA
                               3
#define RLIMIT_STACK
#define RLIMIT_CORE
#define RLIMIT_NOFILE
#define RLIMIT_VMEM
                               6
#define RLIMIT_AS
                               RLIMIT_VMEM
#define ELIM_INFINITY
                               0x7fffffff
typedef unsigned long rlim_t;
struct rlimit{
                       rlim_t rlim_cur;
                       rlim_t rlim_max;
```

Figure 6-30: <pre

```
400
#define MAX_AUTH_BYTES
#define MAXNETNAMELEN
                           255
#define HEXKEYBYTES
                           48
 enum auth_stat{
           AUTH_OK=0,
           AUTH_BADCRED=1,
           AUTH_REJECTEDCRED=2,
           AUTH_BADVERF=3,
           AUTH_REJECTEDVERF=4,
           AUTH_TOOWEAK=5,
           AUTH_INVALIDRESP=6,
           AUTH_FAILED=7
};
union des_block{
           struct {
           unsigned long high;
           unsigned long low;
           } key;
           char c[8];
struct opaque_auth{
                          oa_flavor;
           int
           char
                           *oa_base;
           unsigned int
                          oa_length;
};
```

```
typedef struct {
         struct
                       opaque_auth ah_cred;
                       opaque_auth ah_verf;
         struct
                       des_block ah_key;
         union
         struct auth_ops {
                      (*ah_nextverf)();
         void
         int
                      (*ah_marshal)();
         int
                      (*ah_validate)();
         int
                       (*ah_refresh)();
         void
                       (*ah_destroy)();
         } *ah_ops;
         char *ah_private;
} AUTH;
struct authsys_parms{
         unsigned long aup_time;
         char
                       *aup_machname;
         uid_t
                       aup_uid;
         gid_t
                       aup_gid;
         unsigned int aup_len;
         gid_t
                       *aup_gids;
};
extern struct opaque_auth_null_auth;
#define AUTH_NONE
                       0
#define AUTH_NULL
                       0
#define AUTH_SYS
                       1
#define AUTH_UNIX
                       AUTH_SYS
#define AUTH SHORT
#define AUTH_DES
                       3
```



```
enum clnt_stat{
      RPC_SUCCESS=0,
      RPC_CANTENCODEARGS=1,
      RPC_CANTDECODERES=2,
      RPC_CANTSEND=3,
      RPC_CANTRECV=4,
      RPC_TIMEDOUT=5,
      RPC_INTR=18,
      RPC_UDERROR=23,
      RPC_VERSMISMATCH=6,
      RPC_AUTHERROR=7,
      RPC_PROGUNAVAIL=8,
      RPC_PROGVERSMISMATCH=9,
      RPC_PROCUNAVAIL=10,
      RPC_CANTDECODEARGS=11,
      RPC_SYSTEMERROR=12,
      RPC_UNKNOWNHOST=13,
      RPC_UNKNOWNPROTO=17,
      RPC_UNKNOWNADDR=19,
      RPC_NOBROADCAST=21,
      RPC_RPCBFAILURE=14,
      RPC_PROGNOTREGISTERED=15,
      RPC_N2AXLATEFAILURE=22,
      RPC_TLIERROR=20,
      RPC_FAILED=16
};
#define RPC_PMAPFAILURE RPC_RPCBFAILURE
```

Figure 6-30: c.h> (continued)

```
#define RPC_AYSOCK -1
#define RPC_ANYFD RPC_ANYSOCK
struct rpc_err{
      enum clnt_stat re_status;
      union {
      struct {
      int errno;
      int t_errno;
      } RE_err;
      enum auth_stat RE_why;
      struct {
      unsigned long low;
      unsigned long high;
      } RE_vers;
      struct {
      long s1;
      long s2;
      } RE_lb;
      } ru;
};
```

```
struct rpc_createerr{
     enum clnt_stat cf_stat;
      struct rpc_err cf_error;
};
typedef struct {
     AUTH
              *cl_auth;
      struct clnt_ops {
              enum clnt_stat (*cl_call)();
              void
                                (*cl_abort)();
              void
                                (*cl_geterr)();
              int
                                (*cl_freeres)();
              void
                                (*cl_destroy)();
              int
                                (*cl_control)();
      } *cl_ops;
              *cl_private;
      char
              *cl netid;
      char
      char
              *cl_tp;
} CLIENT;
#define FEEDBACK_REXMIT1
                                1
#define FEEDBACK_OK
#define CLSET_TIMEOUT
                                1
#define CLGET_TIMEOUT
#define CLGET_SERVER_ADDR
                                3
#define CLGET_FD
                                6
#define CLGET_SVC_ADDR
                                7
#define CLSET_FD_CLOSE
                                8
#define CLSET_FD_NCLOSE
                                9
                                4
#define CLSET_RETRY_TIMEOUT
#define CLGET_RETRY_TIMEOUT
```

```
extern struct
rpc_createerr rpc_createerr;
enum xprt_stat{
     XPRT_DIED,
     XPRT MOREREQS,
     XPRT_IDLE
};
typedef struct {
     int xp_fd;
     unsigned short xp_port;
     struct xp_ops {
     int
                          (*xp_recv)();
     enum xprt_stat
                          (*xp_stat)();
     int
                          (*xp_getargs)();
     int
                          (*xp_reply)();
     int
                          (*xp_freeargs)();
     void
                          (*xp_destroy)();
     } *xp_ops;
     int
                          xp_addrlen;
                          *xp_tp;
     char
     char
                          *xp_netid;
     struct netbuf
                          xp_ltaddr;
     struct netbuf
                          xp_rtaddr;
     char
                          xp_raddr[16];
     struct opaque_auth xp_verf;
     char
                          *xp_p1;
     char
                          *xp_p2;
     char
                          *xp_p3;
} SVCXPRT;
```

```
struct svc_req {
         unsigned long rq_prog;
         unsigned long rq_vers;
         unsigned long rq_proc;
         struct opaque_auth rq_cred;
                  *rq_clntcred;
         char
         SVCXPRT
                    *rq_xprt;
};
typedef struct fdset{
         long fds_bits[32];
} fd_set;
 extern fd_set svc_fdset;
 enum msg_type{
         CALL=0,
         REPLY=1
};
enum reply_stat{
         MSG_ACCEPTED=0,
         MSG_DENIED=1
};
 enum accept_stat{
         SUCCESS=0,
         PROG_UNAVAIL=1,
         PROG_MISMATCH=2,
         PROC_UNAVAIL=3,
         GARBAGE_ARGS=4,
         SYSTEM_ERR=5
};
```

Figure 6-30: c.h> (continued)

```
enum reject_stat {
     RPC_MISMATCH=0,
      AUTH_ERROR=1
};
 struct accepted_reply{
      struct opaque_auth ar_verf;
      enum accept_stat ar_stat;
      union {
        struct {
               unsigned long low;
               unsigned long high;
        } AR_versions;
        struct {
               char *where;
               xdrproc_t proc;
        } AR_results;
      } ru;
};
struct rejected_reply{
      enum reject_stat rj_stat;
      union {
        struct {
               unsigned long low;
               unsigned long high;
        } RJ_versions;
        enum auth_stat RJ_why;
      } ru;
};
```

```
struct reply_body{
      enum reply_stat rp_stat;
      union {
        struct accepted_reply RP_ar;
        struct rejected_reply RP_dr;
      } ru;
};
struct call_body{
     unsigned long cb_rpcvers;
      unsigned long cb_prog;
      unsigned long cb_vers;
      unsigned long cb_proc;
      struct opaque_auth cb_cred;
      struct opaque_auth cb_verf;
};
struct rpc_msg{
      unsigned long rm_xid;
      enum msg_type rm_direction;
      union {
        struct call_body RM_cmb;
        struct reply_body RM_rmb;
      } ru;
};
struct rpcb{
      unsigned long r_prog;
      unsigned long r_vers;
      char *r_netid;
      char *r_addr;
      char *r_owner;
};
```

Figure 6-30: c.h> (continued)

```
struct rpcblist{
      struct rpcb rpcb_map;
      struct rpcblist *rpcb_next;
};
 enum xdr_op {
      XDR_ENCODE=0,
      XDR_DECODE=1,
      XDR_FREE=2
};
 struct xdr_discrim{
      int value;
      xdrproc_t proc;
};
enum authdes_namekind {
      ADN_FULLNAME,
      ADN_NICKNAME
};
struct authdes_fullname{
      char *name;
      union des block key;
      unsigned long window;
};
struct authdes_cred{
      enum authdes_namekind adc_namekind;
      struct authdes_fullname adc_fullname;
      unsigned long adc_nickname;
};
```

Figure 6-30: c.h> (continued)

```
typedef struct {
      enum xdr_op
                        x_op;
      struct xdr_ops{
         int
                       (*x_getlong)();
         int
                        (*x_putlong)();
         int
                        (*x_getbytes)();
         int
                       (*x_putbytes)();
         unsigned int
                       (*x_getpostn)();
         int
                        (*x_setpostn)();
         long *
                        (*x_inline)();
         void
                        (*x_destroy)();
      } *x_ops;
      char *x_public;
      char *x_private;
      char *x_base;
            x_handy;
      int
} XDR;
typedef int (*xdrproc_t)()
#define NULL_xdrproc_t ((xdrproc_t)0)
```



```
#define auth_destroy(auth)
  ((*((auth)->ah_ops->ah_destroy))(auth))
#define clnt_call(rh, proc, xargs, argsp, xres, resp, secs)
  ((*(rh)->cl_ops->cl_call)(rh, proc, xargs, \
  argsp, xres, resp, secs))
#define clnt_freeres(rh,xres,resp)
  ((*(rh)->cl_ops->cl_freeres)(rh,xres,resp))
#define clnt_geterr(rh, errp)
  ((*(rh)->cl_ops->cl_geterr)(rh, errp))
#define clnt_control(cl, rq, in)
  ((*(cl)->cl_ops->cl_control)(cl, rq, in))
#define clnt_destroy(rh)
  ((*(rh)->cl_ops->cl_destroy)(rh))
#define svc_destroy(xprt)
  (*(xprt)->xp_ops->xp_destroy)(xprt)
#define svc_freeargs(xprt, xargs, argsp)
  (*(xprt)->xp ops->xp freearqs)((xprt), (xarqs), (arqsp))
#define svc_getargs(xprt, xargs, argsp)
  (*(xprt)->xp_ops->xp_getargs)((xprt), (xargs), (argsp))
#define svc_getrpccaller(x)
  (&(x)->xp_rtaddr)
#define xdr_getpos(xdrs)
  (*(xdrs)->x_ops->x_getpostn)(xdrs)
#define xdr_setpos(xdrs, pos)
  (*(xdrs)->x_ops->x_setpostn)(xdrs, pos)
#define xdr inline(xdrs, len)
  (*(xdrs)->x_ops->x_inline)(xdrs, len)
#define xdr_destroy(xdrs)
  (*(xdrs)->x_ops->x_destroy)(xdrs)
```

Figure 6-31: <search.h>

```
typedef struct entry { char *key; void *data;} ENTRY;
typedef enum { FIND, ENTER} ACTION;
typedef enum { preorder, postorder, endorder, leaf} VISIT;
```

Figure 6-32: <sys/sem.h>

```
#define SEM_UNDO
                    010000
#define GETNCNT
                    3
#define GETPID
                    4
#define GETVAL
                    5
#define GETALL
                    6
#define GETZCNT
                    8
#define SETVAL
#define SETALL
                    9
 struct semid_ds {
         struct ipc_perm
                               sem_perm;
         struct sem
                               *sem_base;
         unsigned short
                               sem_nsems;
         time_t
                               sem_otime;
         long
                               sem_pad1;
         time_t
                               sem_ctime;
         long
                               sem_pad2;
         long
                               sem_pad3[4];
};
 struct sem {
         unsigned short
                               semval;
                               sempid;
         pid_t
         unsigned short
                               semncnt;
         unsigned short
                               semzcnt;
};
 struct sembuf {
         unsigned short
                               sem_num;
         short
                               sem_op;
         short
                               sem_flg;
};
```

Figure 6-33: <setjmp.h>

Figure 6-34: <sys/shm.h>

```
struct shmid_ds{
    struct ipc_perm
                          shm_perm;
    int
                          shm_segsz;
    char
                          *shm_amp;
    unsigned short
                          shm_lkcnt;
    pid_t
                          shm_lpid;
    pid_t
                          shm_cpid;
    unsigned long
                          shm_nattch;
    unsigned long
                          shm_cnattch;
    time_t
                          shm_atime;
    long
                          shm_pad1;
                          shm_dtime;
    time_t
    long
                          shm_pad2;
                          shm_ctime;
    time_t
    long
                          shm_pad3;
    long
                          shm_pad4[4];
};
#define SHM_RDONLY
                          010000
#define SHM_RND
                          020000
```

Figure 6-35: <signal.h>

#define	SIGHUP	1	
#define	SIGINT	2	
#define	SIGQUIT	3	
#define	SIGILL	4	
#define	SIGTRAP	5	
#define	SIGABRT	6	
#define	SIGEMT	7	
#define	SIGFPE	8	
#define	SIGKILL	9	
#define	SIGBUS	10	
#define	SIGSEGV	11	
#define	SIGSYS	12	
#define	SIGPIPE	13	
#define	SIGALRM	14	
#define	SIGTERM	15	
#define	SIGUSR1	16	
#define	SIGUSR2	17	
#define	SIGCHLD	18	
#define	SIGPWR	19	
#define	SIGWINCH	20	
#define	SIGURG	21	
#define	SIGPOLL	22	
#define	SIGSTOP	23	
#define	SIGTSTP	24	
#define	SIGCONT	25	
#define	SIGTTIN	26	
#define	SIGTTOU	27	
#define	SIGXCPU	30	
#define	SIGXFSZ	31)

Figure 6-35: <signal.h (continued)

```
#define SIG BLOCK
                          2
#define SIG_UNBLOCK
#define SIG_SETMASK
                          3
#define SIG ERR
                          (void(*)())-1
#define SIG_IGN
                          (void(*)())1
#define SIG_HOLD
                          (void(*)())2
#define SIG_DFL
                          (void(*)())0
#define SS_ONSTACK
                          0x0000001
                          0x00000002
#define SS_DISABLE
struct sigaltstack {
         char
                          *ss_sp;
         int
                          ss_size;
         int
                          ss_flags;
};
typedef struct sigaltstackstack_t;
typedef struct { unsigned long sigbits[4];} sigset_t;
struct sigaction{
         int
                          sa_flags;
         void
                          (*sa_handler)();
                          sa_mask;
         sigset_t
         int
                          sa_resv[2];
};
#define SA_ONSTACK
                          0x0000001
#define SA_RESETHAND
                          0x0000002
#define SA RESTART
                          0 \times 000000004
#define SA_SIGINFO
                          0x00000008
#define SA_NOCLDWAIT
                          0x00010000
#define SA_NOCLDSTOP
                          0x00020000
```

Figure 6-36: <sys/siginfo.h>

```
#define ILL_ILLOPC
                        1
#define ILL_ILLOPN
                        2
#define ILL_ILLADR
                        3
                        4
#define ILL_ILLTRP
#define ILL_PRVOPC
                        5
#define ILL_PRVREG
                        6
                        7
#define ILL_COPROC
#define ILL_BADSTK
                        8
```

Figure 6-36: <sys/siginfo.h> (continued)

```
#define FPE_INTDIV
                        1
#define FPE_INTOVF
                         2
                         3
#define FPE_FLTDIV
#define FPE_FLTOVF
                         4
#define FPE_FLTUND
                         5
#define FPE_FLTRES
#define FPE_FLTINV
                        7
#define FPE_FLTSUB
                         8
                         1
#define SEGV_MAPERR
#define SEGV_ACCERR
                         2
#define BUS_ADRALN
                         1
#define BUS_ADRERR
#define BUS_OBJERR
                         3
#define TRAP_BRKPT
                         1
#define TRAP_TRACE
                         2
#define CLD_EXITED
                        1
#define CLD KILLED
                         2
#define CLD_DUMPED
                         3
#define CLD_TRAPPED
#define CLD_STOPPED
                         5
#define CLD_CONTINUED
                         6
#define POLL_IN
                        1
#define POLL_OUT
                         2
#define POLL_MSG
                         3
                         4
#define POLL_ERR
                         5
#define POLL_PRI
#define POLL_HUP
                         6
#define SI_MAXSZ
                         128
#define SI_PAD ((SI_MAXSZ/sizeof(int)) - 3)
```

Figure 6-36: <sys/siginfo.h> (continued)

```
typedef struct siginfo{
     int si_signo;
     int si_code;
     int si_errno;
     union {
        int _pad[SI_PAD];
        struct {
           pid_t
                  _pid;
            union {
            struct { uid_t _uid;} _kill;
            struct {
                  clock_t _utime;
                  int _status;
                  clock_t _stime;
            } _cld;
        } _pdata;
        } _proc;
       struct { char *_addr;} _fault;
        struct {
            int
                  _fd;
            long
                 _band;
        } _file;
      } _data;
} siginfo_t;
#define si_pid
                  _data._proc._pid
#define si_uid
                  _data._proc._pdata._kill._uid
#define si_addr
                 _data._fault._addr
#define si_stime _data._proc._pdata._cld._stime
#define si_utime
                 _data._proc._pdata._cld._utime
#define si_status _data._proc._pdata._cld._status
#define si_band
                 _data._file._band
#define si_fd
                  _data._file._fd
```

Figure 6-37: <sys/stat.h>

```
#define _ST_FTYPSZ 16
                stat {
    struct
                st_dev;
        dev_t
        long
                 st_pad1[3];
        ino_t
                 st_ino;
                st_mode;
        mode_t
        nlink_t st_nlink;
               st_uid;
        uid_t
        gid_t
                st_gid;
        dev_t
                st rdev;
                 st_pad2[2];
        long
        off_t
                 st_size;
        long
                 st_pad3;
        timestruc_t st_atim;
        timestruc_t st_mtim;
        timestruc_t st_ctim;
        long
                 st_blksize;
        long
                st_blocks;
        char
                 st_fstype[_ST_FSTYPSZ];
        long
                 st_pad4[8];
};
#define st_atime st_atim.tv_sec
#define st_mtime st_mtim.tv_sec
#define st_ctime st_ctim.tv_sec
```

Figure 6-37: <sys/stat.h> (continued)

```
#define S_IFMT
                        0xF000
#define S_IFIFO
                        0x1000
#define S_IFCHR
                        0x2000
#define S_IFDIR
                        0x4000
#define S_IFBLK
                        0x6000
#define S_IFREG
                        0x8000
#define S_IFLNK
                        0xA000
#define S_ISUID
                        04000
#define S_ISGID
                        02000
#define S_ISVTX
                        01000
#define S_IRWXU
                        00700
#define S_IRUSR
                        00400
#define S_IWUSR
                        00200
#define S_IXUSR
                        00100
#define S IRWXG
                        00070
#define S_IRGRP
                        00040
#define S_IWGRP
                        00020
#define S_IXGRP
                        00010
#define S_IRWXO
                        00007
#define S_IROTH
                        00004
#define S_IWOTH
                        00002
#define S_IXOTH
                        00001
#define S_ISFIFO(mode)
                         ((mode&S_IFMT) == S_IFIFO)
#define S_ISCHR(mode)
                         ((mode&S_IFMT) == S_IFCHR)
#define S_ISDIR(mode)
                         ((mode&S_IFMT) == S_IFDIR)
#define S_ISBLK(mode)
                         ((mode&S_IFMT) == S_IFBLK)
#define S_ISREG(mode)
                         ((mode&S_IFMT) == S_IFREG)
```

Figure 6-38: <sys/statvfs.h>

```
#define FSTYPSZ
                    16
typedef struct statvfs {
        unsigned long f_bsize;
        unsigned long f_frsize;
        unsigned long f_blocks;
        unsigned long f_bfree;
        unsigned long f_bavail;
        unsigned long f_files;
        unsigned long f_ffree;
        unsigned long f_favail;
        unsigned long f_fsid;
         char
                       f_basetype[FSTYPSZ];
        unsigned long f_flag;
         unsigned long f_namemax;
                       f_fstr[32];
        unsigned long f_filler[16];
} statvfs_t;
#define ST_RDONLY
                              0x01
#define ST_NOSUID
                              0x02
```

Figure 6-39: <stdarg.h>



The construction &... is a syntactic extension to ANSI C and may not be supported by all C compilers. The intended semantics are to set list to the address on the stack of the first incoming argument in the variable part of the argument list. See "Function Calling Sequence" in Chapter 3.

Figure 6-40: <stddef.h>

#define NULL 0
typedef int ptrdiff_t;
typedef unsigned int size_t;
typedef long wchar_t;

‡ The _file member of the FILE struct is moved to Level 2 as of Jan. 1, 1993.

6-56 MIPS ABI SUPPLEMENT

```
Figure 6-41: <stdio.h>
```

```
typedef unsigned int
                           size_t;
typedef long
                           fpos_t;
#define _NFILE
                           100
#define NULL
#define BUFSIZ
                           4096
#define _IOFBF
                           0000
#define _IOLBF
                           0100
#define _IONBF
                           0004
#define _IOEOF
                           0020
#define _IOERR
                           0040
#define EOF
                           (-1)
#define FOPEN_MAX
                          60
#define FILENAME_MAX
                          1024
#define stdin
                           (&__iob[0])
#define stdout
                           (&__iob[1])
#define stderr
                           (&__iob[2])
#define clearerr(p)
                           ((void)((p)->_flag &= ~(_IOERR|_I-
OEOF)))†
#define feof(p)
                           ((p)->_flag & _IOEOF)
#define ferror(p)
                           ((p)->_flag & _IOERR) †
#define fileno(p)
                           (p) \rightarrow _file
#define L_ctermid
#define L_cuserid
#define P_tmpdir
                           "/var/tmp/"
```

[†] These macro definitions are moved to Level 2 as of Jan. 1, 1993.

Figure 6-41: <stdio.h> (continued)



The macros clearerr, and fileno will be removed as a source interface in a future release supporting multi-processing. Applications should transition to the function equivalents of these macros in libc. Binary portability will be supported for existing applications.



The constant _NFILE has been removed. It should still appear in stdio.h, but may be removed in a future version of the header file. Applications may not be able to depend on fopen() failing on an attempt to open more than _NFILE files.

Figure 6-42: <stdlib.h>

```
typedef struct {
     int quot;
     int rem;
} div_t;
typedef struct {
    long quot;
     long rem;
} ldiv_t;
#define NULL
                    0
#define EXIT_FAILURE 1
#define EXIT_SUCCESS
#define RAND_MAX
                  32767
extern unsigned char __ctype[];
#define MB_CUR_MAX
                    __ctype[520]
```

Figure 6-43: <stropts.h>

	#define	SNDZERO	0x001
	#define	RNORM	0x000
	#define	RMSGD	0x001
	#define	RMSGN	0x002
	#define	RMODEMASK	0x003
	#define	RPROTDAT	0x004
	#define	RPROTDIS	0x008
	#define	RPROTNORM	0x010
	#define	FLUSHR	0x01
	#define	FLUSHW	0x02
	#define	FLUSHRW	0x03
	#define	S_INPUT	0x0001
	#define	S_HIPRI	0x0002
	#define	S_OUTPUT	0x0004
	#define	S_MSG	8000x0
	#define	S_ERROR	0x0010
	#define	S_HANGUP	0x0020
	#define	S_RDNORM	0x0040
	#define	S_WRNORM	S_OUTPUT
	#define	S_RDBAND	0x0080
	#define	S_WRBAND	0x0100
	#define	S_BANDURG	0x0200
	#define	RS_HIPRI	1
	#define	MSG_HIPRI	0x01
	#define	MSG_ANY	0x02
	#define	MSG_BAND	0x04
	#define	MORECTL	1
	#define	MOREDATA	2
	#define	MUXID_ALL	(-1)
\			

Figure 6-43: <stropts.h> (continued)

```
#define STR
                          ('S'<<8)
#define I_NREAD
                          (STR | 01)
#define I_PUSH
                          (STR | 02)
#define I_POP
                          (STR | 03)
#define I_LOOK
                          (STR | 04)
#define I_FLUSH
                          (STR | 05)
#define I_SRDOPT
                          (STR | 06)
#define I_GRDOPT
                          (STR | 07)
#define I_STR
                          (STR | 010)
#define I_SETSIG
                          (STR | 011)
#define I_GETSIG
                          (STR | 012)
#define I_FIND
                          (STR | 013)
#define I_LINK
                          (STR | 014)
#define I_UNLINK
                          (STR | 015)
#define I_PEEK
                          (STR | 017)
#define I_FDINSERT
                          (STR | 020)
#define I_SENDFD
                          (STR | 021)
#define I_RECVFD
                          (STR | 016)
#define I_SWROPT
                          (STR | 023)
#define I_GWROPT
                          (STR | 024)
#define I_LIST
                           (STR | 025)
#define I_PLINK
                          (STR | 026)
#define I_PUNLINK
                          (STR | 027)
#define I_FLUSHBAND
                          (STR | 034)
#define I_CKBAND
                          (STR | 035)
#define I_GETBAND
                          (STR | 036)
#define I_ATMARK
                          (STR | 037)
#define I_SETCLTIME
                          (STR | 040)
#define I_GETCLTIME
                          (STR | 041)
#define I_CANPUT
                           (STR | 042)
```

Figure 6-43: <stropts.h> (continued)

```
struct strioctl {
      int
               ic_cmd;
               ic_timout;
      int
      int
              ic_len;
      char
               *ic_dp;
};
struct strbuf {
      int
               maxlen;
      int
               len;
      char
               *buf;
};
struct strpeek {
      struct strbuf ctlbuf;
      struct strbuf databuf;
      long
                        flags;
};
struct strfdinsert {
      struct strbuf ctlbuf;
      struct strbuf databuf;
             flags;
      long
      int
               fildes;
      int
               offset;
};
struct strrecvfd {
      int
               fd;
      uid_t
               uid;
      gid_t
               gid;
      char
               fill[8];
};
```

Figure 6-43: <stropts.h> (continued)

```
#define FMNAMESZ
                           8
 struct str_mlist{
         char l_name[FMNAMESZ+1];
};
struct str_list{
                            sl_nmods;
         int
         struct str_mlist *sl_modlist;
};
#define ANYMARK
                           0 \times 01
#define LASTMARK
                           0 \times 02
struct bandinfo{
         unsigned char
                           bi_pri;
         int
                           bi_flag;
};
```

Figure 6-44: <termios.h>

```
#define NCCS
                          23
#define CTRL(c)
                         ((c)&037)
#define IBSHIFT
                         16
#define _POSIX_VDISABLE
typedef unsigned long tcflag_t;
typedef unsigned char
                        cc_t;
typedef unsigned long
                         speed_t;
#define VINTR
#define VQUIT
                         1
#define VERASE
                         2
#define VKILL
                          3
#define VEOF
                          4
#define VEOL
                         5
#define VEOL2
                         6
#define VMIN
                         4
#define VTIME
                         5
#define VSWTCH
                         7
#define VSTART
                         8
#define VSTOP
                         9
#define VSUSP
                         10
#define VDSUSP
                         11
#define VREPRINT
                         12
#define VDISCARD
                         13
#define VWERASE
                         14
#define VLNEXT
                         15
```

Elements 16-22 of the C_CC array are undefined and reserved for future use.

Figure 6-44: <termios.h> (continued)

```
#define CNUL
                        0
#define CDEL
                       0377
#define CESC
                       '\\'
#define CINTR
                       0177
#define CQUIT
                       034
#define CERASE
                       ′#′
#define CKILL
                       '@'
#define CEOT
                       04
#define CEOL
#define CEOL2
#define CEOF
                       04
#define CSTART
                       021
#define CSTOP
                       023
#define CSWTCH
                       032
#define CNSWTCH
#define CSUSP
                       CTRL('z')
#define CDSUSP
                       CTRL('y')
                       CTRL('r')
#define CRPRNT
                       CTRL('o')
#define CFLUSH
                       CTRL('w')
#define CWERASE
#define CLNEXT
                       CTRL('v')
                       0000001
#define IGNBRK
                       0000002
#define BRKINT
#define IGNPAR
                       0000004
#define PARMRK
                       0000010
#define INPCK
                       0000020
#define ISTRIP
                       0000040
#define INLCR
                       0000100
#define IGNCR
                       0000200
#define ICRNL
                       0000400
#define IUCLC
                       0001000
#define IXON
                       0002000
#define IXANY
                       0004000
#define IXOFF
                       0010000
```

Figure 6-44: <termios.h> (continued)

#define	OPOST	0000001
#define	OLCUC	0000002
#define	ONLCR	0000004
#define	OCRNL	0000010
#define	ONOCR	0000020
#define	ONLRET	0000040
#define	OFILL	0000100
#define	OFDEL	0000200
#define	NLDLY	0000400
#define	NL0	0
#define	NL1	0000400
#define	CRDLY	0003000
#define	CR0	0
#define	CR1	0001000
#define	CR2	0002000
#define	CR3	0003000
#define	TABDLY	0014000
#define	TAB0	0
#define	TAB1	0004000
#define	TAB2	0010000
#define	TAB3	0014000
#define	BSDLY	0020000
#define	BS0	0
#define	BS1	0020000
#define	VTDLY	0040000
#define	VT0	0
#define	VT1	0040000
#define	FFDLY	0100000
#define	FF0	0
#define	FF1	0100000

Figure 6-44: <termios.h> (continued)

/ #define	CBAUD	0000017
#define	в0	0
#define	B50	000001
#define	B75	0000002
#define	B110	0000003
#define	B134	000004
#define	B150	0000005
#define	B200	000006
#define	B300	000007
#define	В600	0000010
#define	B1200	0000011
#define	B1800	0000012
#define	B2400	0000013
#define	B4800	0000014
#define	В9600	0000015
#define	B19200	0000016
#define	EXTA	0000016
#define	B38400	0000017
#define	EXTB	0000017
#define	CSIZE	0000060
#define	CS5	0
#define	CS6	0000020
#define	CS7	0000040
#define	CS8	0000060
#define	CSTOPB	0000100
#define	CREAD	0000200
#define	PARENB	0000400
#define	PARODD	0001000
#define	HUPCL	0002000
#define	CLOCAL	0004000

Figure 6-44: <termios.h> (continued)

			_
#define	ISIG	0000001	
#define	ICANON	0000002	
#define	XCASE	000004	
#define	ECHO	0000010	
#define	ECHOE	0000020	
#define	ECHOK	0000040	
#define	ECHONL	0000100	
#define	NOFLSH	0000200	
#define	TOSTOP	0100000	
#define	ECHOCTL	0001000	
#define	ECHOPRT	0002000	
#define	ECHOKE	0004000	
#define	FLUSHO	0020000	
#define	PENDIN	0040000	
#define	IEXTEN	0000400	
#define	TIOC	('T'<<8)	
#define	TCSANOW	(TIOC 14)	
#define	TCSADRAIN	(TIOC 15)	
#define	TCSAFLUSH	(TIOC 16)	
#define	TCIFLUSH	0	
#define	TCOFLUSH	1	
#define	TCIOFLUSH	2	
#define	TCOOFF	0	
#define	TCOON	1	
#define	TCIOFF	2	
#define	TCION	3	

Figure 6-44: <termios.h> (continued)

Figure 6-45: <sys/ticlts.h>

```
#define TCL_BADADDR 1
#define TCL_BADOPT 2
#define TCL_NOPEER 3
#define TCL_PEERBADSTATE 4

#define TCL_DEFAULTADDRSZ 4
```

Figure 6-46: <sys/ticots.h>

```
#define TCO_NOPEER ECONNREFUSED
#define TCO_PEERNOROOMONQ ECONNREFUSED
#define TCO_PEERBADSTATE ECONNREFUSED
#define TCO_PEERINITIATED ECONNRESET
#define TCO_PROVIDERINITIATED ECONNABORTED

#define TCO_DEFAULTADDRSZ 4
```

Figure 6-47: <sys/ticotsord.h>

	#define TCOO_NOPEER	1	
	#define TCOO_PEERNOROOMONQ	2	
	#define TCOO_PEERBADSTATE	3	
	#define TCOO_PEERINITIATED	4	
	#define TCOO_PROVIDERINITIATED	5	
	#define TCOO_DEFAULTADDRSZ	4)
_			

Figure 6-48: <sys/time.h>

```
#define CLK_TCK
#define CLOCKS_PER_SEC
                           1000000
#define NULL
typedef long clock_t;
typedef long time_t;
struct tm{
      int tm_sec;
      int tm_min;
      int tm_hour;
      int tm_mday;
      int tm_mon;
      int tm_year;
      int tm_wday;
      int tm_yday;
      int tm_isdst;
};
struct timeval{
      time_t
                  tv_sec;
       long
                  tv_usec;
};
extern long timezone;
extern int daylight;
extern char *tzname[2];
typedef struct timestruc{
      time_t
                tv_sec;
      long
                  tv_nsec;
} timestruc_t;
/* starred values may vary and should be
      retrieved with sysconf() of pathconf() */
```

Figure 6-49: <sys/times.h>

```
struct tms{
    clock_t         tms_utime;
    clock_t         tms_stime;
    clock_t         tms_cutime;
    clock_t         tms_cstime;
};
```

Figure 6-50: <sys/tiuser.h>, Service Types

```
#define T_CLTS 3
#define T_COTS 1
#define T_COTS_ORD 2
```

Figure 6-51: <sys/tiuser.h>, Transport Interface States

```
#define T_DATAXFER
                        5
#define T_IDLE
                        2
#define T_INCON
                        4
#define T_INREL
                        7
#define T_OUTCON
                        3
                        6
#define T_OUTREL
                        1
#define T_UNBND
#define T_UNINIT
                        0
```

Figure 6-52: <sys/tiuser.h>, User-level Events

```
#define T_ACCEPT1
                        12
#define T_ACCEPT2
                         13
#define T_ACCEPT3
                        14
#define T_BIND
                         1
#define T_CLOSE
                         4
#define T_CONNECT1
                         8
#define T_CONNECT2
                         9
#define T_LISTN
                         11
#define T_OPEN
                         0
#define T_OPTMGMT
                         2
#define T_PASSCON
                         24
#define T_RCV
                         16
#define T_RCVCONNECT
                         10
#define T_RCVDIS1
                         19
#define T_RCVDIS2
                         20
#define T_RCVDIS3
                         21
#define T_RCVREL
                         23
#define T_RCVUDATA
                         6
#define T_RCVUDERR
                         7
#define T_SND
                         15
#define T_SNDDIS1
                         17
#define T_SNDDIS2
                         18
#define T_SNDREL
                         22
#define T_SNDUDATA
                         5
#define T_UNBIND
                         3
```

Figure 6-53: <sys/tiuser.h>, Error Return Values

#d	lefine	TACCES	3	\
#0	lefine	TBADADDR	1	
#0	lefine	TBADDATA	10	
#0	lefine	TBADF	4	
#0	lefine	TBADFLAG	16	
#0	lefine	TBADOPT	2	
#d	lefine	TBADSEQ	7	
#d	lefine	TBUFOVFLW	11	
#d	lefine	TFLOW	12	
#d	lefine	TLOOK	9	
#d	lefine	TNOADDR	5	
#d	lefine	TNODATA	13	
#6	lefine	TNODIS	14	
#d	lefine	TNOREL	17	
#6	lefine	TNOTSUPPORT	18	
#d	lefine	TNOUDERR	15	
#6	lefine	TOUTSTATE	6	
#6	lefine	TSTATECHNG	19	
, #d	lefine	TSYSERR	8	
				Ϊ

Figure 6-54: <sys/tiuser.h>, Transport Interface Data Structures

```
struct netbuf{
         unsigned int
                            maxlen;
         unsigned int
                             len;
         char
                             *buf;
};
struct t_bind{
         struct netbuf
                             addr;
         unsigned int
                            qlen;
};
 struct t_call{
         struct netbuf
                             addr;
         struct netbuf
                             opt;
         struct netbuf
                            udata;
         int
                             sequence;
};
 \verb|struct t_discon|| \\
         struct netbuf
                            udata;
         int
                             reason;
         int
                             sequence;
};
```

Figure 6-54: <sys/tiuser.h>, Transport Interface Data Structures (continued)

```
struct t_info {
         long
                    addr;
                    options;
         long
         long
                    tsdu;
         long
                    etsdu;
         long
                    connect;
         long
                    discon;
         long
                    servtype;
};
struct t_optmgmt{
         struct netbuf
                           opt;
         long
                           flags;
};
 struct t_uderr{
         struct netbuf
                           addr;
         struct netbuf
                           opt;
         long
                           error;
};
 struct t_unitdata{
         struct netbuf
                           addr;
         struct netbuf
                           opt;
                           udata;
         struct netbuf
};
```

Figure 6-55: <sys/tiuser.h>, Structure Types

```
#define T_BIND 1
#define T_CALL 3
#define T_DIS 4
#define T_INFO 7
#define T_OPTMGMT 2
#define T_UDERROR 6
#define T_UNITDATA 5
```

Figure 6-56: <sys/tiuser.h>, Fields of Structures

Figure 6-57: <sys/tiuser.h>, Events Bitmasks

```
#define T_LISTEN
                         0x01
#define T_CONNECT
                         0x02
#define T_DATA
                         0x04
#define T_EXDATA
                         0x08
#define T_DISCONNECT
                         0x10
#define T_ERROR
                         0x20
#define T_UDERR
                         0x40
#define T_ORDREL
                         0 \times 80
#define T_EVENTS
                         0xff
```

Figure 6-58: <sys/tiuser.h>, Flags

```
#define T_MORE 0x01
#define T_EXPEDITED 0x02
#define T_NEGOTIATE 0x04
#define T_CHECK 0x08
#define T_DEFAULT 0x10
#define T_SUCCESS 0x20
#define T_FAILURE 0x40
```

Figure 6-59: <sys/types.h>

```
typedef long
                        time t;
typedef long
                        daddr_t;
typedef unsigned long
                       dev_t;
typedef long
                       gid_t;
typedef unsigned long
                      ino_t;
typedef int
                       key_t;
typedef long
                        pid_t;
typedef unsigned long
                       mode_t;
typedef unsigned long
                       nlink_t;
typedef long
                        off_t;
typedef long
                        uid_t;
typedef long
                        clock_t
typedef unsigned int
                        size_t
```

Figure 6-60: <sys/ucontext.h>

```
typedef unsigned int greg_t;
#define NGREG 36
typedef greg_t gregset_t[NGREG];
typedef struct fpregset {
         union {
         double
                       fp_dregs[16];
         float
                       fp freqs [32];
         unsigned int fp_regs[32];
         } fp_r;
         unsigned int fp_csr;
         unsigned int fp_pad;
} fpregset_t;
 typedef struct {
         gregset_t gregs;
         fpregset_t fpregs;
} mcontext_t;
 typedef struct ucontext{
         unsigned long
                            uc flags;
                             *uc_link;
         struct ucontext
         sigset_t
                       uc_sigmask;
         stack_t
                       uc_stack;
         mcontext_t
                       uc_mcontext;
                       uc_filler[48];
         long
} ucontext_t;
```

The size of the ucontext sruct is 128 words according to the alignment rules in Chapter 3. Specifically, the fpregset struct is double word aligned, forcing the mcontext_t and ucontext structures to also be double word aligned.

Figure 6-60: <sys/ucontext.h> (continued)

#define CXT R0	0	
#define CXT AT	1	
#define CXT V0	2	
#define CXT_V1	3	
#define CXT A0	4	
#define CXT A1	5	
#define CXT A2	6	
#define CXT A3	7	
#define CXT T0	8	
#define CXT T1	9	
#define CXT T2	10	
#define CXT T3	11	
#define CXT T4	12	
#define CXT T5	13	
#define CXT T6	14	
#define CXT_T7	15	
#define CXT_S0	16	
#define CXT_S1	17	
#define CXT_S2	18	
#define CXT_S3	19	
#define CXT_S4	20	
#define CXT_S5	21	
#define CXT_S6	22	
#define CXT_S7	23	
#define CXT_T8	24	
#define CXT_T9	25	
#define CXT_K0	26	
#define CXT_K1	27	
#define CXT_GP	28	
#define CXT_SP	29	

Figure 6-60: <sys/ucontext.h> (continued)

```
#define CXT_S8 30
#define CXT_RA 31
#define CXT_MDLO 32
#define CXT_MDHI 33
#define CXT_CAUSE 34
#define CXT_EPC 35
```

Figure 6-61: <sys/uio.h>

```
typedef struct iovec{
         char *iov_base;
         int iov_len;
} iovec_t;
```

Figure 6-62: <ulimit.h>

```
#define UL_GETFSIZE 1
#define UL_SETFSIZE 2
```

Figure 6-63: <unistd.h>

```
#define R_OK
#define W_OK
                              2
#define X_OK
                              1
                              0
#define F_OK
#define F_ULOCK
#define F_LOCK
                              1
#define F_TLOCK
#define F_TEST
#define SEEK_SET
#define SEEK_CUR
                              1
#define SEEK_END
#define _POSIX_JOB_CONTROL
#define _POSIX_SAVED_IDS
                              1
#define _POSIX_VDISABLE
#define _POSIX_VERSION
#define _XOPEN_VERSION
/* starred values vary and should be
  retrieved using sysconf() or pathconf() */
```

Figure 6-63: <unistd.h> (continued)

#define	SC ARG MAX	1
#define	SC CHILD MAX	2
#define	SC CLK TCK	3
#define	 _SC_NGROUPS_MAX	4
#define	_SC_OPEN_MAX	5
#define	_SC_JOB_CONTROL	6
#define	_SC_SAVED_IDS	7
#define	_SC_VERSION	8
#define	_SC_PASS_MAX	9
#define	_SC_PAGESIZE	11
#define	_SC_XOPEN_VERSION	12
#define	PC LINK MAX	1
#define	PC MAX CANON	2
#define	PC MAX INPUT	3
#define	 _PC_NAME_MAX	4
#define	PC PATH MAX	5
#define	PC PIPE BUF	6
#define	 _PC_CHOWN_RESTRICTED	7
#define	_PC_NO_TRUNC	8
	_PC_VDISABLE	9
#define	STDIN_FILENO	0
#define	STDOUT_FILENO	1
#define	STDERR FILENO	2

Figure 6-64: <utime.h>

```
struct utimbuf{
    time_t actime;
    time_t modtime;
};
```

Figure 6-65: <sys/utsname.h>

```
#define SYS_NMLN
                  257
struct utsname{
      char
                  sysname[SYS_NMLN];
      char
                  nodename[SYS_NMLN];
      char
                  release[SYS_NMLN];
      char
                  version[SYS_NMLN];
                  machine[SYS_NMLN];
      char
      char
                  m_type[SYS_NMLN];
      char
                  base_rel[SYS_NMLN];
      char
                  reserve5[SYS_NMLN];
      char
                  reserve4[SYS_NMLN];
      char
                  reserve3[SYS_NMLN];
      char
                  reserve2[SYS_NMLN];
      char
                  reserve1[SYS_NMLN];
      char
                  reserve0[SYS_NMLN];
};
```

The fields m_{type} , base_rel, reserve5, reserve4, reserve3, reserve2, reserve1, and reserve0 are not defined in the SVID and are reserved for future use.

Figure 6-66: <wait.h>

```
#define WEXITED
                       0001
#define WTRAPPED
                       0002
#define WSTOPPED
                       0004
#define WCONTINUED
                       0010
#define WUNTRACED
                       WSTOPPED
#define WNOHANG
                       0100
#define WNOWAIT
                       0200
#define WSTOPFLG
                       0177
#define WCONTFLG
                       0177777
#define WCOREFLG
                       0200
#define WSIGMASK
                       0177
#define WWORD(stat)
                          ((int)((stat))&0177777)
#define WIFEXITED(stat)
                           ((int) ((stat) & 0377) = = 0)
#define WIFSIGNALED(stat)\
  (((int)((stat)\&0377)>0)\&\&(((int)(((stat)>>8)\&0377))==0))
#define WIFSTOPPED(stat)\
  (((int)((stat)\&0377)==WSTOPFLAG)\&\&(((int)(((stat)>>8))
  &0377))!=0))
#define WIFCONTINUED(stat) (WWORD(stat)==WCONTFLG)
#define WEXITSTATUS(stat)
                            (((int)(((stat>>8)&0377))
#define WTERMSIG(stat)
                            (((int)((stat)&0377)&WSIGMASK))
#define WSTOPSIG(stat)
                            ((int)(((stat)>>8)&0377))
#define WCOREDUMP(stat)
                            ((stat)&WCOREFLG)
```

Figure 6-67: <varargs.h>

X Window Data Definitions

This section is new, but will not be diffmarked.

This section contains standard data definitions that describe system data for the optional X Window windowing libraries. These data definitions are referred to by their names in angle brackets: <name.h> and <sys/name.h>. Included in these data definitions are macro definitions and structure definitions. While an ABI-conforming system may provide X11 and X Toolkit Intrinsics interfaces, it need not contain the actual data definitions referenced here. Programmers should observe that the sources of the structures defined in these data definitions are defined in SVID or the appropriate X Consortium documentation (see chapter 10 in the Generic ABI).

Figure 6-1: <X11/Atom.h>

	#define	XA_PRIMARY	((Atom) 1)
	#define	XA_SECONDARY	((Atom) 2)
	#define	XA_ARC	((Atom) 3)
	#define	XA_ATOM	((Atom) 4)
	#define	XA_BITMAP	((Atom) 5)
	#define	XA_CARDINAL	((Atom) 6)
	#define	XA_COLORMAP	((Atom) 7)
	#define	XA_CURSOR	((Atom) 8)
	#define	XA_CUT_BUFFER0	((Atom) 9)
	#define	XA_CUT_BUFFER1	((Atom) 10)
	#define	XA_CUT_BUFFER2	((Atom) 11)
	#define	XA_CUT_BUFFER3	((Atom) 12)
	#define	XA_CUT_BUFFER4	((Atom) 13)
	#define	XA_CUT_BUFFER5	((Atom) 14)
	#define	XA_CUT_BUFFER6	((Atom) 15)
	#define	XA_CUT_BUFFER7	((Atom) 16)
	#define	XA_DRAWABLE	((Atom) 17)
	#define	XA_FONT	((Atom) 18)
	#define	XA_INTEGER	((Atom) 19)
	#define	XA_PIXMAP	((Atom) 20)
	#define	XA_POINT	((Atom) 21)
	#define	XA_RECTANGLE	((Atom) 22)
	#define	XA_RESOURCE_MANAGER	((Atom) 23)
	#define	XA_RGB_COLOR_MAP	((Atom) 24)
	#define	XA_RGB_BEST_MAP	((Atom) 25)
	#define	XA_RGB_BLUE_MAP	((Atom) 26)
	#define	XA_RGB_DEFAULT_MAP	((Atom) 27)
	#define	XA_RGB_GRAY_MAP	((Atom) 28)
	#define	XA_RGB_GREEN_MAP	((Atom) 29)
	#define	XA_RGB_RED_MAP	((Atom) 30)
	• •	XA_STRING	((Atom) 31)
	#define	XA_VISUALID	((Atom) 32)
_			

Figure 6-1: <X11/Atom.h> (continued)

/	• •	XA_WINDOW	((Atom)	•
	**	XA_WM_COMMAND	((Atom)	34)
	#define	XA_WM_HINTS	((Atom)	35)
	#define	XA_WM_CLIENT_MACHINE	((Atom)	36)
		XA_WM_ICON_NAME	((Atom)	37)
	#define	XA_WM_ICON_SIZE	((Atom)	38)
	#define	XA_WM_NAME	((Atom)	39)
	#define	XA_WM_NORMAL_HINTS	((Atom)	40)
	#define	XA_WM_SIZE_HINTS	((Atom)	41)
	#define	XA_WM_ZOOM_HINTS	((Atom)	42)
	#define	XA_MIN_SPACE	((Atom)	43)
	#define	XA_NORM_SPACE	((Atom)	44)
	#define	XA_MAX_SPACE	((Atom)	45)
	#define	XA_END_SPACE	((Atom)	46)
	#define	XA_SUPERSCRIPT_X	((Atom)	47)
	#define	XA_SUPERSCRIPT_Y	((Atom)	48)
	#define	XA_SUBSCRIPT_X	((Atom)	49)
	#define	XA_SUBSCRIPT_Y	((Atom)	50)
	#define	XA_UNDERLINE_POSITION	((Atom)	51)
	#define	XA_UNDERLINE_THICKNESS	((Atom)	52)
	#define	XA_STRIKEOUT_ASCENT	((Atom)	53)
	#define	XA_STRIKEOUT_DESCENT	((Atom)	54)
	#define	XA_ITALIC_ANGLE	((Atom)	55)
	#define	XA_X_HEIGHT	((Atom)	56)
	#define	XA_QUAD_WIDTH	((Atom)	57)
	#define	XA_WEIGHT	((Atom)	58)
	#define	XA_POINT_SIZE	((Atom)	59)
	#define	XA_RESOLUTION	((Atom)	60)
	#define	XA_COPYRIGHT	((Atom)	61)
	#define	XA_NOTICE	((Atom)	62)
	#define	XA_FONT_NAME	((Atom)	63)
	#define	XA_FAMILY_NAME	((Atom)	64)

Figure 6-1: <X11/Atom.h> (continued)

```
#define XA_FULL_NAME ((Atom) 65)
#define XA_CAP_HEIGHT ((Atom) 66)
#define XA_WM_CLASS ((Atom) 67)
#define XA_WM_TRANSIENT_FOR ((Atom) 68)
#define XA_LAST_PREDEFINED ((Atom) 68)
```

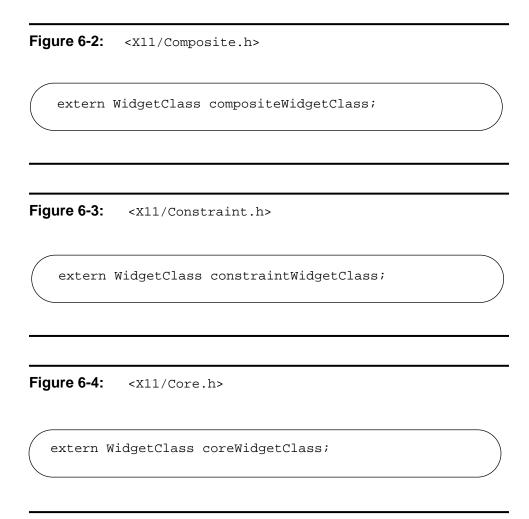


Figure 6-5: <X11/cursorfont.h>

/ #defi	ne XC_num_glyphs	154
#defi	ne XC_X_cursor	0
#defi	ne XC_arrow	2
#defi	ne XC_based_arrow_down	4
#defi	ne XC_based_arrow_up	6
#defi	ne XC_boat	8
#defi	ne XC_bogosity	10
#defi	ne XC_bottom_left_corner	12
#defi	ne XC_bottom_right_corner	14
#defi	ne XC_bottom_side	16
#defi	ne XC_bottom_tee	18
#defi	ne XC_box_spiral	20
#defi	ne XC_center_ptr	22
#defi	ne XC_circle	24
#defi	ne XC_clock	26
#defi	ne XC_coffee_mug	28
#defi	ne XC_cross	30
#defi	ne XC_cross_reverse	32
#defi	ne XC_crosshair	34
#defi	ne XC_diamond_cross	36
#defi	ne XC_dot	38
#defi	ne XC_dotbox	40
#defi	ne XC_double_arrow	42
#defi	ne XC_draft_large	44
#defi	ne XC_draft_small	46
#defi	ne XC_draped_box	48
#defi	ne XC_exchange	50
#defi	ne XC_fleur	52
#defi	ne XC_gobbler	54
#defi	ne XC_gumby	56
#defi	ne XC_hand1	58
#defi	ne XC_hand2	60

Figure 6-5: <X11/cursorfont.h> (continued)

	XC_heart	62
	XC_icon	64
	XC_iron_cross	66
	XC_left_ptr	68
#define	XC_left_side	70
#define	XC_left_tee	72
#define	XC_leftbutton	74
#define	XC_ll_angle	76
#define	XC_lr_angle	78
#define	XC_man	80
#define	XC_middlebutton	82
#define	XC_mouse	84
#define	XC_pencil	86
#define	XC_pirate	88
#define	XC_plus	90
#define	XC_question_arrow	92
#define	XC_right_ptr	94
#define	XC_right_side	96
#define	XC_right_tee	98
#define	XC_rightbutton	100
#define	XC_rtl_logo	102
#define	XC_sailboat	104
#define	XC_sb_down_arrow	106
#define	XC_sb_h_double_arrow	108
#define	XC_sb_left_arrow	110
#define	XC_sb_right_arrow	112
#define	XC_sb_up_arrow	114
#define	XC_sb_v_double_arrow	116
#define	XC_shuttle	118
#define	XC_sizing	120
#define	XC_spider	122
#define	XC_spraycan	124

Figure 6-5: <X11/cursorfont.h> (continued)

#define	XC_star	126
#define	XC_target	128
#define	XC_tcross	130
#define	<pre>XC_top_left_arrow</pre>	132
#define	XC_top_left_corner	134
#define	XC_top_right_corner	136
#define	XC_top_side	138
#define	XC_top_tee	140
#define	XC_trek	142
#define	XC_ul_angle	144
#define	XC_umbrella	146
#define	XC_ur_angle	148
#define	XC_watch	150
#define	XC_xterm	152

Figure 6-6: <X11/Intrinsic.h>

```
typedef char
                            *String;
#define XtNumber(arr)\
       ((Cardinal) (sizeof(arr) / sizeof(arr[0])))
typedef void
                           Widget;
typedef Widget
                           *WidgetList;
typedef void
                           CompositeWidget;
                          XtActionList;
typedef XtActionsRec
typedef void
                          XtAppContext;
typedef unsigned long
                          XtValueMask;
typedef unsigned long
                          XtIntervalId;
typedef unsigned long
                          XtInputId;
typedef unsigned long
                          XtWorkProcId;
typedef unsigned int
                          XtGeometryMask;
typedef unsigned long
                          XtGCMask;
typedef unsigned long
                          Pixel;
typedef int
                          XtCacheType;
#define
             XtCacheNone 0x001
#define
            XtCacheAll
                          0x002
#define
            XtCacheByDisplay
                                    0x003
#define
            XtCacheRefCount
                                    0x100
typedef char
                          Boolean;
typedef long
                          XtArqVal;
typedef unsigned char
                          XtEnum;
typedef unsigned int
                           Cardinal;
typedef unsigned short
                          Dimension;
typedef short
                          Position;
typedef char
                           *XtPointer;
```

Figure 6-6: <X11/Intrinsic.h> (continued)

```
typedef void
                         XtTranslations;
typedef void
                        XtAccelerators;
typedef unsigned int
                        Modifiers;
#define XtCWQueryOnly
                      (1 << 7)
#define XtSMDontChange
typedef void
                         XtCacheRef;
typedef void
                        XtActionHookId;
typedef unsigned long EventMask;
typedef enum {XtListHead, XtListTail } XtListPosition;
typedef struct {
        String
                         string;
        XtActionProc
                        proc;
} XtActionsRec;
typedef enum {
  XtAddress,
  XtBaseOffset,
  XtImmediate,
  XtResourceString,
  XtResourceQuark,
  XtWidgetBaseOffset,
  XtProcedureArg
} XtAddressMode;
typedef struct {
  XtAddressMode
                         address_mode;
  XtPointer
                        address_id;
  Cardinal
                        size;
} XtConvertArgRec, *XtConvertArgList;
```

Figure 6-6: <X11/Intrinsic.h> (continued)

```
#define XtInputNoneMask
                                 0L
#define XtInputReadMask
                                1L<<0)
#define XtInputWriteMask
                                (1L << 1)
#define XtInputExceptMask
                                (1L << 2)
typedef struct {
     XtGeometryMask
                                request_mode;
     Position x, y;
     Dimension width, height, border_width;
     Widget
                 sibling;
} XtWidgetGeometry;
typedef struct {
     String
                  name;
     XtArqVal
                  value;
} Arg, *ArgList;
typedef XtPointer XtVarArgsList;
typedef struct {
     XtCallbackProc
                              callback;
      XtPointer
                              closure;
} XtCallbackRec, *XtCallbackList;
typedef enum {
     XtCallbackNoList,
     XtCallbackHasNone,
     XtCallbackHasSome
} XtCallbackStatus;
typedef struct {
     Widget
                  shell_widget;
     Widget
                 enable widget;
} XtPopdownIDRec, *XtPopdownID;
```

Figure 6-6: <X11/Intrinsic.h> (continued)

```
typedef enum {
      XtGeometryYes,
      XtGeometryNo,
      XtGeometryAlmost,
      XtGeometryDone
} XtGeometryResult;
typedef enum {
      XtGrabNone,
      XtGrabNonexclusive,
      XtGrabExclusive
} XtGrabKind;
typedef struct {
      String resource_name;
String resource_class;
String resource_type;
      Cardinal resource_size;
      Cardinal resource_offset;
      String
                  default_type;
      XtPointer
                  default_addr;
} XtResource, *XtResourceList;
typedef struct {
      char
                         match;
      String
                         substitution;
} SubstitutionRec,
                        *Substitution;
                         (*XtFilePredicate);
typedef Boolean
typedef XtPointer
                         XtRequestId;
extern XtConvertArgRec const colorConvertArgs[];
extern XtConvertArgRec const screenConvertArg[];
```

Figure 6-6: <X11/Intrinsic.h> (continued)

```
#define XtAllEvents
                                 ((EventMask) -1L)
#define XtIMXEvent
#define XtIMTimer
                                 2
#define XtIMAlternateInput
#define XtIMAll (XtIMXEvent | XtIMTimer | XtIMAlternateInput)
#define XtOffsetOf(s_type,field) XtOffset(s_type*,field)
#define XtNew(type) ((type *) XtMalloc((unsigned sizeof(type))
#define XT_CONVERT_FAIL
                                 (Atom)0x8000001
#define XtIsRectObj(object) \
(_XtCheckSubclassFlag(object,(XtEnum)0x02))
#define XtIsWidget(object) \
(_XtCheckSubclassFlag(object,(XtEnum)0x04))
#define XtIsComposite(widget) \
(_XtCheckSubclassFlag(widget,(XtEnum)0x08))
#define XtIsConstraint(widget) \
(_XtCheckSubclassFlag(widget,(XtEnum)0x10))
#define XtIsShell(widget) \
 (_XtCheckSubclassFlag(widget,(XtEnum)0x20))
#define XtIsOverrideShell(widget) \
(_XtIsSubclassOf(widget,(Widge Class)overrideShellWidgetClass,\
 (WidgetClass)shellWidgetClass, (XtEnum)0x20))
#define XtIsWMShell(widget) \
(_XtCheckSubclassFlag(widget,(XtEnum)0x40))
#define XtIsVendorShell(widget)\
 (_XtIsSubclassOf(widget,(WidgetClass)vendorShellWidgetClass,
\#define XtIsTopLevelShell(widget)\
(_XtCheckSubclassFlag(widget, (XtEnum)0x80))
#define XtIsApplicationShell(widget)\
(_XtIsSubclassOf(widget,(WidgetClass)appliationShellWidgetClass)
 (WidgetClass)topLevelShellWidgetClass, (XtEum)0x80))
```

Figure 6-6: <X11/Intrinsic.h> (continued)

```
#define XtSetArg(arg,n,d)\
      ((void)((arg).name = (n), (arg).value =
(XtArgVal)(d) ))
#define XtOffset(p_type,field)\
      ((Cardinal) (((char *) (&(((p_type)NULL)-
>field)))\
       - ((char *) NULL)))
#define XtVaNestedList
                                 "XtVaNestedList"
#define XtVaTypedArg
                                 "XtVaTypedArg"
#define XtUnspecifiedPixmap
                                 ((Pixmap)2)
#define XtUnspecifiedShellInt
                                 (-1)
#define XtUnspecifiedWindow
                                ((Window)2)
#define XtUnspecifiedWindowGroup ((Window)3)
#define XtDefaultForeground
                                "XtDefaultForeground"
#define XtDefaultBackground
                                 "XtDefaultBackground"
#define XtDefaultFont
                                 "XtDefaultFont"
#define XtDefaultFontSet
                                 "XtDefaultFontSet"
```

Figure 6-7: <X11/Object.h>

```
extern WidgetClass objectClass;
```

Figure 6-8: <X11/RectObj.h>

extern WidgetClass rectObjClass;

Figure 6-9: <X11/Shell.h>

```
extern WidgetClass shellWidgetClass;
extern WidgetClass overrideShellWidgetClass;
extern WidgetClass wmShellWidgetClass;
extern WidgetClass transientShellWidgetClass;
extern WidgetClass topLevelShellWidgetClass;
extern WidgetClass applicationShellWidgetClass;
```

Figure 6-10: <X11/Vendor.h>

extern WidgetClass vendorShellWidgetClass;

Figure 6-11: <x11/x.h>

```
typedef unsigned long XID;
typedef XID Window;
typedef XID Drawable;
typedef XID Font;
typedef XID Pixmap;
typedef XID Cursor;
typedef XID Colormap;
typedef XID GContext;
typedef XID KeySym;
typedef unsigned long Atom;
typedef unsigned long VisualID;
typedef unsigned long Time;
typedef unsigned char KeyCode;
#define AllTemporary
                           0L
#define AnyButton
                           0L
#define AnyKey
                           0L
#define AnyPropertyType
                           0L
#define CopyFromParent
                           0L
#define CurrentTime
                           0Ъ
#define InputFocus
                           1L
#define NoEventMask
                           0L
#define None
                           0Ъ
#define NoSymbol
                           0L
#define ParentRelative
                           1L
#define PointerWindow
                           0L
#define PointerRoot
                           1L
```

Figure 6-11: <X11/X.h> (continued)

```
#define KeyPressMask
                                    (1L << 0)
#define KeyReleaseMask
                                    (1L << 1)
#define ButtonPressMask
                                    (1L << 2)
#define ButtonReleaseMask
                                    (1L << 3)
#define EnterWindowMask
                                    (1L << 4)
#define LeaveWindowMask
                                    (1L < < 5)
#define PointerMotionMask
                                    (1L < < 6)
#define PointerMotionHintMask
                                    (1L << 7)
#define Button1MotionMask
                                    (1L << 8)
#define Button2MotionMask
                                    (1L << 9)
#define Button3MotionMask
                                    (1L << 10)
#define Button4MotionMask
                                    (1L << 11)
#define Button5MotionMask
                                    (1L << 12)
#define ButtonMotionMask
                                    (1L << 13)
#define KeymapStateMask
                                    (1L << 14)
#define ExposureMask
                                    (1L << 15)
#define VisibilityChangeMask
                                    (1L << 16)
#define StructureNotifyMask
                                    (1L << 17)
#define ResizeRedirectMask
                                    (1L << 18)
#define SubstructureNotifyMask
                                    (1L << 19)
#define SubstructureRedirectMask (1L<<20)</pre>
#define FocusChangeMask
                                    (1L << 21)
#define PropertyChangeMask
                                    (1L << 22)
#define ColormapChangeMask
                                    (1L << 23)
#define OwnerGrabButtonMask
                                    (1L << 24)
```

Figure 6-11: <X11/X.h> (continued)

#define	KeyPress	2	
	KeyRelease	3	
#define	ButtonPress	4	
#define	ButtonRelease	5	
#define	MotionNotify	6	
#define	EnterNotify	7	
#define	LeaveNotify	8	
#define	FocusIn	9	
#define	FocusOut	10	
#define	KeymapNotify	11	
#define	Expose	12	
#define	GraphicsExpose	13	
#define	NoExpose	14	
#define	VisibilityNotify	15	
#define	CreateNotify	16	
#define	DestroyNotify	17	
#define	UnmapNotify	18	
#define	MapNotify	19	
#define	MapRequest	20	
#define	ReparentNotify	21	
#define	ConfigureNotify	22	
#define	ConfigureRequest	23	
#define	GravityNotify	24	
#define	ResizeRequest	25	
#define	CirculateNotify	26	
#define	CirculateRequest	27	
#define	PropertyNotify	28	
#define	SelectionClear	29	
#define	SelectionRequest	30	
#define	SelectionNotify	31	
#define	ColormapNotify	32	
#define	ClientMessage	33	
#define	MappingNotify	34	,

Figure 6-11: <X11/X.h> (continued)

	#define	ShiftMask	(1<<0)
	#define	LockMask	(1<<1)
	#define	ControlMask	(1<<2)
	#define	Mod1Mask	(1<<3)
	#define	Mod2Mask	(1<<4)
	#define	Mod3Mask	(1<<5)
	#define	Mod4Mask	(1<<6)
	#define	Mod5Mask	(1<<7)
			(1 0)
		Button1Mask	(1<<8)
		Button2Mask	(1<<9)
		Button3Mask	(1<<10)
		Button4Mask	(1<<11)
		Button5Mask	(1<<12)
	#define	AnyModifier	(1<<15)
	#define	Button1	1
		Button2	2
		Button3	3
		Button4	4
		Button5	5
	WOL 1110	2400010	
	#define	NotifyNormal	0
	#define	NotifyGrab	1
	#define	NotifyUngrab	2
	#define	NotifyWhileGrabbed	3
	#define	NotifyHint	1
	#define	NotifyAncestor	0
	#define	NotifyVirtual	1
	#define	NotifyInferior	2
	#define	NotifyNonlinear	3
	#define	NotifyNonlinearVirtual	4
		NotifyPointer	5
	#define	NotifyPointerRoot	6
	#define	NotifyDetailNone	7
\			

Figure 6-11: <X11/X.h> (continued)

#define	VisibilityUnobscured	0
#define	VisibilityPartiallyObscured	1
#define	VisibilityFullyObscured	2
u 1 . C.'	7.1	0
	PlaceOnTop	0
#define	PlaceOnBottom	1
#define	PropertyNewValue	0
#define	PropertyDelete	1
#define	ColormapUninstalled	0
#define	ColormapInstalled	1
#define	GrabModeSync	0
	GrabModeAsync	1
παCIIIIC	Ol abriode Abylic	1
#define	GrabSuccess	0
#define	AlreadyGrabbed	1
#define	GrabInvalidTime	2
#define	GrabNotViewable	3
#define	GrabFrozen	4
#define	AsyncPointer	0
	SyncPointer	1
	ReplayPointer	2
	AsyncKeyboard	3
	SyncKeyboard	4
	ReplayKeyboard	5
	AsyncBoth	6
	SyncBoth	7
HACTING		,
	RevertToNone	(int)None
#define	RevertToPointerRoot	(int)PointerRoot
#define	RevertToParent	2

Figure 6-11: <X11/X.h> (continued)

#de	efine	Success	0
#de	efine	BadRequest	1
#de	efine	BadValue	2
#de	efine	BadWindow	3
#de	efine	BadPixmap	4
#de	efine	BadAtom	5
#de	efine	BadCursor	6
#de	efine	BadFont	7
#de	efine	BadMatch	8
#de	efine	BadDrawable	9
#de	efine	BadAccess	10
#de	efine	BadAlloc	11
#de	efine	BadColor	12
#de	efine	BadGC	13
#de	efine	BadIDChoice	14
#de	efine	BadName	15
#de	efine	BadLength	16
#de	efine	BadImplementation	17
#de	efine	InputOutput	1
#de	efine	InputOnly	2
#de	efine	CWBackPixmap	(1L<<0)
#de	efine	CWBackPixel	(1L<<1)
#de	efine	CWBorderPixmap	(1L<<2)
		CWBorderPixel	(1L<<3)
#de	efine	CWBitGravity	(1L<<4)
#de	efine	CWWinGravity	(1L<<5)
#de	efine	CWBackingStore	(1L<<6)
#de	efine	CWBackingPlanes	(1L<<7)
#de	efine	CWBackingPixel	(1L<<8)
		CWOverrideRedirect	(1L<<9)
#de	efine	CWSaveUnder	(1L<<10)
1		CWEventMask	(1L<<11)
#de	efine	CWDontPropagate	(1L<<12)
#de	efine	CWColormap	(1L<<13)
\ #de	efine	CWCursor	(1L<<14)

Figure 6-11: <X11/X.h> (continued)

_			
	#define	CWX	(1<<0)
	#define	CWY	(1<<1)
	#define	CWWidth	(1<<2)
	#define	CWHeight	(1<<3)
	#define	CWBorderWidth	(1<<4)
	#define	CWSibling	(1<<5)
	#define	CWStackMode	(1<<6)
		ForgetGravity	0
	#define	NorthWestGravity	1
		NorthGravity	2
	#define	NorthEastGravity	3
		WestGravity	4
	#define	CenterGravity	5
	#define	EastGravity	6
	#define	SouthWestGravity	7
	#define	SouthGravity	8
	#define	SouthEastGravity	9
	#define	StaticGravity	10
	#define	UnmapGravity	0
	#dofino	NotUseful	0
		WhenMapped	1
	#define		2
	#deline	Always	2
	#define	IsUnmapped	0
	#define	IsUnviewable	1
	#define	IsViewable	2
	#define	SetModeInsert	0
	#define	SetModeDelete	1
	#define	DestroyAll	0
		RetainPermanent	1
(#define	RetainTemporary	2
\			/

Figure 6-11: <X11/X.h> (continued)

#define	Above	0
#define		1
#define	-	2
#define	BottomIf	3
#define	Opposite	4
#define	RaiseLowest	0
#define	LowerHighest	1
#define	PropModeReplace	0
#define	PropModePrepend	1
#define	PropModeAppend	2
#define	GXclear	0x0
#define	GXand	0x1
#define	GXandReverse	0x2
#define	GXcopy	0x3
#define	GXandInverted	0x4
#define	GXnoop	0x5
#define	GXxor	0x6
#define	GXor	0x7
#define	GXnor	0x8
#define	GXequiv	0x9
#define	GXinvert	0xa
#define	GXorReverse	0xb
#define	GXcopyInverted	0xc
#define	GXorInverted	0xd
#define	GXnand	0xe
#define	GXset	0xf
#define	LineSolid	0
#define	LineOnOffDash	1
#define	LineDoubleDash	2
#define	CapNotLast	0
#define	CapButt	1
#define	CapRound	2
	CapProjecting	3

Figure 6-11: <X11/X.h> (continued)

#define	JoinMiter	0	
#define	JoinRound	1	
#define	JoinBevel	2	
#define	FillSolid	0	
#define	FillTiled	1	
#define	FillStippled	2	
#define	FillOpaqueStippled	3	
#define	EvenOddRule	0	
#define	WindingRule	1	
#define	ClipByChildren	0	
#define	IncludeInferiors	1	
	Unsorted	0	
#define	YSorted	1	
#define	YXSorted	2	
#define	YXBanded	3	
#define	CoordModeOrigin	0	
#define	CoordModePrevious	1	
#define	Complex	0	
#define	Nonconvex	1	
#define	Convex	2	
#define	ArcChord	0	
#define	ArcPieSlice	1	/

Figure 6-11: <X11/X.h> (continued)

#def:	ine GCFunction	(1L<<0)
#def:	ine GCPlaneMask	(1L<<1)
#defi	ine GCForeground	(1L<<2)
#def:	ine GCBackground	(1L<<3)
#def:	ine GCLineWidth	(1L<<4)
#def:	ine GCLineStyle	(1L<<5)
#def:	ine GCCapStyle	(1L<<6)
#def:	ine GCJoinStyle	(1L<<7)
#def:	ine GCFillStyle	(1L<<8)
#def:	ine GCFillRule	(1L<<9)
#def:	ine GCTile	(1L<<10)
#def:	ine GCStipple	(1L<<11)
#def:	ine GCTileStipXOrig	in (1L<<12)
#def:	ine GCTileStipYOrig	in (1L<<13)
#def:	ine GCFont	(1L<<14)
#def:	ine GCSubwindowMode	(1L<<15)
#def:	ine GCGraphicsExpos	ures (1L<<16)
#def:	ine GCClipXOrigin	(1L<<17)
#def:	ine GCClipYOrigin	(1L<<18)
#def:	ine GCClipMask	(1L<<19)
#def:	ine GCDashOffset	(1L<<20)
#def:	ine GCDashList	(1L<<21)
#def:	ine GCArcMode	(1L<<22)
#def:	ine FontLeftToRight	0
#def:	ine FontRightToLeft	1
#def:	ine XYBitmap	0
#defi	ine XYPixmap	1
#def:	ine ZPixmap	2
#def:	ine AllocNone	0
#def:	ine AllocAll	1
#def:	ine DoRed	(1<<0)
#def:	ine DoGreen	(1<<1)
\ #def	ine DoBlue	(1<<2)

Figure 6-11: <X11/X.h> (continued)

#define	CursorShape	0
#define	TileShape	1
#define	StippleShape	2
#define	AutoRepeatModeOff	0
	AutoRepeatModeOn	1
	AutoRepeatModeDefault	2
GG	11400116F 040110402 014410	_
#define	LedModeOff	0
#define	LedModeOn	1
#define	KBKeyClickPercent	(1L<<0)
	KBBellPercent	(1L<<1)
#define	KBBellPitch	(1L<<2)
#define	KBBellDuration	(1L<<3)
#define	KBLed	(1L<<4)
#define	KBLedMode	(1L<<5)
#define	KBKey	(1L<<6)
#define	KBAutoRepeatMode	(1L<<7)
#define	MappingSuccess	0
	MappingBusy	1
	MappingFailed	2
#GCI IIIC	mappingi dired	2
#define	MappingModifier	0
#define	MappingKeyboard	1
#define	MappingPointer	2
#define	DontPreferBlanking	0
#define	PreferBlanking	1
#define	DefaultBlanking	2
#define	DontAllowExposures	0
	AllowExposures	1
	DefaultExposures	2
	_	,

Figure 6-11: <X11/X.h> (continued)

#define	ScreenSaverReset	0
#define	ScreenSaverActive	1
#define	EnableAccess	1
#define	DisableAccess	0
#define	StaticGray	0
#define	GrayScale	1
#define	StaticColor	2
#define	PseudoColor	3
#define	TrueColor	4
#define	DirectColor	5
#define	LSBFirst	0
#define	MSBFirst	1

Figure 6-12: <X11/Xcms.h>

```
#define XcmsFailure
                                 0
#define XcmsSuccess
                                 1
#define XcmsSuccessWithCompression
#define XcmsUndefinedFormat
     (XcmsColorFormat)0x0000000
#define XcmsCIEXYZFormat
      (XcmsColorFormat)0x0000001
#define XcmsCIEuvYFormat
      (XcmsColorFormat)0x0000002
#define XcmsCIExyYFormat
      (XcmsColorFormat) 0x00000003
#define XcmsCIELabFormat
      (XcmsColorFormat)0x0000004
#define XcmsCIELuvFormat
     (XcmsColorFormat)0x0000005
#define XcmsTekHVCFormat
      (XcmsColorFormat)0x0000006
#define XcmsRGBFormat
      (XcmsColorFormat)0x8000000
#define XcmsRGBiFormat
     (XcmsColorFormat)0x80000001
#define XcmsInitNone
                                 0x00
#define XcmsInitSuccess
                                 0x01
typedef unsigned int XcmsColorFormat;
typedef double XcmsFloat;
typedef struct {
     unsigned short red;
     unsigned short green;
     unsigned short blue;
} XcmsRGB;
```

Figure 6-12: <X11/Xcms.h> (continued)

```
typedef struct {
    XcmsFloat red;
    XcmsFloat green;
    XcmsFloat blue;
} XcmsRGBi;
typedef struct {
    XcmsFloat X;
    XcmsFloat Y;
    XcmsFloat Z;
} XcmsCIEXYZ;
typedef struct {
    XcmsFloat u_prime;
    XcmsFloat v_prime;
    XcmsFloat Y;
} XcmsCIEuvY;
typedef struct {
    XcmsFloat x;
    XcmsFloat y;
    XcmsFloat Y;
} XcmsCIExyY;
typedef struct {
    XcmsFloat L_star;
    XcmsFloat a_star;
    XcmsFloat b_star;
} XcmsCIELab;
```

Figure 6-12: <X11/Xcms.h> (continued)

```
typedef struct {
        XcmsFloat L_star;
        XcmsFloat u_star;
        XcmsFloat v_star;
} XcmsCIELuv;
typedef struct {
        XcmsFloat H;
        XcmsFloat V;
        XcmsFloat C;
} XcmsTekHVC;
typedef struct {
        XcmsFloat pad0;
        XcmsFloat pad1;
        XcmsFloat pad2;
        XcmsFloat pad3;
} XcmsPad;
```

Figure 6-12: <X11/Xcms.h> (continued)

```
typedef struct {
    union {
           XcmsRGB
                              RGB;
           XcmsRGBi
                              RGBi;
           XcmsCIEXYZ
                              CIEXYZ;
           XcmsCIEuvY
                              CIEuvY;
           XcmsCIExyY
                              CIExyY;
           XcmsCIELab
                              CIELab;
           XcmsCIELuv
                              CIELuv;
           XcmsTekHVC
                              TekHVC;
           XcmsPad
                              Pad;
           spec;
    unsigned long
                              pixel;
    XcmsColorFormat
                              format;
} XcmsColor;
typedef struct {
           XcmsColor
                              screenWhitePt;
           XPointer
                              functionSet;
           XPointer
                             screenData;
           unsigned char
                             state;
           char
                              pad[3];
} XcmsPerScrnInfo;
typedef void *XcmsCCC;
typedef Status (*XcmsConversionProc)();
typedef XcmsConversionProc *XcmsFuncListPtr;
```

Figure 6-12: <X11/Xcms.h> (continued)

```
typedef struct {
    char
                             *prefix;
    XcmsColorFormat
                             id;
    XcmsParseStringProc
                             parseString;
    XcmsFuncListPtr
                             to_CIEXYZ;
    XcmsFuncListPtr
                             from_CIEXYZ;
    int
                             inverse_flag;
} XcmsColorSpace;
typedef struct {
    XcmsColorSpace
                             **DDColorSpaces;
    XcmsScreenInitProc
                             screenInitProc;
    XcmsScreenFreeProc
                             screenFreeProc;
} XcmsFunctionSet;
```

Figure 6-13: <X11/Xlib.h>

```
typedef char *XPointer;

#define Bool int
#define Status int
#define True 1
#define False 0
#define QueuedAlready 0
#define QueuedAfterReading 1
#define QueuedAfterFlush 2

#define AllPlanes ((unsigned long)~OL)
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typdef void XExtData;

typdef void XExtCodes;

typedef struct {
    int depth;
    int bits_per_pixel;
    int scanline_pad;
} XPixmapFormatValues;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
   int function;
   unsigned long plane_mask;
   unsigned long foreground;
   unsigned long background;
   int line_width;
   int line_style;
   int cap_style;
   int join_style;
   int fill_style;
   int fill_rule;
   int arc_mode;
   Pixmap tile;
   Pixmap stipple;
   int ts_x_origin;
   int ts_y_origin;
   Font font;
   int subwindow_mode;
   Bool graphics_exposures;
   int clip_x_origin;
   int clip_y_origin;
   Pixmap clip_mask;
   int dash_offset;
   char dashes;
} XGCValues;
typedef void GC;
typedef void Visual;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef void Screen;
typedef struct {
      Pixmap background_pixmap;
      unsigned long background_pixel;
      Pixmap border_pixmap;
      unsigned long border_pixel;
      int bit_gravity;
      int win_gravity;
      int backing_store;
      unsigned long backing_planes;
      unsigned long backing_pixel;
      Bool save_under;
      long event_mask;
      long do_not_propagate_mask;
      Bool override_redirect;
      Colormap colormap;
      Cursor cursor;
} XSetWindowAttributes;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
     XExtData *ext_data;
      int depth;
      int bits_per_pixel;
      int scanline_pad;
} ScreenFormat;
typedef struct {
      int x, y;
      int width, height;
      int border_width;
      int depth;
      Visual *visual;
      Window root;
      int class;
      int bit_gravity;
      int win_gravity;
      int backing_store;
      unsigned long backing_planes;
      unsigned long backing_pixel;
      Bool save_under;
      Colormap colormap;
      Bool map_installed;
      int map_state;
      long all_event_masks;
      long your_event_mask;
      long do_not_propagate_mask;
      Bool override_redirect;
      Screen *screen;
} XWindowAttributes;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int family;
      int length;
      char *address;
} XHostAddress;
typedef struct _XImage {
      int width, height;
      int xoffset;
      int format;
      char *data;
      int byte_order;
      int bitmap_unit;
      int bitmap_bit_order;
      int bitmap_pad;
      int depth;
      int bytes_per_line;
      int bits_per_pixel;
      unsigned long red_mask;
      unsigned long green_mask;
      unsigned long blue_mask;
      XPointer obdata;
      struct funcs {
            struct _XImage *(*create_image)();
            int (*destroy_image)();
            unsigned long (*get_pixel)();
            int (*put_pixel)();
            struct _XImage *(*sub_image)();
            int (*add_pixel)();
      } f;
} XImage;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int x, y;
      int width, height;
      int border_width;
      Window sibling;
      int stack_mode;
} XWindowChanges;
typedef struct {
     unsigned long pixel;
      unsigned short red, green, blue;
      char flags;
      char pad;
} XColor;
typedef struct {
     short x1, y1, x2, y2;
} XSegment;
typedef struct {
     short x, y;
} XPoint;
typedef struct {
      short x, y;
      unsigned short width, height;
} XRectangle;
typedef struct {
     short x, y;
     unsigned short width, height;
      short angle1, angle2;
} XArc;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
     int key_click_percent;
      int bell_percent;
      int bell_pitch;
      int bell_duration;
      int led;
      int led_mode;
      int key;
      int auto_repeat_mode;
} XKeyboardControl;
typedef struct {
     int key_click_percent;
      int bell_percent;
     unsigned int bell_pitch, bell_duration;
     unsigned long led_mask;
      int global_auto_repeat;
      char auto_repeats[32];
} XKeyboardState;
typedef struct {
     Time time;
     short x, y;
} XTimeCoord;
typedef struct {
      int
                  max_keypermod;
                  *modifiermap;
     KeyCode
} XModifierKeymap;
typedef void Display;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
     unsigned long serial;
      Bool send event;
      Display *display;
      Window window;
      Window root;
      Window subwindow;
      Time time;
      int x, y;
      int x_root, y_root;
      unsigned int state;
      unsigned int keycode;
     Bool same_screen;
} XKeyEvent;
typedef XKeyEvent XKeyPressedEvent;
typedef XKeyEvent XKeyReleasedEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Window root;
      Window subwindow;
      Time time;
      int x, y;
      int x_root, y_root;
      unsigned int state;
      unsigned int button;
      Bool same_screen;
} XButtonEvent;
typedef XButtonEvent XButtonPressedEvent;
typedef XButtonEvent XButtonReleasedEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
     int type;
     unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Window root;
      Window subwindow;
     Time time;
      int x, y;
      int x_root, y_root;
      unsigned int state;
      char is_hint;
     Bool same_screen;
} XMotionEvent;
typedef XMotionEvent XPointerMovedEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Window root;
      Window subwindow;
     Time time;
      int x, y;
      int x_root, y_root;
      int mode;
      int detail;
      Bool same_screen;
      Bool focus;
      unsigned int state;
} XCrossingEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef XCrossingEvent XEnterWindowEvent;
typedef XCrossingEvent XLeaveWindowEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      int mode;
      int detail;
} XFocusChangeEvent;
typedef XFocusChangeEvent XFocusInEvent;
typedef XFocusChangeEvent XFocusOutEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      char key_vector[32];
} XKeymapEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      int x, y;
      int width, height;
      int count;
} XExposeEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
Drawable drawable;
      int x, y;
      int width, height;
      int count;
      int major_code;
      int minor_code;
} XGraphicsExposeEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Drawable drawable;
      int major_code;
      int minor_code;
} XNoExposeEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      int state;
} XVisibilityEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window parent;
      Window window;
      int x, y;
      int width, height;
      int border_width;
      Bool override_redirect;
} XCreateWindowEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
} XDestroyWindowEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      Bool from_configure;
} XUnmapEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      Bool override_redirect;
} XMapEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window parent;
      Window window;
} XMapRequestEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      Window parent;
      int x, y;
      Bool override_redirect;
} XReparentEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      int x, y;
      int width, height;
      int border width;
      Window above;
      Bool override_redirect;
} XConfigureEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      int x, y;
} XGravityEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      int width, height;
} XResizeRequestEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window parent;
      Window window;
      int x, y;
      int width, height;
      int border_width;
      Window above;
      int detail;
      unsigned long value_mask;
} XConfigureRequestEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window event;
      Window window;
      int place;
} XCirculateEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window parent;
      Window window;
      int place;
} XCirculateRequestEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Atom atom;
      Time time;
      int state;
} XPropertyEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Atom selection;
      Time time;
} XSelectionClearEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window owner;
      Window requestor;
      Atom selection;
      Atom target;
      Atom property;
      Time time;
} XSelectionRequestEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
     Display *display;
     Window requestor;
     Atom selection;
     Atom target;
     Atom property;
     Time time;
} XSelectionEvent;
typedef struct {
      int type;
     Display *display;
     XID resourceid;
     unsigned long serial;
     unsigned char error_code;
     unsigned char request_code;
     unsigned char minor_code;
} XErrorEvent;
typedef struct {
      int type;
     unsigned long serial;
      Bool send_event;
     Display *display;
      Window window;
     Atom message_type;
      int format;
     union {
            char b[20];
            short s[10];
            long 1[5];
      } data;
} XClientMessageEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      Colormap colormap;
      Bool new;
      int state;
} XColormapEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
      int request;
      int first_keycode;
      int count;
} XMappingEvent;
typedef struct {
      int type;
      unsigned long serial;
      Bool send_event;
      Display *display;
      Window window;
} XAnyEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef union _XEvent {
   int
                           type;
  XAnyEvent
                           xany;
  XKeyEvent
                           xkey;
   XButtonEvent
                           xbutton;
  XMotionEvent
                           xmotion;
  XCrossingEvent
                           xcrossing;
   XFocusChangeEvent
                           xfocus;
  XExposeEvent
                           xexpose;
  XGraphicsExposeEvent
                           xgraphicsexpose;
   XNoExposeEvent
                           xnoexpose;
  XVisibilityEvent
                           xvisibility;
  XCreateWindowEvent
                           xcreatewindow;
   XDestroyWindowEvent
                           xdestroywindow;
   XUnmapEvent
                           xunmap;
   XMapEvent
                           xmap;
  XMapRequestEvent
                           xmaprequest;
   XReparentEvent
                           xreparent;
   XConfigureEvent
                           xconfigure;
  XGravityEvent
                           xgravity;
   XResizeRequestEvent
                           xresizerequest;
  XConfigureRequestEvent xconfigurerequest;
  XCirculateEvent
                           xcirculate;
   XCirculateRequestEvent xcirculaterequest;
  XPropertyEvent
                           xproperty;
  XSelectionClearEvent
                           xselectionclear;
   XSelectionRequestEvent xselectionrequest;
  XSelectionEvent
                           xselection;
  XColormapEvent
                           xcolormap;
                           xclient;
   XClientMessageEvent
  XMappingEvent
                           xmapping;
   XErrorEvent
                           xerror;
                           xkeymap;
   XKeymapEvent
               long
                           pad[24];
   } XEvent;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
#define XAllocID(dpy) ((*(dpy)->resource_alloc)((dpy)))
typedef struct {
      short
                  lbearing;
      short
                  rbearing;
                  width;
      short
      short
                  ascent;
      short
                  descent;
      unsigned short attributes;
} XCharStruct;
typedef struct {
      Atom name;
      unsigned long card32;
} XFontProp;
typedef struct {
      XExtData
                  *ext_data;
      Font
                        fid;
                  direction;
      unsigned
      unsigned min_char_or_byte2;
      unsigned
                 max_char_or_byte2;
      unsigned
                  min_byte1;
      unsigned
                  max_bytel;
      Bool
                        all_chars_exist;
      unsigned
                  default_char;
      int
                        n_properties;
      XFontProp
                        *properties;
      XCharStruct
                        min_bounds;
      XCharStruct
                        max_bounds;
                        *per_char;
      XCharStruct
      int
                        ascent;
      int
                        descent;
} XFontStruct;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      char *chars;
      int nchars;
      int delta;
      Font font;
} XTextItem;
typedef struct {
      unsigned char byte1;
      unsigned char byte2;
} XChar2b;
typedef struct {
      XChar2b *chars;
      int nchars;
      int delta;
      Font font;
} XTextItem16;
typedef union {
      Display *display;
      GC gc;
      Visual *visual;
      Screen *screen;
      ScreenFormat *pixmap_format;
      XFontStruct *font;
} XEDataObject;
typedef struct {
                      max_ink_extent;
max_logical_extent;
      XRectangle
      XRectangle
} XFontSetExtents;
typedef void XFontSet;
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct {
      char
                       *chars;
      int
                       nchars;
                       delta;
      XFontSet *font_set;
} XmbTextItem;
typedef struct {
     wchar_t *chars;
      int
                   nchars;
      int
                       delta;
      XFontSet font_set;
} XwcTextItem;
typedef void (*XIMProc)();
typedef void
                  XIM;
typedef void
                  XIC;
typedef unsigned long XIMStyle;
typedef struct {
     unsigned short count_styles;
      XIMStyle *supported_styles;
} XIMStyles;
#define XIMPreeditArea
                                0x0001L
#define XIMPreeditCallbacks
                                0x0002L
#define XIMPreeditPosition
                                0 \times 0004 L
#define XIMPreeditNothing
                               0 \times 00008 L
#define XIMPreeditNone
                                0x0010L
#define XIMStatusArea
                                0x0100L
#define XIMStatusCallbacks
                               0 \times 0200 L
#define XIMStatusNothing
                                0x0400L
#define XIMStatusNone
                                0x0800L
```

Figure 6-13: <X11/Xlib.h> (continued)

```
#define XNVaNestedList
                                  "XNVaNestedList"
#define XNQueryInputStyle
                                  "queryInputStyle"
#define XNClientWindow
                                  "clientWindow"
#define XNInputStyle
                                  "inputStyle"
#define XNFocusWindow
                                  "focusWindow"
#define XNResourceName
                                  "resourceName"
#define XNResourceClass
                                  "resourceClass"
#define XNGeometryCallback
                                  "geometryCallback"
#define XNFilterEvents
                                  "filterEvents"
#define XNPreeditStartCallback
                                  "preeditStartCallback"
#define XNPreeditDoneCallback
                                  "preeditDoneCallback"
#define XNPreeditDrawCallback
                                  "preeditDrawCallback"
#define XNPreeditCaretCallback
                                  "preeditCaretCallback"
#define XNPreeditAttributes
                                  "preeditAttributes"
#define XNStatusStartCallback
                                  "statusStartCallback"
#define XNStatusDoneCallback
                                  "statusDoneCallback"
#define XNStatusDrawCallback
                                 "statusDrawCallback"
#define XNStatusAttributes
                                 "statusAttributes"
#define XNArea
                                 "area"
#define XNAreaNeeded
                                  "areaNeeded"
                                  "spotLocation"
#define XNSpotLocation
#define XNColormap
                                  "colorMap"
#define XNStdColormap
                                  "stdColorMap"
#define XNForeground
                                  "foreground"
#define XNBackground
                                  "background"
#define XNBackgroundPixmap
                                 "backgroundPixmap"
#define XNFontSet
                                 "fontSet"
#define XNLineSpace
                                  "lineSpace"
#define XNCursor
                                  "cursor"
```

Figure 6-13: <X11/Xlib.h> (continued)

```
#define XBufferOverflow
                                 -1
#define XLookupNone
                                 1
#define XLookupChars
                                 2
#define XLookupKeySym
                                 3
#define XLookupBoth
                                 4
typedef XPointer XVaNestedList;
typedef struct {
      XPointer client_data;
      XIMProc callback;
} XIMCallback;
typedef unsigned long XIMFeedback;
#define XIMReverse
#define XIMUnderline
                                 (1 << 1)
#define XIMHighlight
                                 (1 << 2)
#define XIMPrimary
                                 (1 < < 5)
#define XIMSecondary
                                 (1 < < 6)
#define XIMTertiary
                                 (1 << 7)
typedef struct _XIMText {
      unsigned short length;
      XIMFeedback *feedback;
      Bool encoding_is_wchar;
      union {
            char *multi_byte;
            wchar_t *wide_char;
      } string;
} XIMText
```

Figure 6-13: <X11/Xlib.h> (continued)

```
typedef struct _XIMPreeditDrawCallbackStruct {
      int caret;
      int chg_first;
      int chg_length;
      XIMText *text;
} XIMPreeditDrawCallbackStruct;
typedef enum {
     XIMForwardChar, XIMBackwardChar,
     XIMForwardWord, XIMBackwardWord,
     XIMCaretUp, XIMCaretDown,
     XIMNextLine, XIMPreviousLine,
     XIMLineStart, XIMLineEnd,
     XIMAbsolutePosition,
     XIMDontChange
} XIMCaretDirection;
typedef enum {
     XIMIsInvisible,
     XIMIsPrimary,
     XIMIsSecondary
} XIMCaretStyle;
typedef struct _XIMPreeditCaretCallbackStruct {
     int position;
     XIMCaretDirection direction;
     XIMCaretStyle style;
} XIMPreeditCaretCallbackStruct;
```

Figure 6-14: <X11/Xlib.h> (continued)

```
typedef enum {
    XIMTextType,
    XIMBitmapType
} XIMStatusDataType;

typedef struct _XIMStatusDrawCallbackStruct {
    XIMStatusDataType type;
    union {
        XIMText *text;
        Pixmap bitmap;
    } data;
} XIMStatusDrawCallbackStruct;
```

Figure 6-15: <X11/Xresource.h>

```
typedef int
                  XrmQuark, *XrmQuarkList;
#define NULLQUARK ((XrmQuark) 0)
typedef enum {XrmBindTightly, XrmBindLoosely} \
      XrmBinding, *XrmBindingList;
typedef XrmQuark
                        XrmName;
typedef XrmQuarkList
                        XrmNameList;
typedef XrmQuark
                        XrmClass;
typedef XrmQuarkList
                        XrmClassList;
typedef XrmQuark
                        XrmRepresentation;
#define XrmStringToName(string)
     XrmStringToQuark(string)
#define XrmStringToNameList(str, name) \
      XrmStringToQuarkList(str, name)
#define XrmStringToClass(class)
      XrmStringToQuark(class)
#define XrmStringToClassList(str,class) \
      XrmStringToQuarkList(str, class)
#define XrmStringToRepresentation(string) \
      XrmStringToQuark(string)
typedef struct {
      unsigned int
                        size;
      XPointer
                        addr;
} XrmValue, *XrmValuePtr;
typedef void
                        XrmHashBucket;
typedef XrmHashBucket
                        *XrmHashTable;
typedef XrmHashTable
                        XrmSearchList[];
typedef void
                        XrmDatabase;
#define XrmEnumAllLevels
                                 0
#define XrmEnumOneLevel
                                 1
```

Figure 6-15: <X11/Xresource.h> (continued)

```
typedef enum {
     XrmoptionNoArg,
     XrmoptionIsArg,
     XrmoptionStickyArg,
     XrmoptionSepArg,
     XrmoptionResArg,
     XrmoptionSkipArg,
     XrmoptionSkipLine,
     XrmoptionSkipNArgs
} XrmOptionKind;
typedef struct {
     char
                                 *option;
      char
                                 *specifier;
      XrmOptionKind
                                 argKind;
     XPointer
                                 value;
} XrmOptionDescRec, *XrmOptionDescList;
```

Figure 6-16: <X11/Xutil.h>

```
#define NoValue
                         0x0000
#define XValue
                         0 \times 0001
                         0x0002
#define YValue
#define WidthValue
                         0 \times 0004
#define HeightValue
                         0x0008
#define AllValues
                         0x000F
#define XNegative
                         0 \times 0010
#define YNegative
                         0 \times 0020
typedef struct {
      long flags;
      int x, y;
      int width, height;
      int min_width, min_height;
      int max_width, max_height;
      int width_inc, height_inc;
      struct {
             int x;
            int y;
      } min_aspect, max_aspect;
      int base_width, base_height;
      int win_gravity;
} XSizeHints;
#define USPosition
                        (1L << 0)
#define USSize
                         (1L << 1)
#define PPosition
                        (1L << 2)
#define PSize
                        (1L << 3)
#define PMinSize
                        (1L << 4)
                        (1L << 5)
#define PMaxSize
#define PResizeInc
                        (1L << 6)
#define PAspect
                         (1L << 7)
#define PBaseSize (1L << 8)
#define PWinGravity (1L << 9)</pre>
#define PAllHints (PPosition|PSize|PMinSize| \
      PMaxSize|PResizeInc|PAspect)
```

Figure 6-16: <X11/Xutil.h> (continued)

```
typedef
                        struct {
                        flags;
            long
            Bool
                        input;
            int
                        initial_state;
            Pixmap
                        icon_pixmap;
            Window
                        icon_window;
            int
                        icon_x, icon_y;
                        icon_mask;
            Pixmap
            XID
                        window_group;
      } XWMHints;
      #define InputHint
                                       (1L << 0)
      #define StateHint
                                       (1L << 1)
      #define IconPixmapHint
                                       (1L << 2)
                                       (1L << 3)
      #define IconWindowHint
      #define IconPositionHint
                                      (1L << 4)
      #define IconMaskHint
                                       (1L << 5)
      #define WindowGroupHint
                                       (1L << 6)
      #define AllHints (InputHint|StateHint|
            IconPixmapHint | IconWindowHint |
            IconPositionHint | Icon-
MaskHint | WindowGroupHint)
      #define WithdrawnState
                                       0
      #define NormalState
                                       1
      #define IconicState
      typedef struct {
                                       *value;
            unsigned char
            Atom
                                       encoding;
            int
                                       format;
            unsigned long
                                       nitems;
      } XTextProperty;
                                       -1
      #define XNoMemory
      #define XLocaleNotSupported
                                       -2
      #define XConverterNotFound
                                        -3
```

Figure 6-16: <X11/Xutil.h> (continued)

```
typedef int XContext;
  typedef enum {
    XStringStyle,
   XCompoundTextStyle,
   XTextStyle,
   XStdICCTextStyle
  } XICCEncodingStyle;
  typedef struct {
    int min_width, min_height;
    int max_width, max_height;
    int width_inc, height_inc;
  } XIconSize;
  typedef struct {
    char *res name;
    char *res class;
  } XClassHint;
  #define XDestroyImage(ximage)
    ((*((ximage)->f.destroy_image))((ximage)))
  #define XGetPixel(ximage, x, y)
    ((*((ximage)->f.get_pixel))((ximage), (x), (y)))
  #define XPutPixel(ximage, x, y, pixel)
    ((*((ximage)->f.put_pixel))((ximage), (x),
    (y), (pixel)))
  #define XSubImage(ximage, x, y, width, height)
    ((*((ximage)->f.sub_image))((ximage),
    (x), (y), (width), (height)))
  #define XAddPixel(ximage, value)
    ((*((ximage)->f.add_pixel))((ximage),
    (value)))
  typedef struct _XComposeStatus {
    XPointer compose ptr;
    int chars_matched;
  } XComposeStatus;
```

Figure 6-16: <X11/Xutil.h> (continued)

```
#define IsKeypadKey(keysym)
      (((unsigned)(keysym) >= XK_KP_Space) && \
      ((unsigned)(keysym) <= XK_KP_Equal))
#define IsCursorKey(keysym)
      (((unsigned)(keysym) >= XK Home) && \
      ((unsigned)(keysym) < XK_Select))</pre>
#define IsPFKey(keysym)
      (((unsigned)(keysym) >= XK_KP_F1) \
      && ((unsigned)(keysym) <= XK_KP_F4))
#define IsFunctionKey(keysym)
      (((unsigned)(keysym) >= XK_F1) && \
((unsigned)(keysym) <= XK_F35))
#define IsMiscFunctionKey(keysym)
      (((unsigned)(keysym) >= XK_Select) && \
      ((unsigned)(keysym) <= XK_Break))
#define IsModifierKey(keysym)
      ((((unsigned)(keysym) >= XK Shift L) \
      && ((unsigned)(keysym) <= XK_Hyper_R))
      | ((unsigned)(keysym) == XK_Mode_switch)
      ((unsigned)(keysym) == XK_Num_Lock))
typedef void Region;
#define RectangleOut
#define RectangleIn
                        1
#define RectanglePart
typedef struct {
      Visual *visual;
      VisualID visualid;
      int
           screen;
            depth;
      int
      int
           class;
      unsigned long red_mask;
      unsigned long green_mask;
      unsigned long blue_mask;
      int
            colormap_size;
      int
            bits_per_rgb;
 XVisualInfo;
```

Figure 6-16: <X11/Xutil.h> (continued)

```
#define VisualNoMask
                                 0 \times 0
#define VisualIDMask
                                 0x1
#define VisualScreenMask
                                 0x2
#define VisualDepthMask
                                 0x4
#define VisualClassMask
                                 0x8
#define VisualRedMaskMask
                                 0x10
#define VisualGreenMaskMask
                                 0x20
#define VisualBlueMaskMask
                                 0x40
#define VisualColormapSizeMask
                                 0x80
#define VisualBitsPerRGBMask
                                 0x100
#define VisualAllMask
                                 0x1FF
typedef struct {
     Colormap
                                 colormap;
      unsigned long
                                 red max;
     unsigned long
                                red_mult;
      unsigned long
                                green_max;
     unsigned long
                                 green_mult;
     unsigned long
                                 blue max;
      unsigned long
                                 blue_mult;
      unsigned long
                                 base pixel;
      VisualID
                                 visualid;
      XID
                                 killid;
} XStandardColormap;
#define ReleaseByFreeingColormap ((XID) 1L)
#define BitmapSuccess
#define BitmapOpenFailed
                                 1
#define BitmapFileInvalid
                                 2
#define BitmapNoMemory
                                 3
#define XCSUCCESS
                                 0
#define XCNOMEM
                                 1
#define XCNOENT
                                 2
```

TCP/IP Data Definitions

This section is new, but will not be diffmarked.

Figure 6-17: <netinet/in.h>



This section contains standard data definitions that describe system data for the optional TCP/IP Interfaces. These data definitions are referred to by their names in angle brackets: <name.h> and <sys/name.h>. Included in these data definitions are macro definitions and structure definitions. While an ABI-conforming system may provide TCP/IP interfaces, it need not contain the actual data definitions referenced here. Programmers should observe that the sources of the structures defined in these data definitions are defined in SVID.

```
#define
          INADDR_ANY
                           (u_long)0x00000000
#define
          INADDR_LOOPBACK (u_long)0x7F000001
          INADDR_BROADCAST (u_long)0xfffffff
#define
#define
          IPPROTO_TCP
                           6
#define
          IPPROTO IP
                           0
                           1
#define
          IP_OPTIONS
struct in_addr {
      union {
          struct { u_char s_b1,s_b2,s_b3,s_b4; } S_un_b;
          struct { u_short s_w1,s_w2; } S_un_w;
          u_long S_addr;
      } S_un;
#define IN_SET_LOOPBACK_ADDR(a)\
      {(a)->sin_addr.s_addr=htonl(INADDR_LOOPBACK);
struct sockaddr_in {
               sin_family;
      short
      u_short sin_port;
      struct
               in_addr sin_addr;
      char
               sin_zero[8];
};
```

Figure 6-18: <netinet/ip.h>

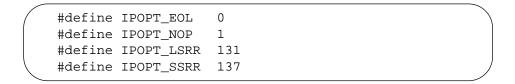
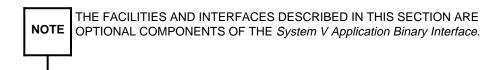


Figure 6-19: <X11/netinet/tcp.h>

#define TCP_NODELAY 0x01

Development Environment

Development Commands





NOTE This chapter is new, but will not be marked with diff-marks.

The Development Environment for MIPS implementations of System V Release 4 will contain all of the development commands required by the System V ABI, namely;

as	CC	ld
m4	lex	vacc

Each command accepts all of the options required by the System V ABI, as defined in the SD_CMD section of the *System V Interface Definition, Third Edition*.

PATH Access to Development Tools

The development environment for the MIPS System V implementations is accessible using the system default value for PATH. The default if no options are given to the cc command is to use the libraries and object file formats that are required for ABI compliance.

Software Packaging Tools

The development environment for MIPS implementations of the System V ABI shall include each of the following commands as defined in the AS_CMD section of the *System V Interface Definition, Third Edition*.

pkgproto pkgtrans pkgmk

System Headers

Systems that do not have an ABI Development Environment may or may not have

DEVELOPMENT ENVIRONMENT

system header files. If an ABI Development Environment is supported, system header files will be included with the Development Environment. The primary source for contents of header files is always the *System V Interface Definition, Third Edition*. In those cases where SVID Third Edition doesn't specify the contents of system headers, Chapter 6 "Data Definitions" of this document shall define the associations of data elements to system headers for compilation. For greatest source portability, applications should depend only on header file contents defined in SVID.

Static Archives

Level 1 interfaces defined in *System V Interface Definition, Third Edition*, for each of the following libraries, may be statically linked safely into applications. The resulting executable will not be made non-compliant to the ABI solely because of the static linkage of such members in the executable.

libm

The archive libm.a is located in /usr/lib on conforming MIPS development environments.

Execution Environment

Application Environment



This chapter is new, but will not be marked with diff-marks.

This section specifies the execution environment information available to application programs running on a MIPS ABI-conforming computer.

The /dev Subtree

All networking device files described in the Generic ABI shall be supported on all MIPS ABI-conforming computers. In addition, the following device files are required to be present on all MIPS ABI-conforming computers.

/dev/null This device file is a special "null" device that may be used to

test programs or provide a data sink. This file is writable by all

processes.

/dev/tty This device file is a special one that directs all output to the

controlling TTY of the current process group. This file is read-

able and writable by all processes.

/dev/sxtXX /dev/ttyXX

These device files, where XX represents a two-digit integer, represent device entries for terminal sessions. All these device files must be examined by the ttyname() call. Applications must not have the device names of individual terminals hard-coded within them. The sxt entries are optional in the system but, if present must be included in the library routine's search.